

Deep Learning

Aymeric DIEULEVEUT

May 2023

1 Goals

2 The first Neural Network: the Perceptron

3 Convolutional Neural Networks

4 Applications

- Image classification
- Pose, action detection
- Object detection and segmentation
- Scene labeling - Semantic segmentation
- Object tracking - videos

Outline

1 Goals

2 The first Neural Network: the Perceptron

3 Convolutional Neural Networks

4 Applications

- Image classification
- Pose, action detection
- Object detection and segmentation
- Scene labeling - Semantic segmentation
- Object tracking - videos

Deep Learning in one slide + Goals

- **How does it work:**

- ▶ Automatically learn representations of observations
- ▶ Learn highly non-linear models.

- **What does it require:**

- ▶ Large datasets with structure
- ▶ Computational power

- **Why now:**

- ▶ Combination of the 2 points above
- ▶ investment !

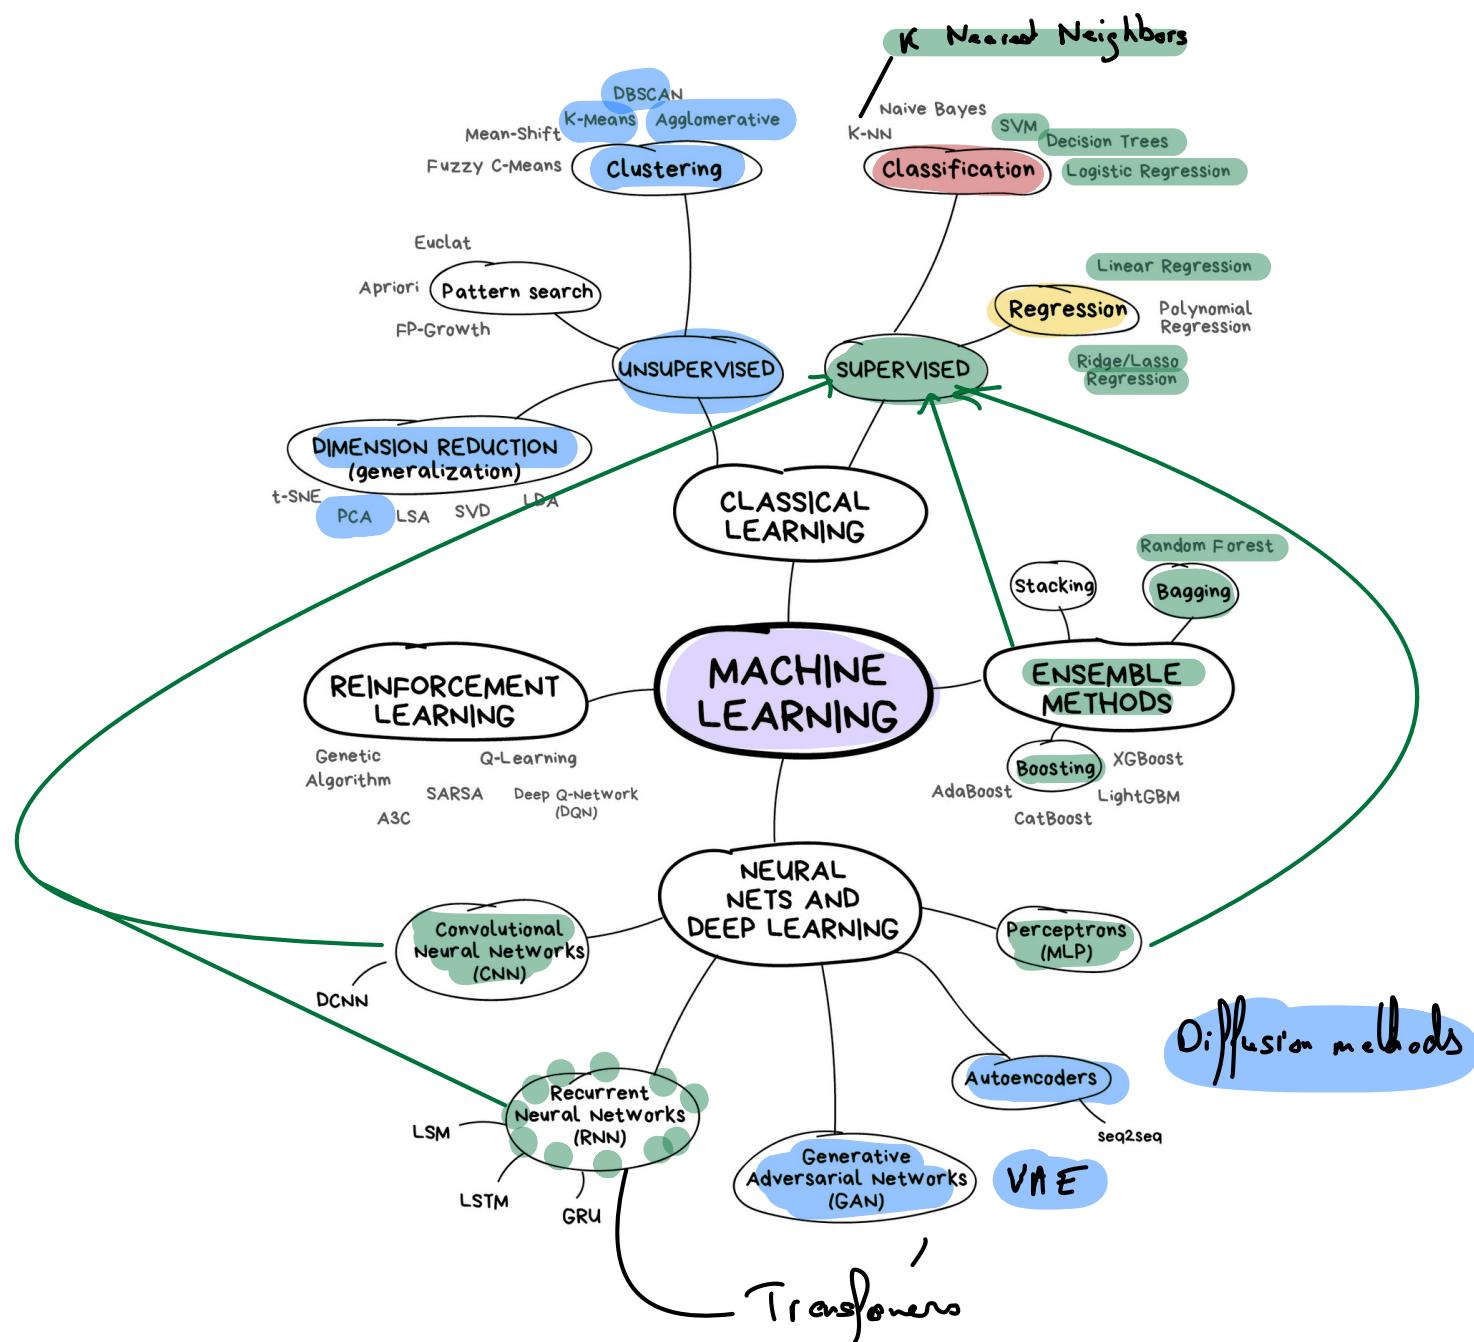
- **Some Applications**

- ▶ Image classification; object / face recognition
- ▶ Self driving cars
- ▶ Automatic Translation, Information extraction
- ▶ Caption Generation
- ▶ Ads, recommendation systems, etc.

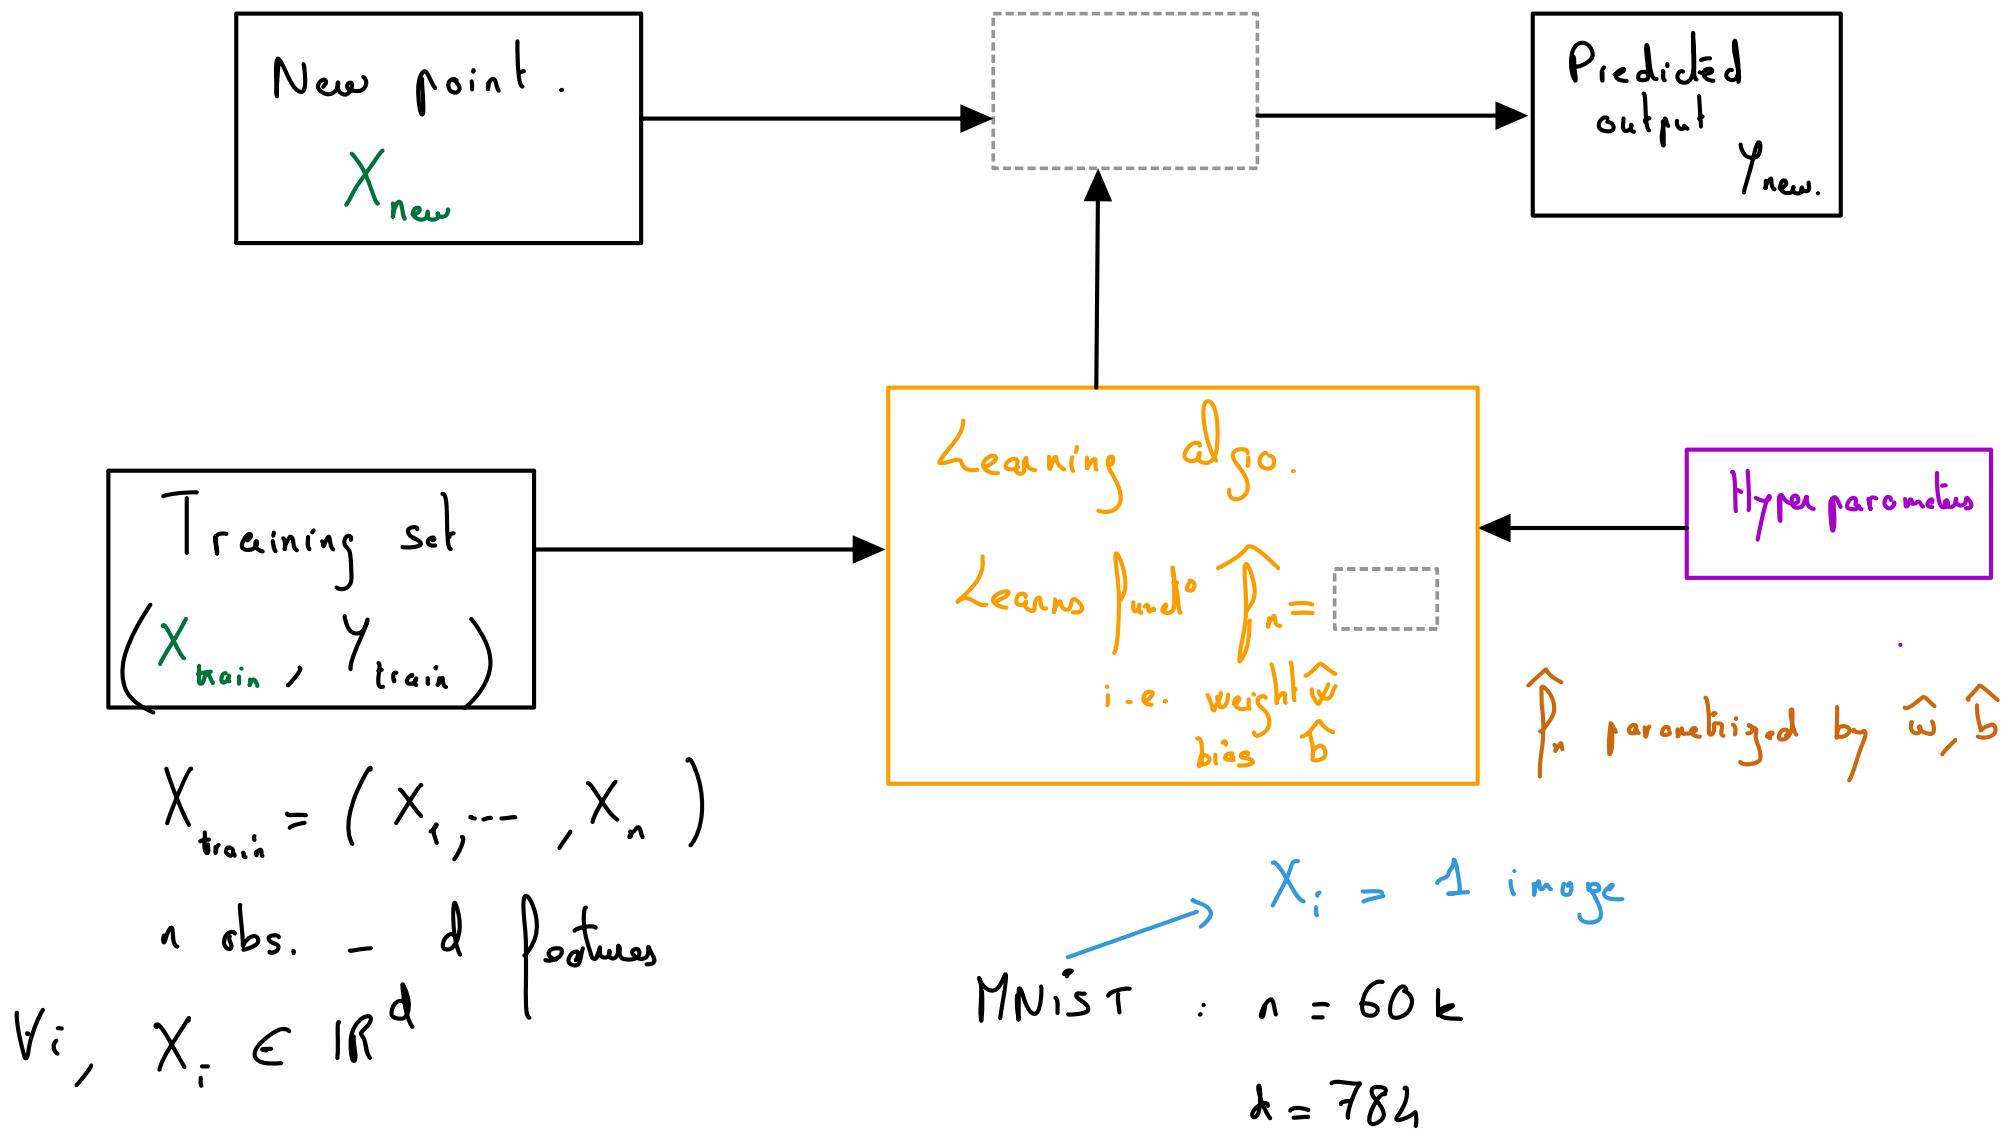
Goals: Understanding core concepts of Deep Learning.

- When and how it works on paper (data, models, architecture)
- How to implement a simple neural network with Python.
- Overview of some of the main applications and challenges.

Machine Learning



Summary of previous lectures on Machine Learning : Supervised ML



set of \square

Objective

Ridge: $LR + \frac{1}{2} \|\mathbf{w}\|_2^2$ reg

Lasso: $LR + \lambda \|\mathbf{w}\|_1$ reg

Linear regression (LR)
 $l(x, y) = (x - y)^2$

Logistic regression
 $l(x, y) = \log(1 + e^{-y})$

SVM

$l(x, y) = \max(0, 1 - y)$

Neural Networks

Linear Functions

$$\hat{f}_{w,b}: X \rightarrow \langle w, X \rangle + b$$

Beyond linear

Functions

Decision Trees

Random Forests

Empirical risk minimization

→ find a function within the class
that performs well on Train!

$$\sum_{i=1}^n l(\hat{f}_{w,b}(x_i), y_i) + \text{Reg}(w)$$

loss. predict output regulariser

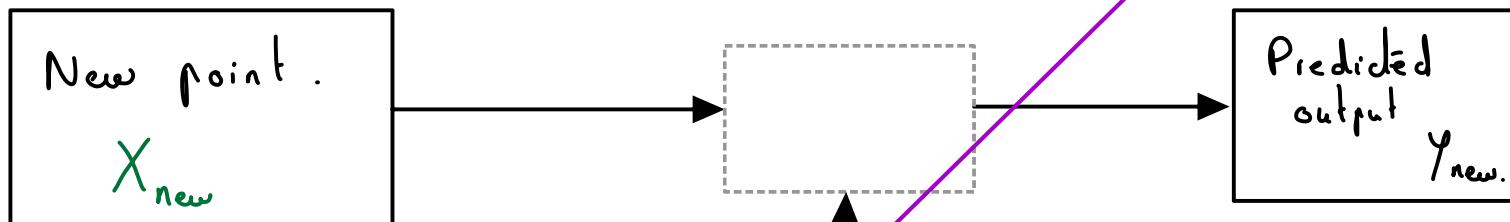
No "global" objective

just incrementally

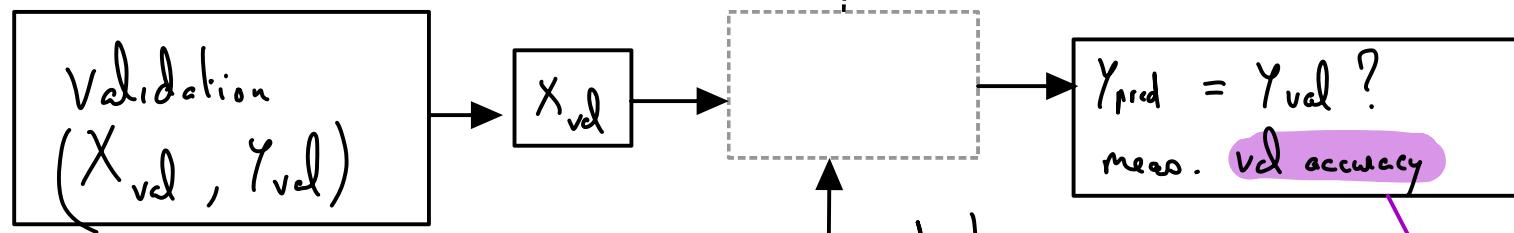
reduce impurity .

Supervised Machine Learning Pipeline

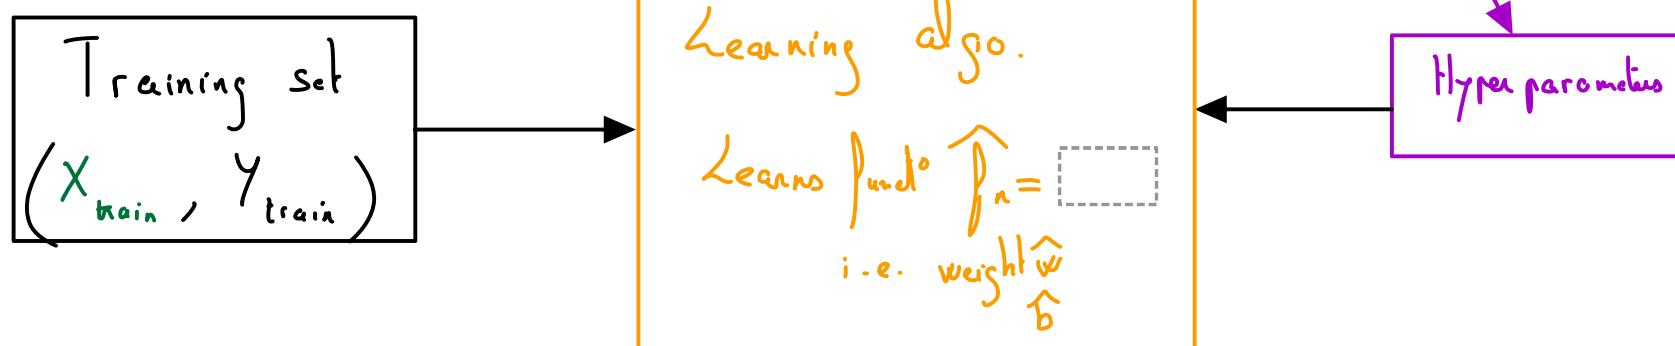
T
E
S
T



V
A
L
I
D



T
R
A
I
N
I
N
G



Outline

1 Goals

2 The first Neural Network: the Perceptron

3 Convolutional Neural Networks

4 Applications

- Image classification
- Pose, action detection
- Object detection and segmentation
- Scene labeling - Semantic segmentation
- Object tracking - videos

From logistic regression to Multi Layer Perceptron 1

Supervised ML problem:

$\in \mathbb{R}$

- ① Data $D_n = (X_i, Y_i)_{i=1,\dots,n}, X_i \in \mathbb{R}^d, Y_i \in \{0, \dots, K\}$
- ② Goal: Predict a the label for a new point.

2 approaches:

Generative approach

- Define a model on your data (e.g., logit model).
- Estimate parameter (likelihood maximization).

Discriminative approach

- No statistical model !
- Define a loss function ℓ
- Goal:

$$\min_{f \in \mathcal{F}} \mathbb{E}[\ell(f(X), Y)]$$

grad output

- Approach:

$$\min_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n [\ell(f(X_i), Y_i)]$$

Empirical risk minimum

From logistic regression to Multi Layer Perceptron 1

Supervised ML problem:

- ① Data $D_n = (X_i, Y_i)_{i=1,\dots,n}, X_i \in \mathbb{R}^d, Y_i \in \{0, \dots, K\}$
- ② Goal: Predict a the label for a new point.

2 approaches:

Generative approach

- Define a model on your data (e.g., logit model).
- Estimate parameter (likelihood maximization).

statistically inspired

Discriminative approach

- No statistical model !
- Define a loss function ℓ
- Goal:

$$\min_{f \in \mathcal{F}} \mathbb{E}[\ell(f(X), Y)]$$

- Approach:

$$\min_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n [\ell(f(X_i), Y_i)]$$

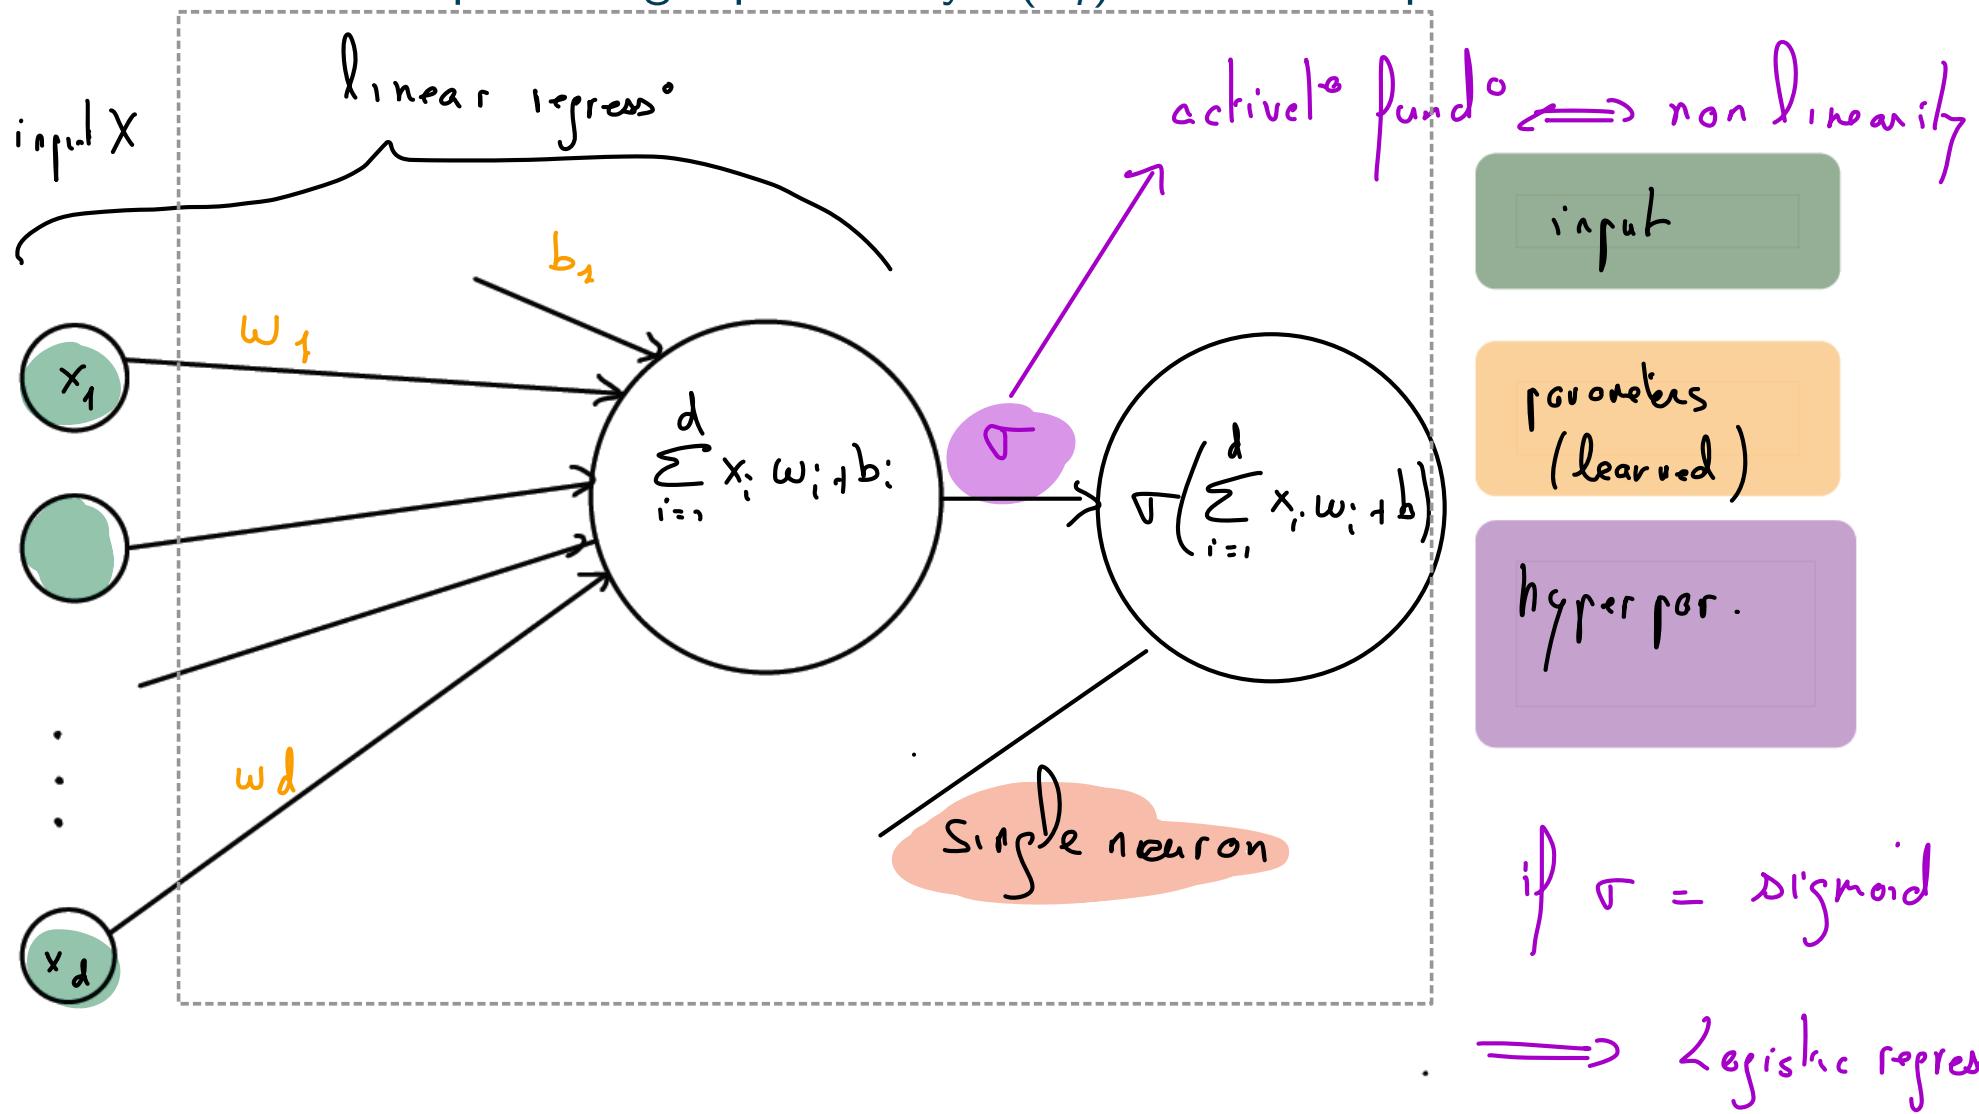
Dominant DL approach

Logistic regression can be seen both ways !

$$\min_{w \in \mathbb{R}^d} \frac{1}{n} \sum_{i=1}^n \log(1 + \exp(-Y_i w^T X_i))$$

From logistic regression to Multi Layer Perceptron 3

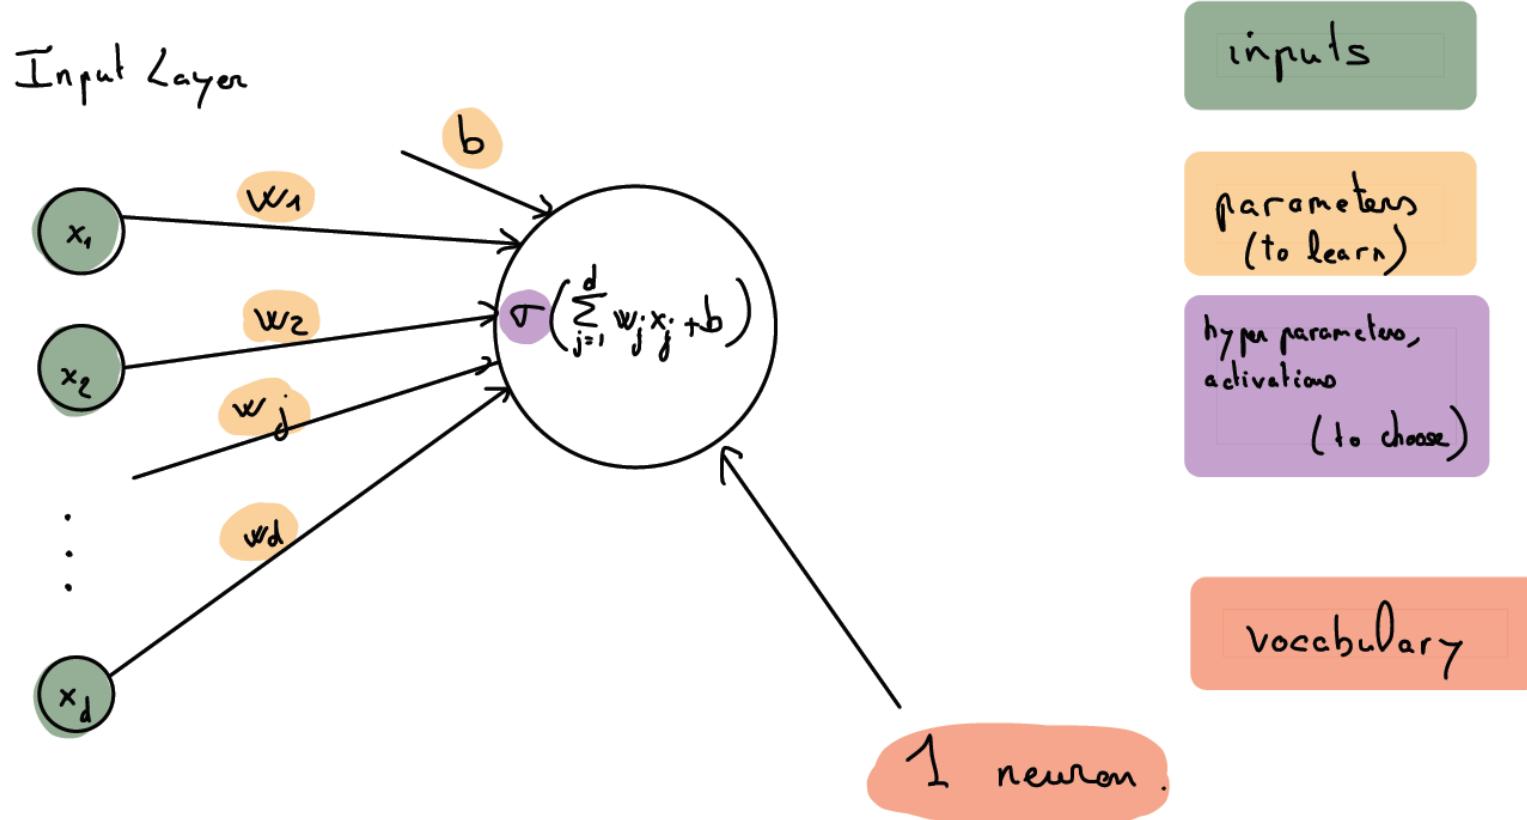
- Class of functions for logistic regression: $\mathcal{F} = \{X \mapsto w^T X, w \in \mathbb{R}^d\}$.
- Prediction for a new point is $1_{\sigma(X_i) > 1/2}$
- Can be seen as predicting a probability $\sigma(X_i)$ that the output is 1



This is a neural network, with one neuron !

From logistic regression to Multi Layer Perceptron 3

- Class of functions for logistic regression: $\mathcal{F} = \{X \mapsto w^T X, w \in \mathbb{R}^d\}$.
- Prediction for a new point is $1_{\sigma(X_i) > 1/2}$
- Can be seen as predicting a probability $\sigma(X_i)$ that the output is 1



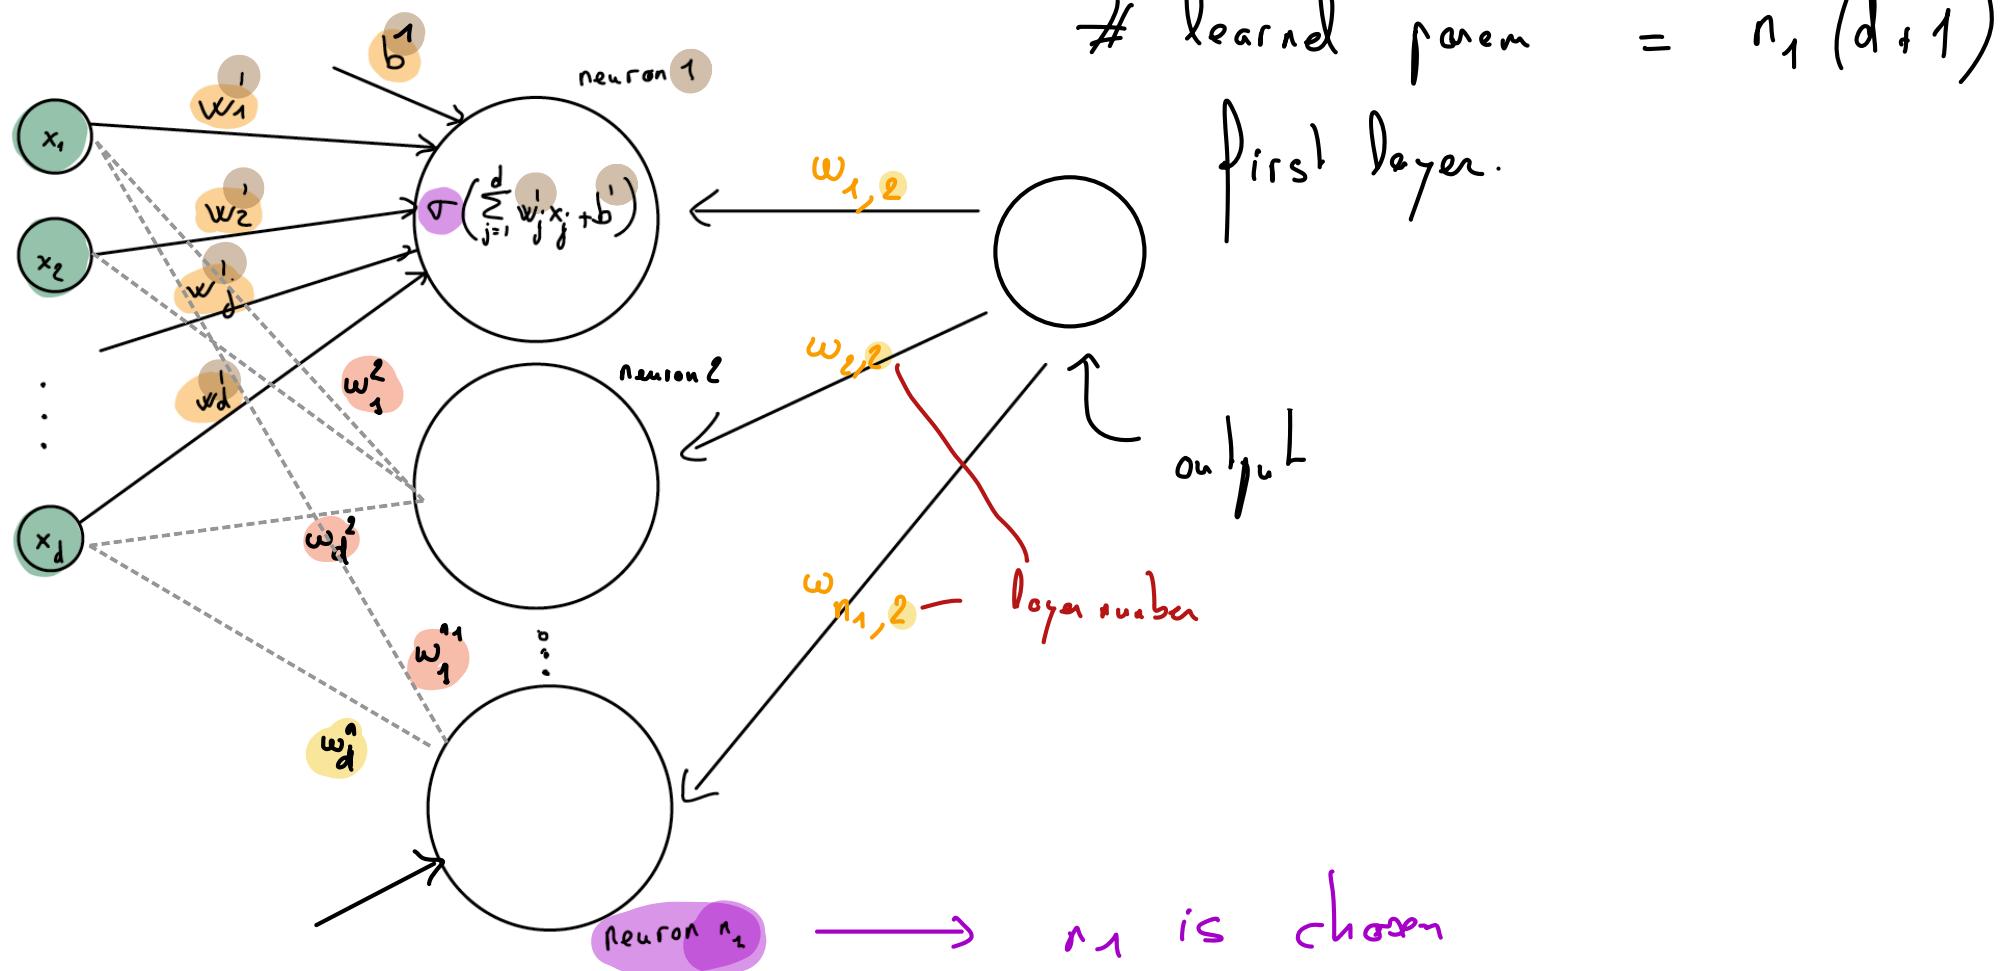
This is a neural network, with one neuron !

From logistic regression to Multi Layer Perceptron 4

Extend to

- Multiple Neurons
- Multiple Layers

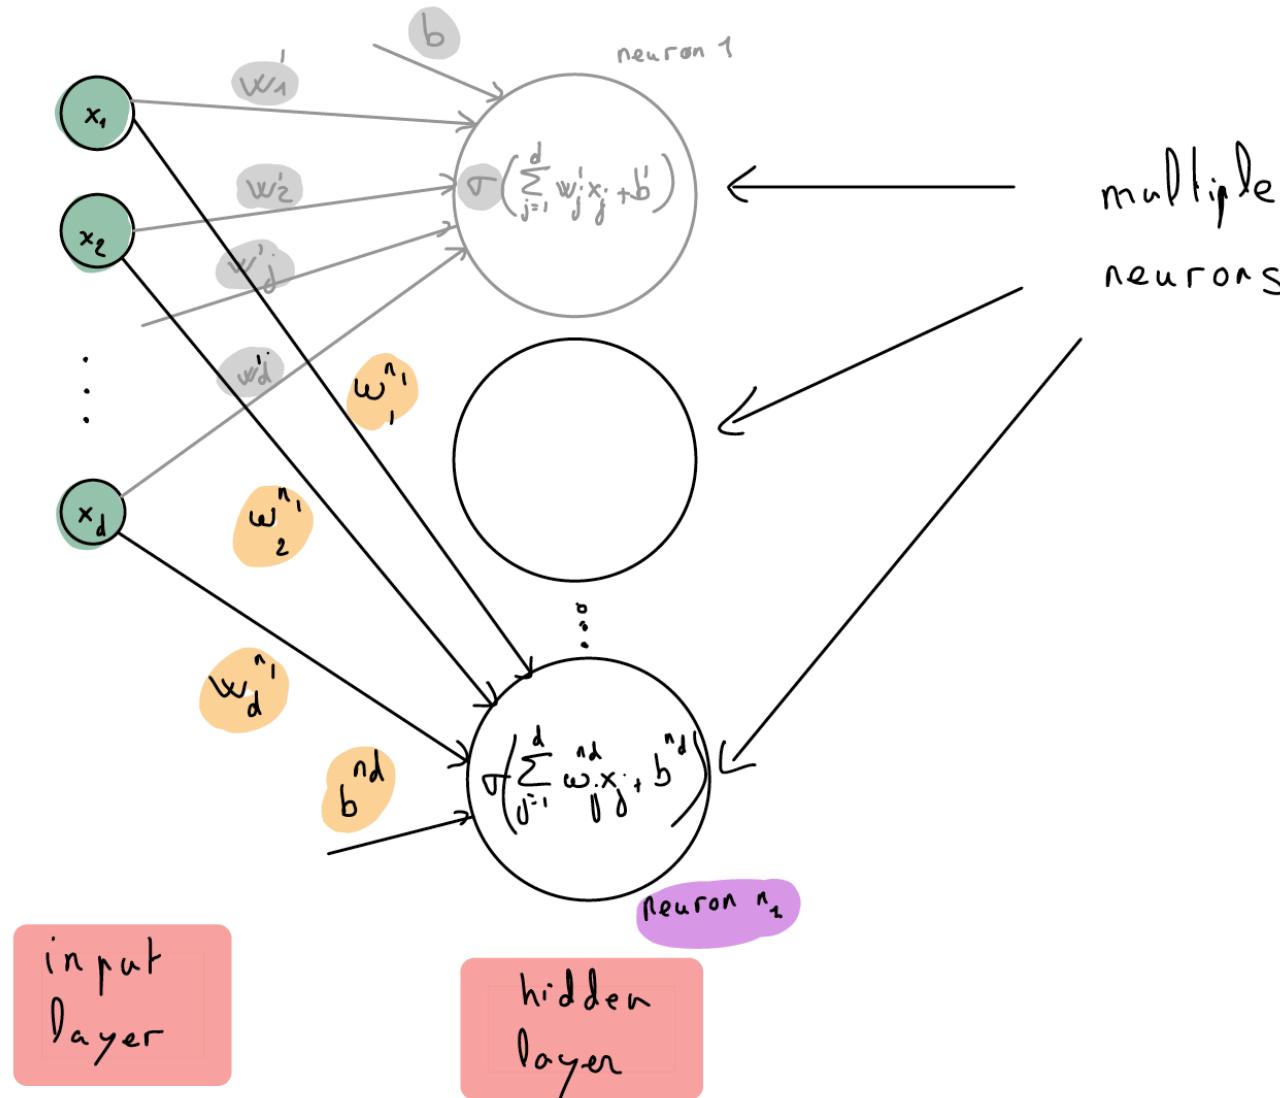
number of biases to be learned n_1
number of weights \dots $n_1 \times d$



From logistic regression to Multi Layer Perceptron 4

Extend to

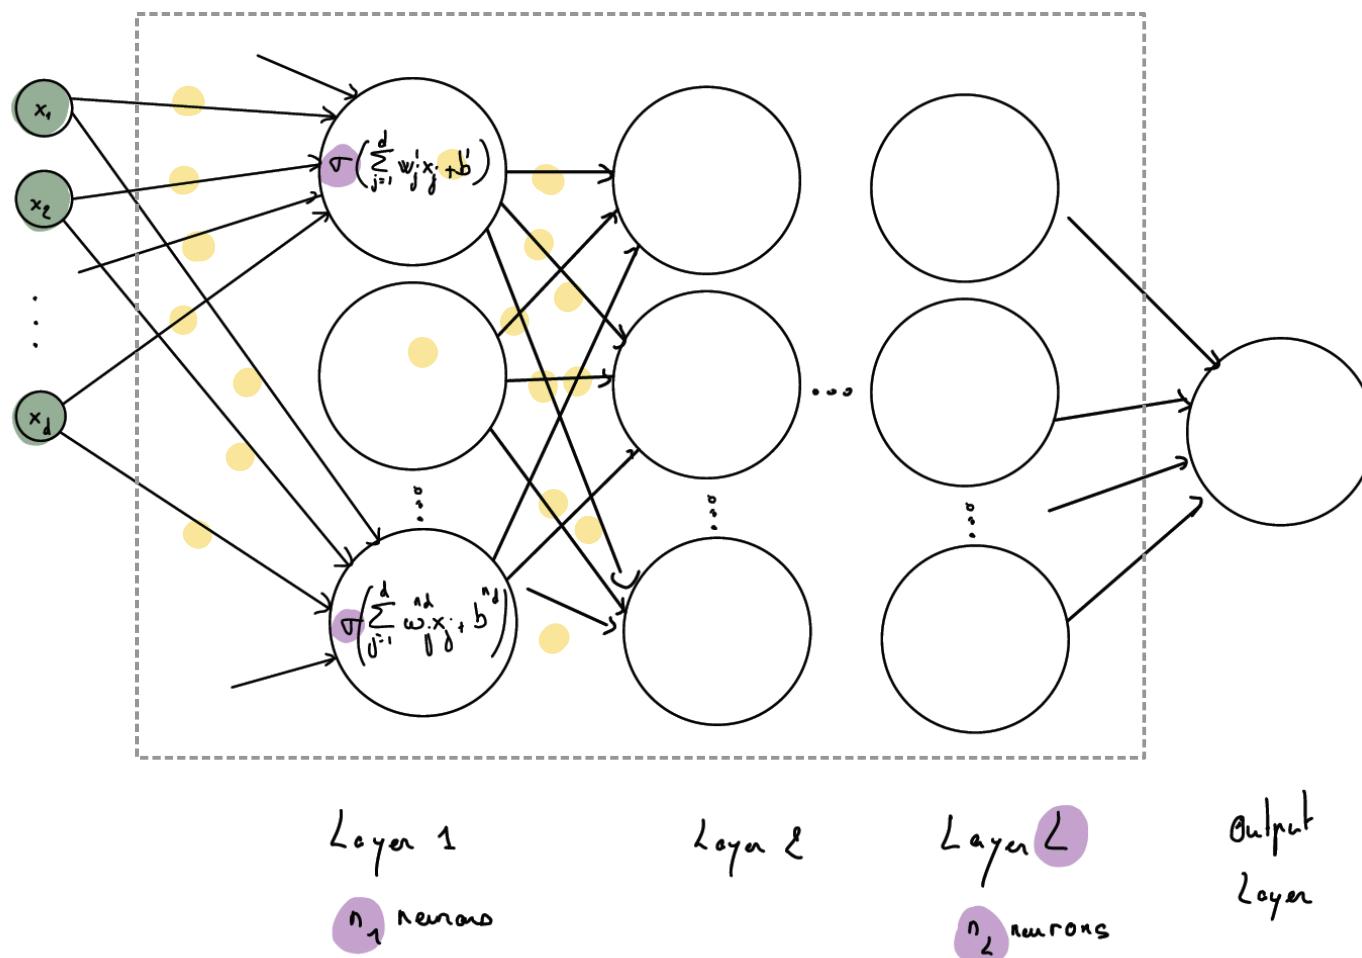
- Multiple Neurons
- Multiple Layers



From logistic regression to Multi Layer Perceptron 5

Extend to

- Multiple Neurons
- Multiple Layers



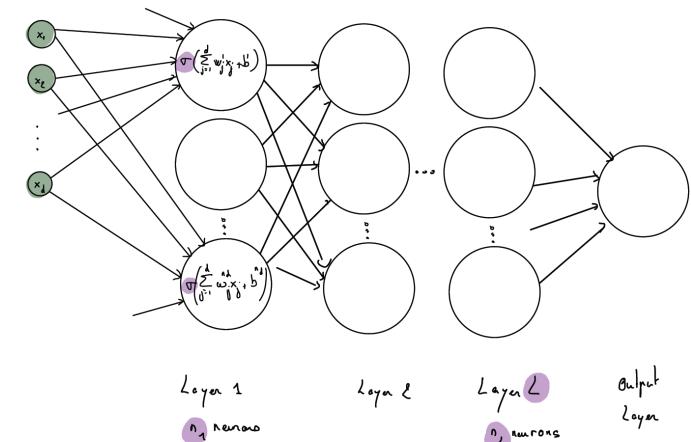
From logistic regression to Multi Layer Perceptron 5

Neural network:

exp. risk minimization

- Goal $\min_{f \in \mathcal{F}_{\text{MLP}}} \frac{1}{n} \sum_{i=1}^n [\ell(f(X_i), Y_i)]$
 - Class of non linear functions.
$$\mathcal{F}_{\text{MLP}} = \{f(X; W, b) = \sigma(b_n + W_n \sigma(\dots (b_1 + W_1 x)))\}$$

$$W_i \in \mathbb{R}^{n_i \times n_{i-1}}, b_i \in \mathbb{R}^{\cdots} \} \quad \text{Learned}$$



Summary: NN generalize regression by looking for candidates in a much larger class of functions than linear ones.

What needs to be learned

What needs to be chosen

Vocabulary

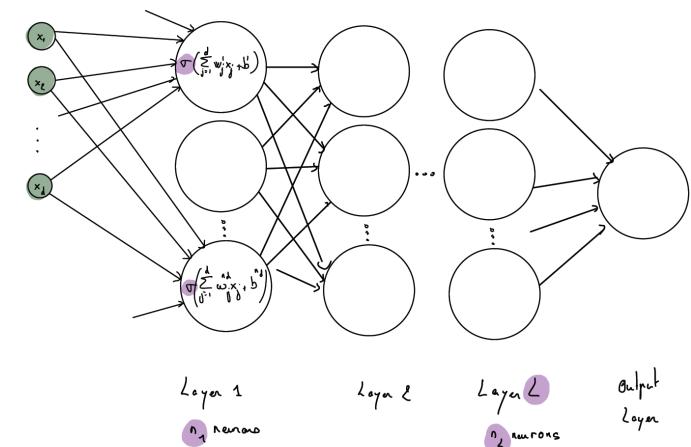
From logistic regression to Multi Layer Perceptron 5

Neural network:

- Goal $\min_{f \in \mathcal{F}_{MLP}} \frac{1}{n} \sum_{i=1}^n [\ell(f(X_i), Y_i)]$

- Class of non linear functions.

$$\mathcal{F}_{MLP} = \{f(X; W, b) = \sigma(b_n + W_n \sigma(\dots(b_1 + W_1 x))), \\ W \in \mathbb{R}^{c \times c}, b \in \mathbb{R}^{c \times 1}\}$$



Summary: NN generalize regression by looking for candidates in a much larger class of functions than linear ones.

What needs to be learned

Parameters W, b

What needs to be chosen

- Activation σ
- Number of layers L
- Number of neurons per hidden layer
 $n_i, i = 1, \dots, L$.

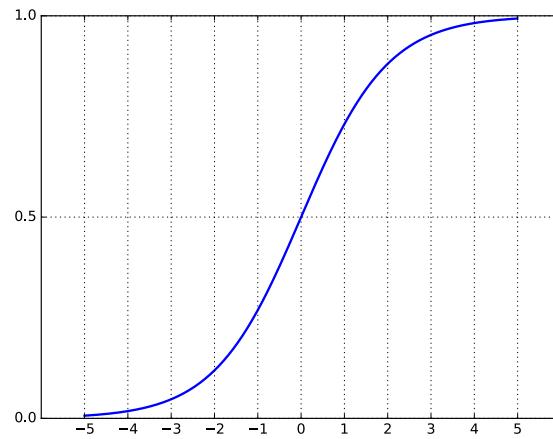
Vocabulary

- Neuron
- Activation
- Hidden Layer
- Width, Depth
- Fully connected layer

2 possible activations: Sigmoid and Rectified Linear Unit

3

x Sigmoid function



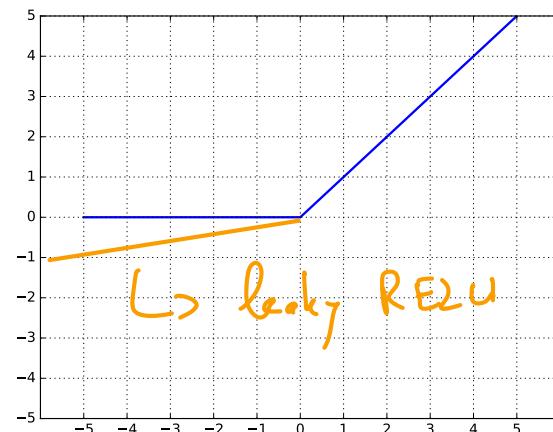
- $x \mapsto \frac{\exp(x)}{1+\exp(x)}$

- Problems:

- ① Saturated function: gradient killer -> (need for rescaling data)

→ dangerous : may fail to learn

x Rectified Linear Unit (ReLU)



- $x \mapsto \max(0, x)$

- Pros & cons:

- ① Not a saturated function. Kills negative values.
- ② Empirically, convergence is faster than sigmoid/tanh.
- ③ Plus: biologically plausible

x Tanh activation \approx Sigmoid
less saturated

How to deal with multi-class output: softmax activation

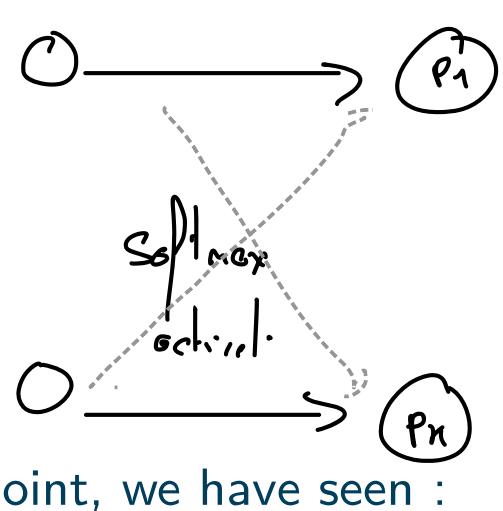
$$\sum_{\text{Last Layer}}^n = k$$

When we do multi-class classification, i.e., $Y_i \in 1, \dots, K$, we output K different values:

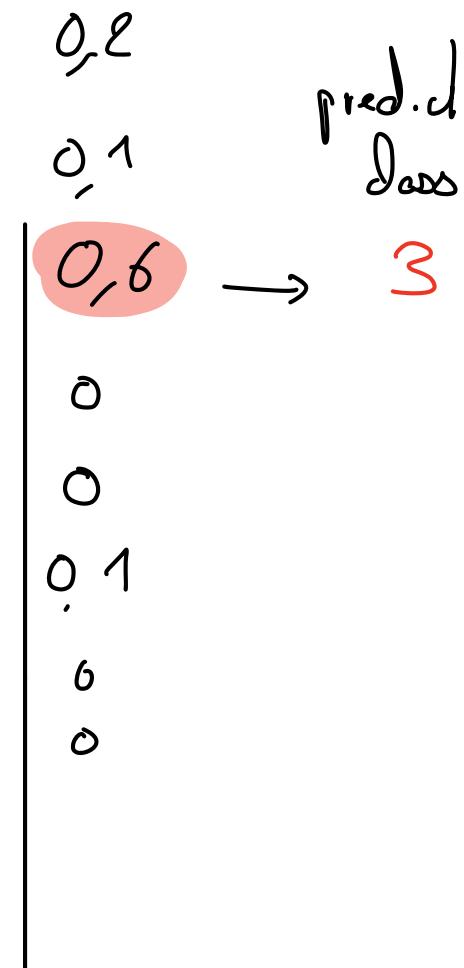
- Each of them corresponds to the probability of belonging to the class j , $1 \leq j \leq K$
- To obtain probabilities (adding up to 1):
 - ① We have K neurons on the last layer
 - ② And use a **Softmax activation** to renormalize.

Softmax output unit, used to predict $\{1, \dots, K\}$:

$$\text{softmax}(z)_i = \frac{e^{z_i}}{\sum_{k=1}^K e^{z_k}}$$



K probabilities



Up to that point, we have seen :

- What a Neural Network was
- Why it is expected to learn better.

Next Question: how to implement it ?

Python - Keras

1. Hardware:

- CPU
- GPU
- TPU



2. Software (Python packages):

- Pytorch/Tensorflow
- → Keras : TF high level API. Ideal for applications.



Keras - A few lines !¹

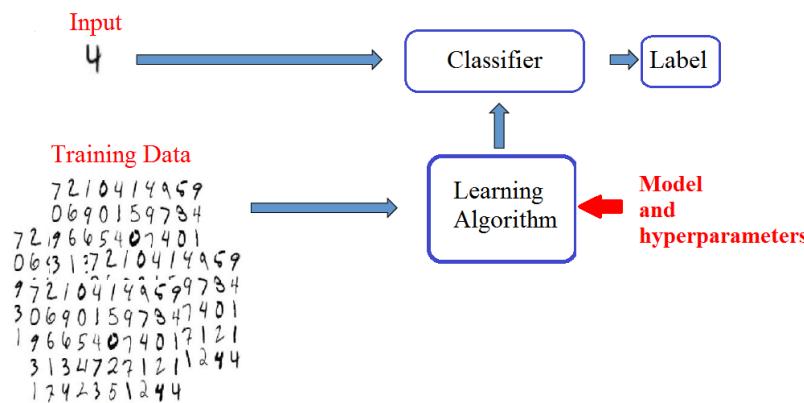
train Fully connected network on image classif.
MNIST

What do we need to specify?

What do we need to do next?

Keras - A few lines !¹

What do we need to specify?



Intuition

Up to that point, we have seen :

- What a Neural Network was.
- How to implement it in Python.

Next Question: why and when does it work ?

You cannot vote anymore

Which of the following things need to be chosen (hyperparameters of the NN)?

1 neuron (13% 5) 5 input size (46% 18)

2 activation functions (95% 37 ✓) 6 number of layers (100% 39 ✓)

3 number of neurons per layer (92% 36 ✓) 7 optimization algorithm (38% 15 ✓)

4 weights (18% 7) 8 biases (8% 3)

wooclap Questions 1 / 10 Messages 100 % Exit Tout afficher

You cannot vote anymore

Which of the following things need to be learned ?

1 neuron (9% 3) 5 input size (9% 3)

2 activation functions (3% 1) 6 number of layers (0% 0)

3 number of neurons per layer (3% 1) 7 optimization algorithm (9% 3)

4 weights (97% 32 ✓) 8 biases (76% 25 ✓)

wooclap Questions 2 / 10 Messages 100 % Exit Tout afficher

The choice of the optim algo matters!

You cannot vote anymore

What happens here?

1 Logistic Regression works (data is linearly separable) (58% 15 ✓) 5 I should change the activation function (4% 1)

2 I should add more neurons (0% 0) 6 Logistic regression fails (21% 5)

3 The network overfits (4% 1) 7 One hidden layer is sufficient (4% 1)

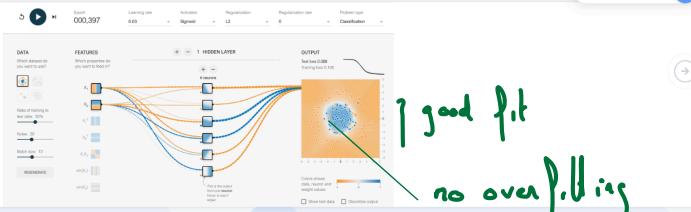
4 I should add more layers (0% 0) 8 I should change to a non-linear activation (25% 6)

wooclap Questions 3 / 10 Messages 100 % Exit Tout afficher

wooclap Questions 4 / 10 Messages 100 % Exit Tout afficher

Train error
 $0.67 = 50\%$
 under fitting!

What happens here?



good fit
no overfitting

Logistic Regression works (data is linearly separable)

9%

3

I should change the activation function

12%

4

I should use more neurons

6%

2

The network overfits

9%

3

I should add more layers

36%

12

Logistic Reg would fail

wooclap

Questions

6 / 10

Messages



100 %



Exit

39% correct

33 / 45



Tout afficher



x



x



x



x



x



x



x



x



x



x



x



x



x



x



x



x



x



x



x



x



x



x



x



x



x



x



x



x



x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

x

What happens here?

You cannot vote anymore

DATA: Which dataset do you want to use?
FEATURES: Which properties do you want to predict?
HIDDEN LAYERS: 4 HIDDEN LAYERS
OUTPUT: Test loss 0.000, Training loss 0.000

1 Logistic Regression works (data is linearly separable)
2 I should use more neurons
3 The network overfits
4 I should add more layers
5 I should change the activation function
6 Logistic regression fails
7 One hidden layer is sufficient
8 I should change to a non-linear activation

wooclap Questions 7 / 10 + Messages 🔍 100 % 🔍 Exit

pb:
saturation gradient killing

What happens here?

You cannot vote anymore

DATA: Which dataset do you want to use?
FEATURES: Which properties do you want to predict?
HIDDEN LAYERS: 4 HIDDEN LAYERS
OUTPUT: Test loss 0.02, Training loss 0.000

1 Logistic Regression fails (data is linearly separable)
2 I should use more neurons
3 The network overfits
4 I should add more layers
5 I should change the activation function
6 Logistic regression fails
7 One hidden layer is sufficient
8 I should change to a non-linear activation

wooclap Questions 8 / 10 + Messages 🔍 100 % 🔍 Exit

Test Loss 16%
Train = 5%

see slide on activation func

What happens here?

You cannot vote anymore

DATA: Which dataset do you want to use?
FEATURES: Which properties do you want to predict?
HIDDEN LAYERS: 8 HIDDEN LAYERS
OUTPUT: Test loss 0.000, Training loss 0.000

1 Logistic Regression works (data is linearly separable)
2 I should use more neurons
3 The network overfits
4 I should add more layers
5 I should change the activation function
6 Logistic regression fails
7 It works!
8 I should change to a non-linear activation

wooclap Questions 10 / 10 + Messages 🔍 100 % 🔍 Exit

Test: 5.8%
Tr: 3%

Intuition

Up to that point, we have seen :

- What a Neural Network was.
- How to implement it in Python.

Next Question: why and when does it work ?

- ➊ Approximation theorem
- ➋ Optimization

Intuition

Up to that point, we have seen :

- What a Neural Network was.
- How to implement it in Python.

Next Question: why and when does it work ?

- ➊ Approximation theorem
- ➋ Optimization

Approximation Theorem

Goal

$$\min_{f \in \mathcal{F}_{\text{MLP}}} \frac{1}{n} \sum_{i=1}^n [\ell(f(X_i), Y_i)].$$

As said before, a MLP can output non-linear functions... But how powerful is it ?

Approximation Theorem

Goal

$$\min_{f \in \mathcal{F}_{\text{MLP}}} \frac{1}{n} \sum_{i=1}^n [\ell(f(X_i), Y_i)].$$

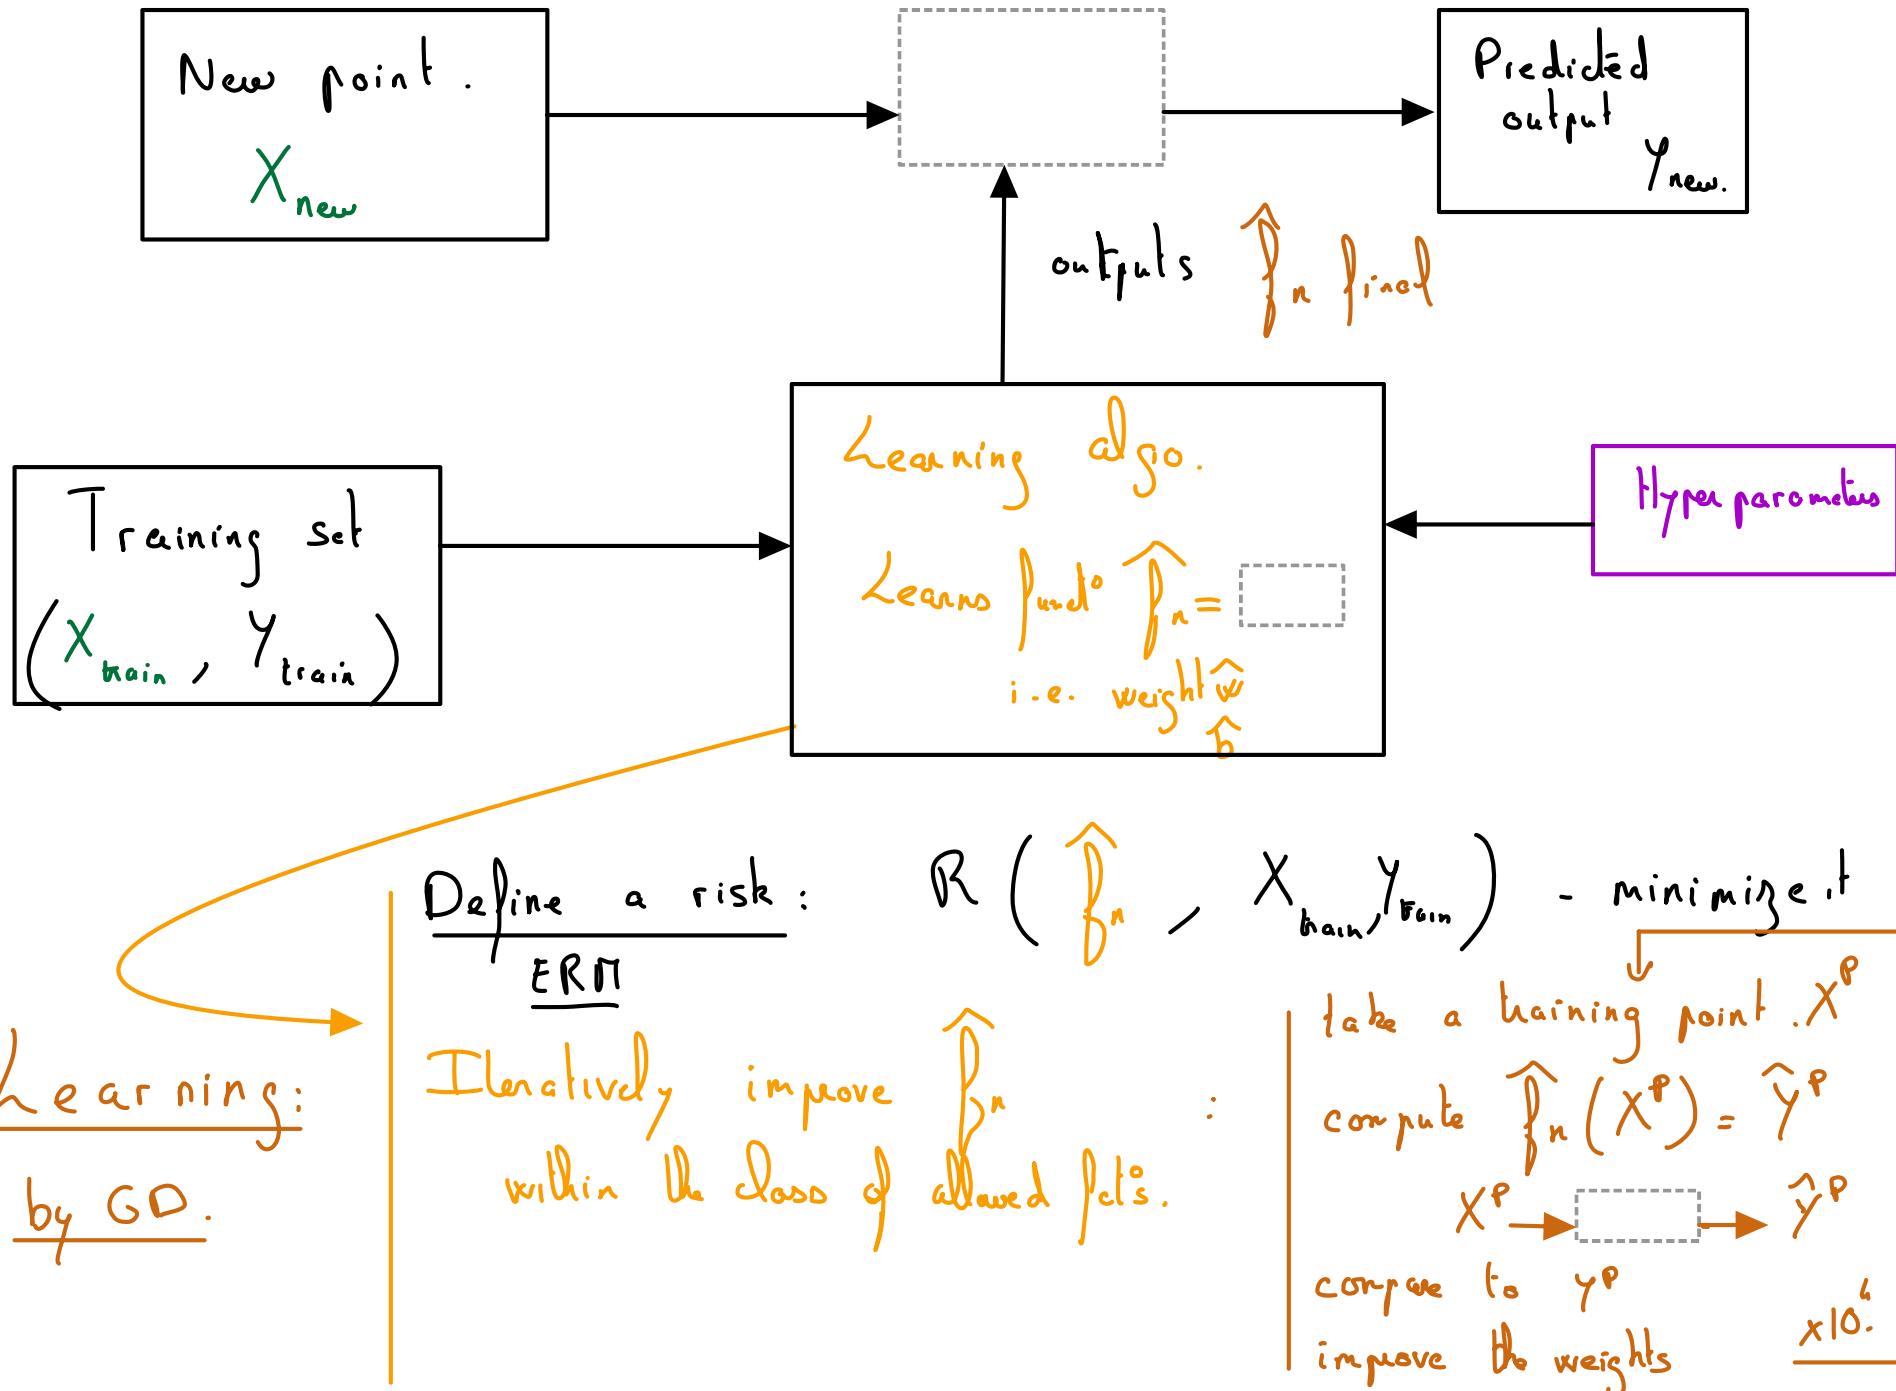
As said before, a MLP can output non-linear functions... But how powerful is it ?

Continuous Neural Networks with one single hidden layer and any bounded and non constant activation function can approximate any function in L^p , provided a sufficient number of hidden units.

in practice, use more than one hidden layer !!

→ Very powerful in terms of approximations.

Not very surprising: non linear function with millions of parameters !



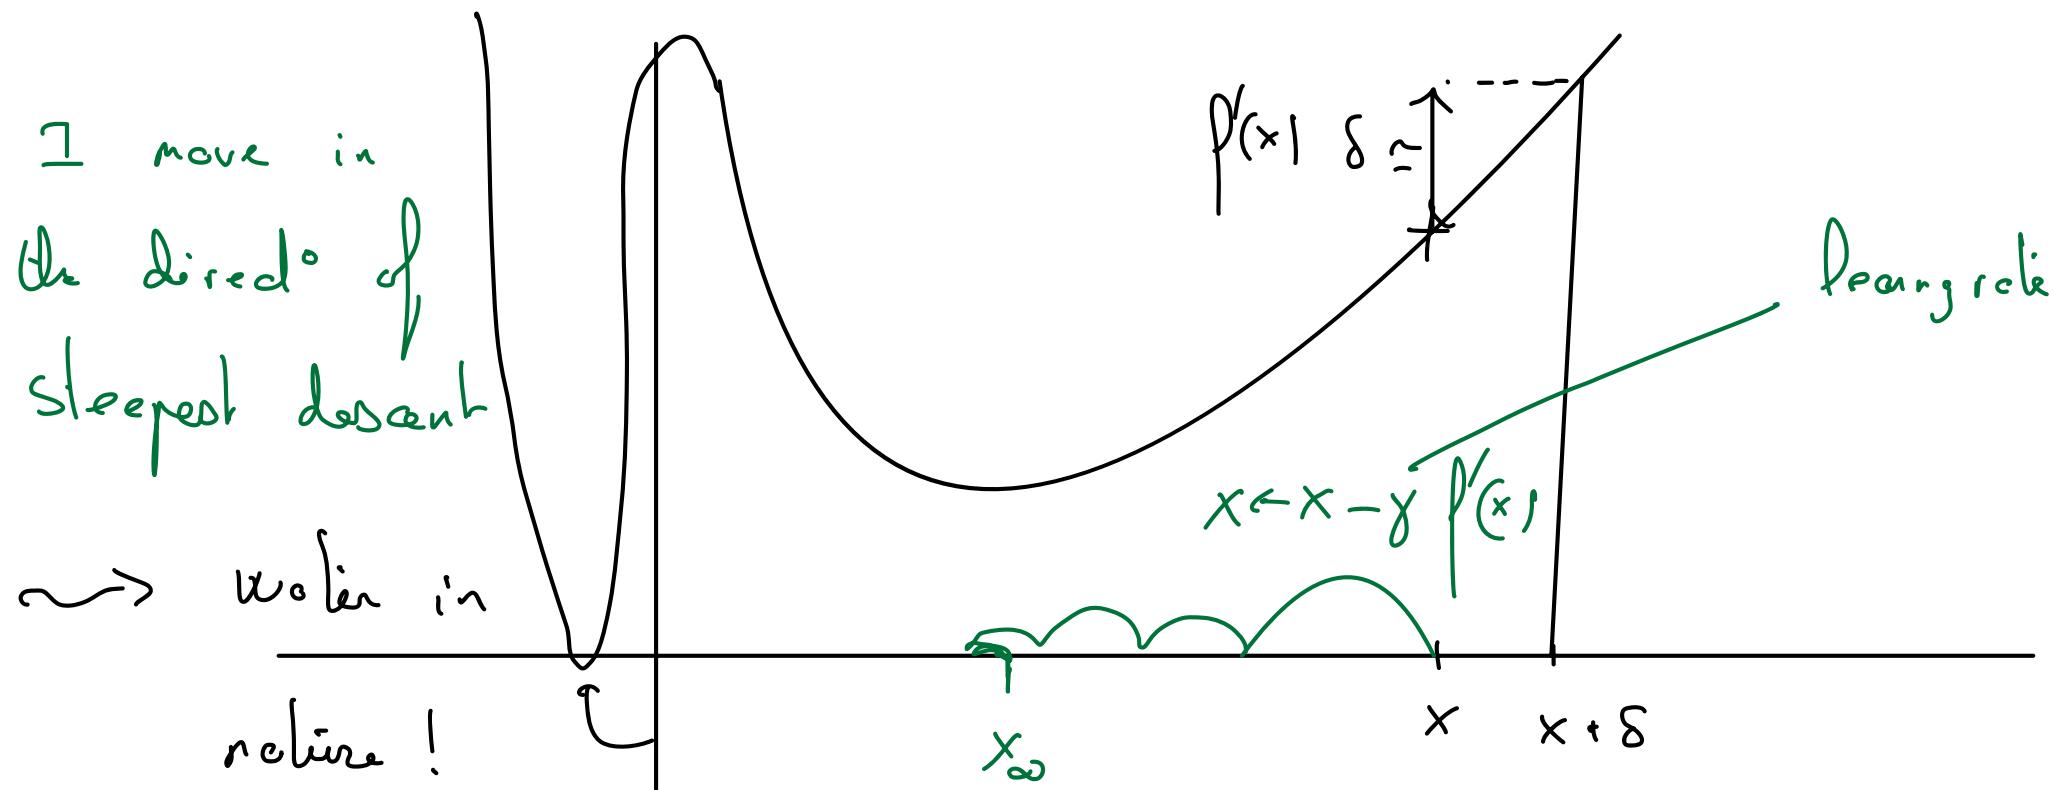
Optimization

Goal

$$\min_{f \in \mathcal{F}_{\text{MLP}}} \frac{1}{n} \sum_{i=1}^n [\ell(f(X_i), Y_i)].$$

How to minimize a function?

Gradient Methods. = cheapest way to update (learn) the weights iteratively to minimize the loss.



Optimization

Goal

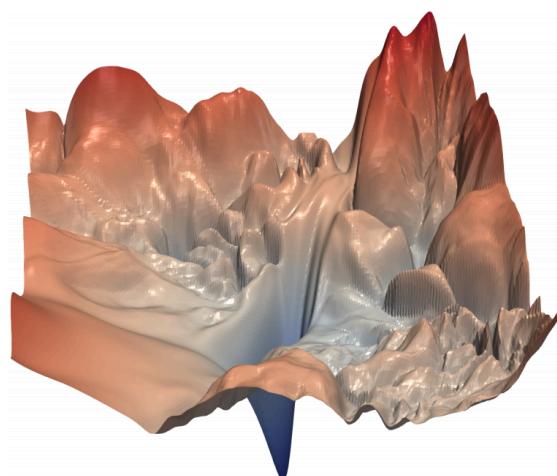
$$\min_{f \in \mathcal{F}_{\text{MLP}}} \frac{1}{n} \sum_{i=1}^n [\ell(f(X_i), Y_i)].$$

Gradient Methods. = cheapest way to update (learn) the weights iteratively to minimize the loss.

Two “miracles”:

- ① Autodifferentiation: even though the function is very complex and high dimensional, it is possible to compute gradients !
- ② Optimization seems to work ok, though the function is non convex - we do not end up in too bad local minima. (specialized optim algos: Adam, AdaDelta, etc.)

(local)



↑
non-convex on
gradient-descent

These 2 miracles make it possible to learn a good regressor/classifier with NN and to benefit from the powerful approximation properties

Conclusion

Neural networks :

- ① learn very complex non linear functions in high dimension
- ② can approximate nearly any function
- ③ can be optimized even though highly non convex.

Are thus expected to work:

- When depth or width increases.
- Thus with very large datasets
- Requiring a lot of computational power.

⇒ Explains why they were so successful since 2010

Conclusion

Neural networks :

- ① learn very complex non linear functions in high dimension
- ② can approximate nearly any function
- ③ can be optimized even though highly non convex.

Are thus expected to work:

- When depth or width increases.
- Thus with very large datasets
- Requiring a lot of computational power.

⇒ **Explains why they were so successful since 2010**

What my layer wants me to say: we haven't talked about many points !!

- Initialization
- Regularization

Conclusion

Neural networks :

- ① learn very complex non linear functions in high dimension
- ② can approximate nearly any function
- ③ can be optimized even though highly non convex.

Are thus expected to work:

- When depth or width increases.
- Thus with very large datasets
- Requiring a lot of computational power.

⇒ **Explains why they were so successful since 2010**

What my layer wants me to say: we haven't talked about many points !!

- Initialization
- Regularization

Next:

- Brief history
- How to use the structure ?

A brief history of NN

- 1943: Neural networks
- 1957: Perceptron
- 1974-86: Backpropagation, RBM, RNN
- 1989-98: CNN, MNIST
- 2009: ImageNet
- 2012: AlexNet,
- 2014: GANss
- 2016: AlphaGo
- 2018: BERT

The Computational power made the major change (+ investment and creativity).

Directions

- When does it work ?

Large or huge datasets. Structured tasks.

↳ **Images and text.**

⚠ do not try to apply DL everywhere

Each application requires a special architecture:

- Images : Convolutional Neural Network (CNN). 
- Text : Recurrent Neural Networks (RNN).

↳ homework

Outline

1 Goals

2 The first Neural Network: the Perceptron

3 Convolutional Neural Networks

4 Applications

- Image classification
- Pose, action detection
- Object detection and segmentation
- Scene labeling - Semantic segmentation
- Object tracking - videos

We consider an **image classification task**: we want to recognize which object is on an image.

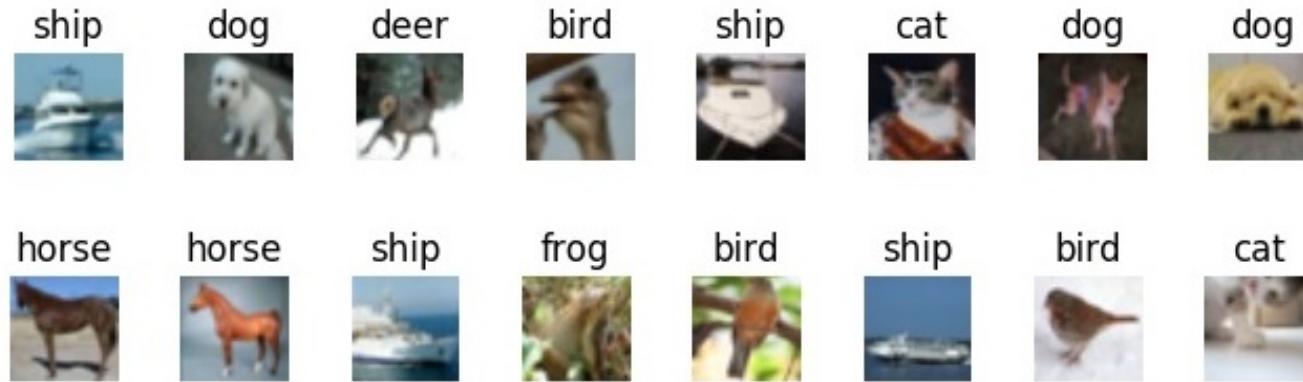


Figure: CIFAR dataset

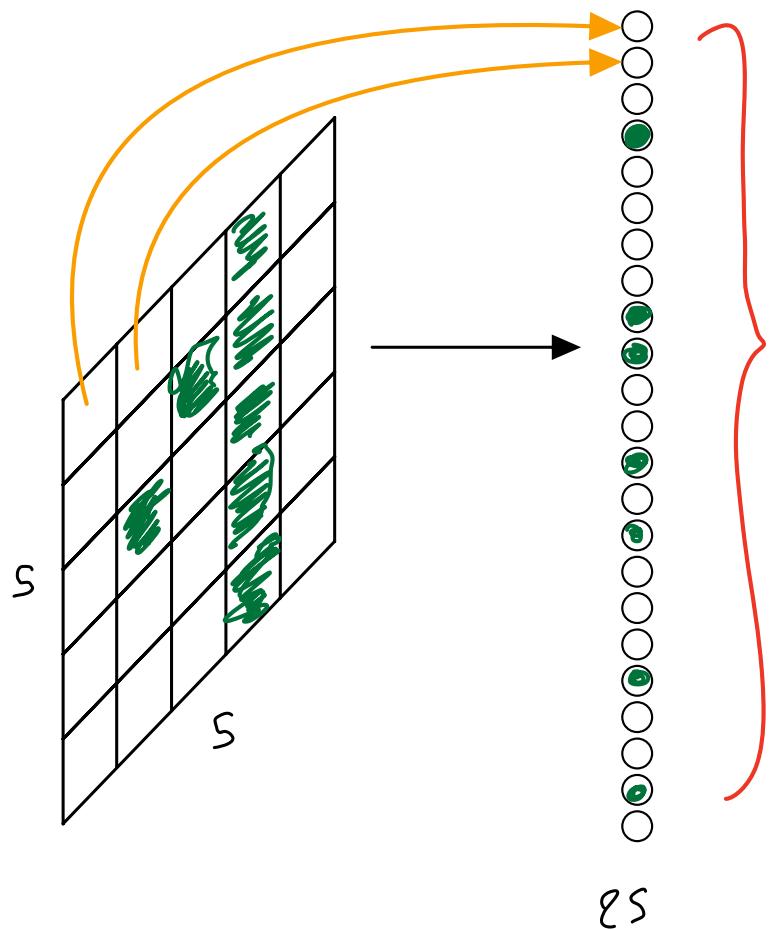
The input is an image: about 10^3 to 10^7 pixels: these are our inputs.

Problems: The multi-layer perceptron:

- Does not take into account local information. It would work similarly on any permutation of the pixels !
- Has, just on the first layer, as many parameters as there are inputs !

Convolutional Layer

Instead of making a linear combination of all the pixels, we consider a weighted average of a moving window of pixels, of size 3x3, or 5x5...

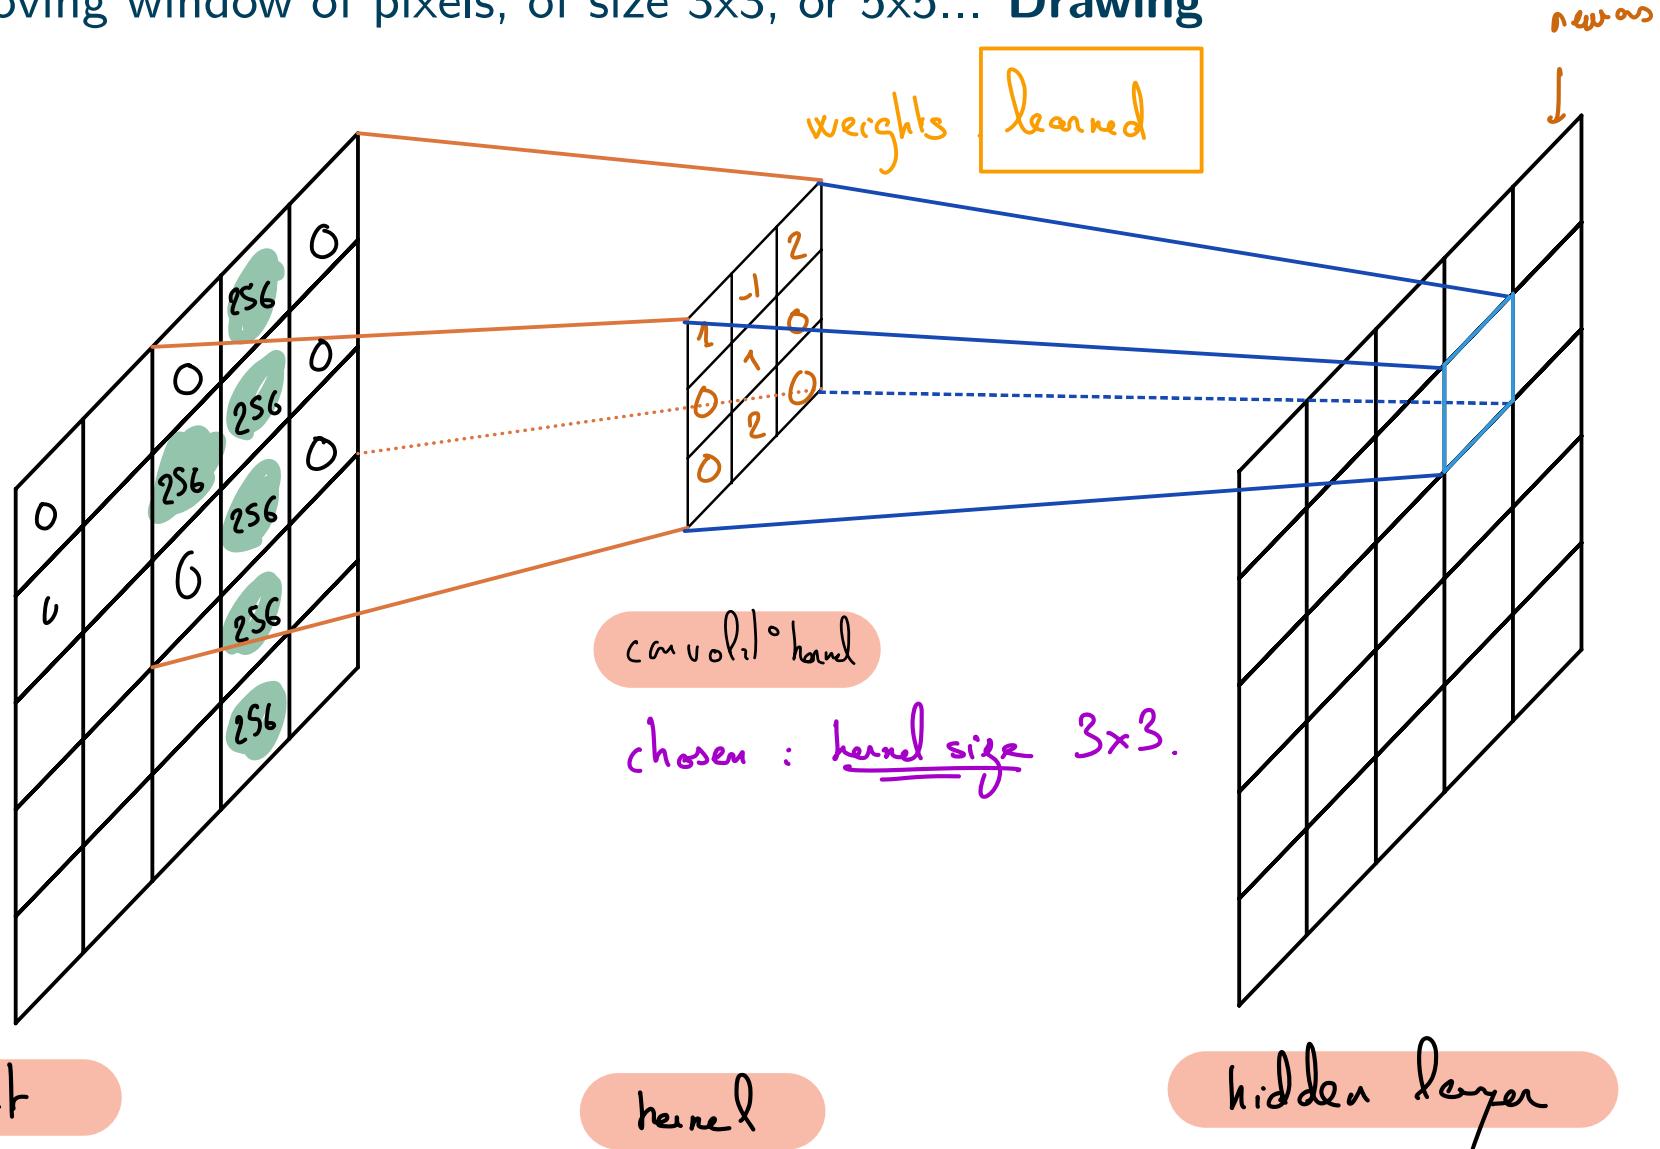


a fully connected layer considers
all the inputs symmetrically
no action of
neighboring pixels.

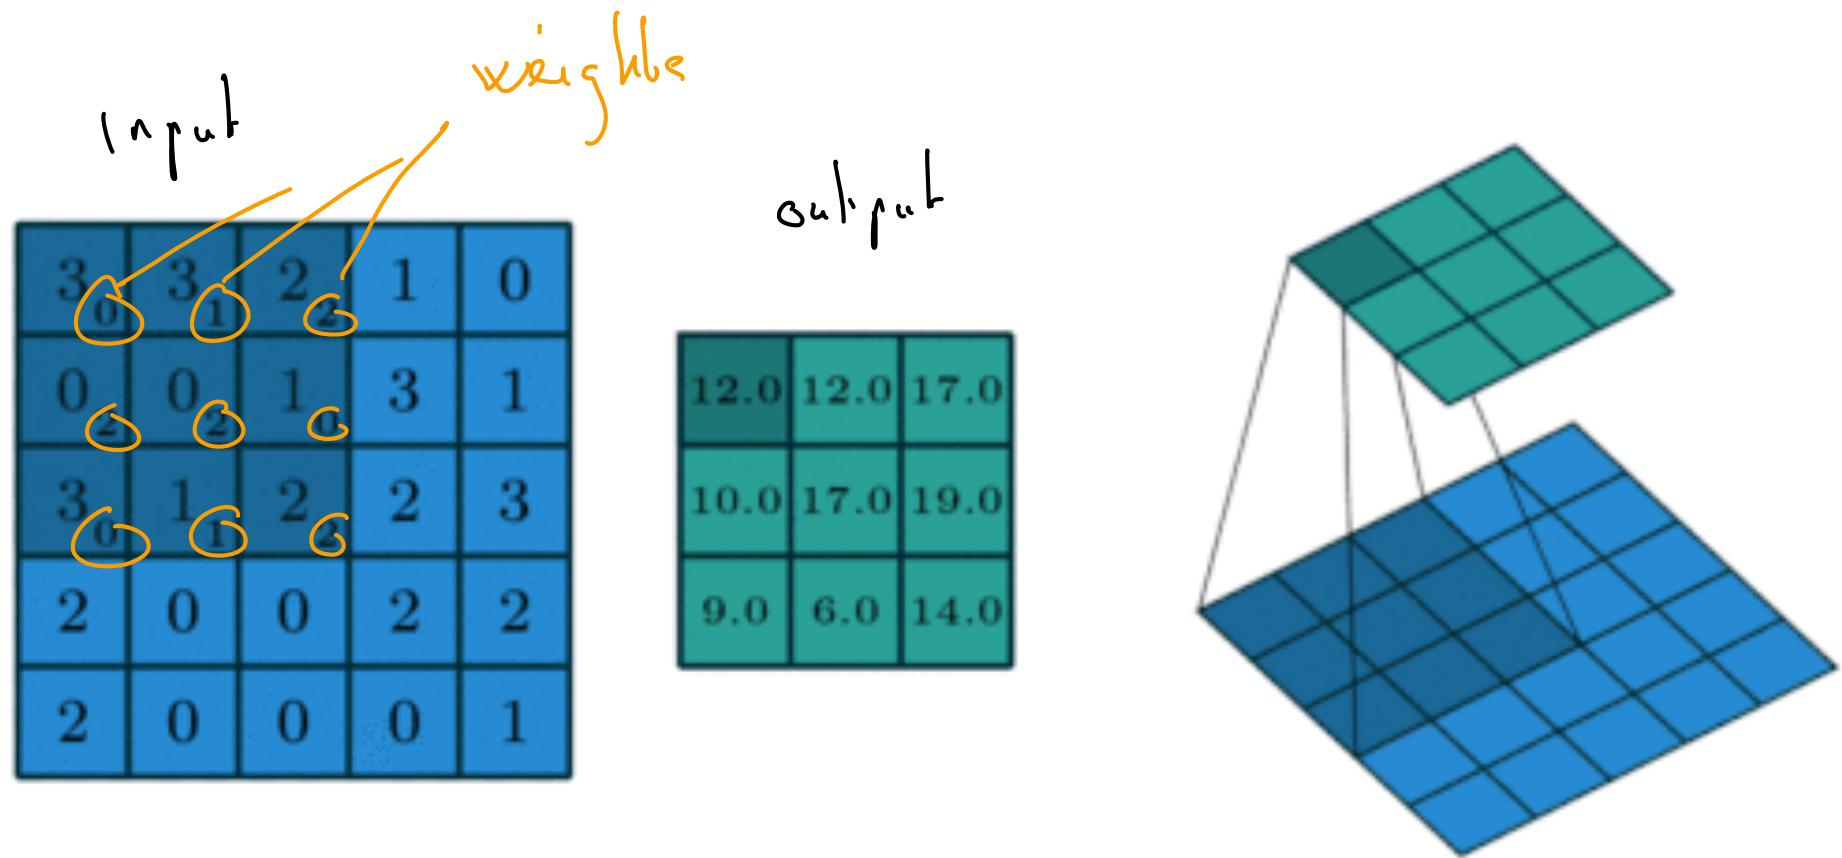
⚠ Take into account the neighbours!

Convolutional Layer

Instead of making a linear combination of all the pixels, we consider a weighted average of a moving window of pixels, of size 3×3 , or 5×5 ... Drawing



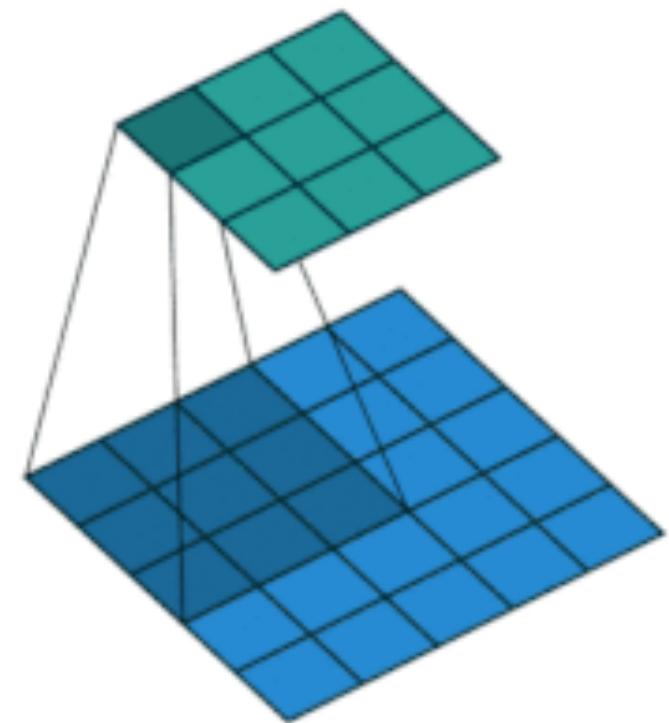
Convolutional Layer



Convolutional Layer

3 ₀	3 ₁	2 ₂	1	0
0 ₂	0 ₂	1 ₀	3	1
3 ₀	1 ₁	2 ₂	2	3
2	0	0	2	2
2	0	0	0	1

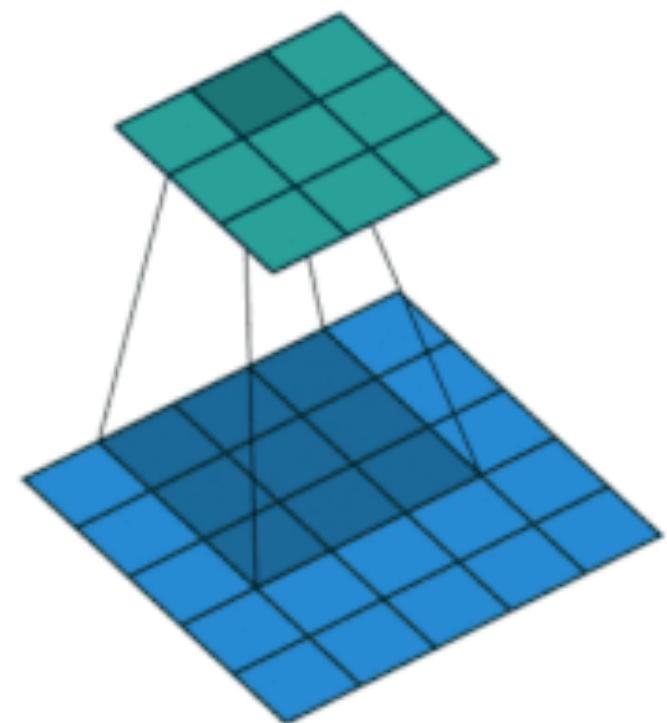
12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0



Convolutional Layer

3	3 ₀	2 ₁	1 ₂	0
0	0 ₂	1 ₂	3 ₀	1
3	1 ₀	2 ₁	2 ₂	3
2	0	0	2	2
2	0	0	0	1

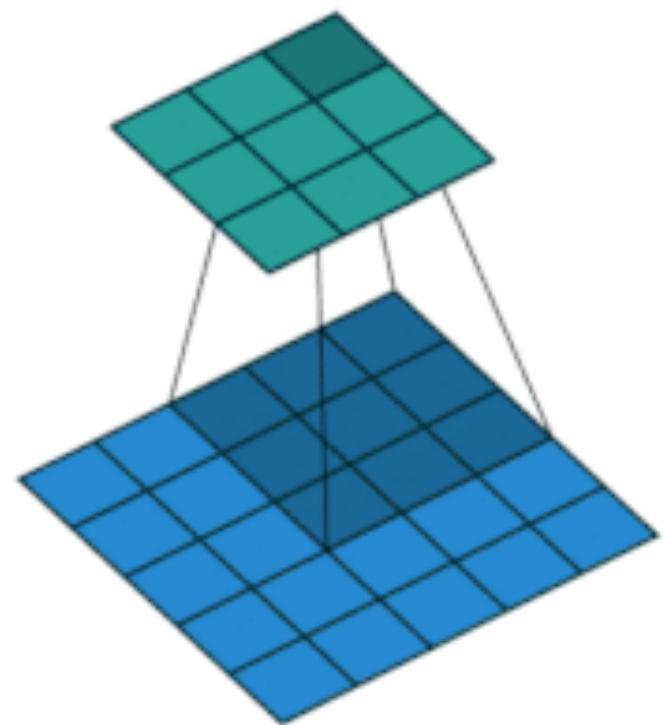
12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0



Convolutional Layer

3	3	2 ₀	1 ₁	0 ₂
0	0	1 ₂	3 ₂	1 ₀
3	1	2 ₀	2 ₁	3 ₂
2	0	0	2	2
2	0	0	0	1

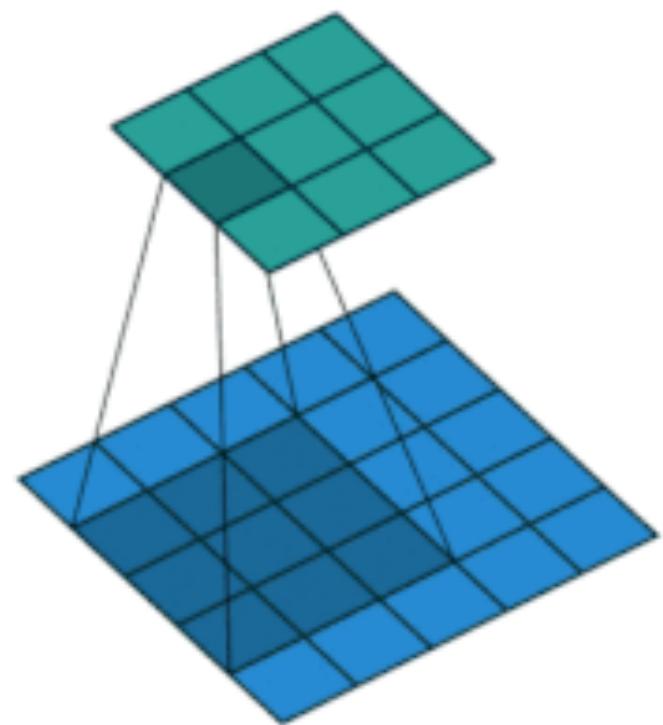
12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0



Convolutional Layer

3	3	2	1	0
0 ₀	0 ₁	1 ₂	3	1
3 ₂	1 ₂	2 ₀	2	3
2 ₀	0 ₁	0 ₂	2	2
2	0	0	0	1

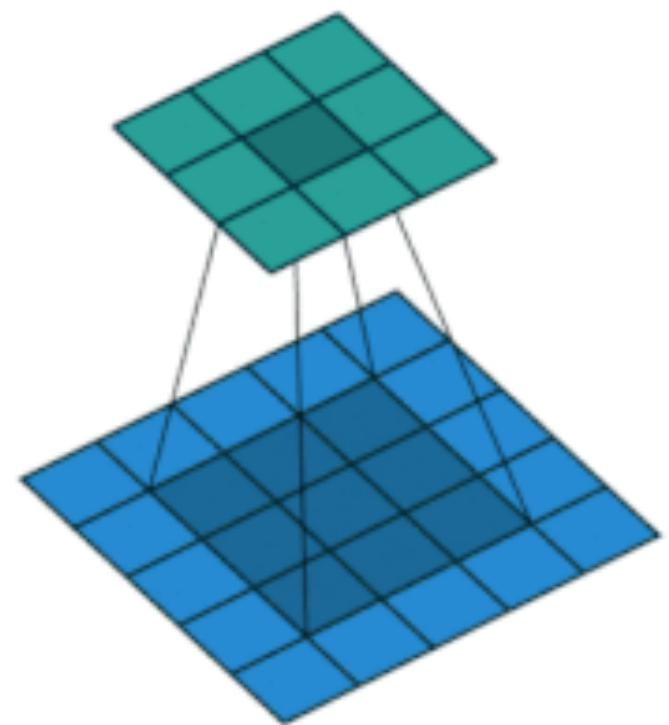
12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0



Convolutional Layer

3	3	2	1	0
0	0_0	1_1	3_2	1
3	1_2	2_2	2_0	3
2	0_0	0_1	2_2	2
2	0	0	0	1

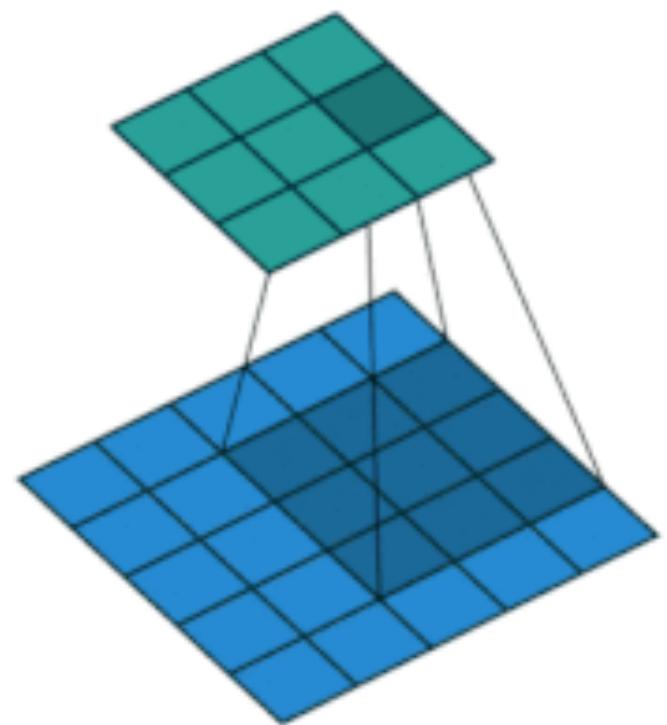
12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0



Convolutional Layer

3	3	2	1	0
0	0	1_0	3_1	1_2
3	1	2_2	2_2	3_0
2	0	0_0	2_1	2_2
2	0	0	0	1

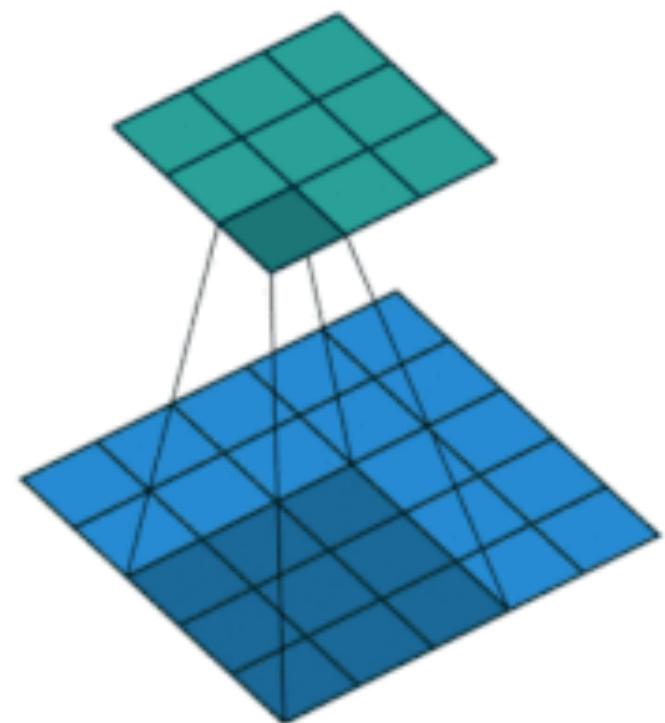
12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0



Convolutional Layer

3	3	2	1	0
0	0	1	3	1
3 ₀	1 ₁	2 ₂	2	3
2 ₂	0 ₂	0 ₀	2	2
2 ₀	0 ₁	0 ₂	0	1

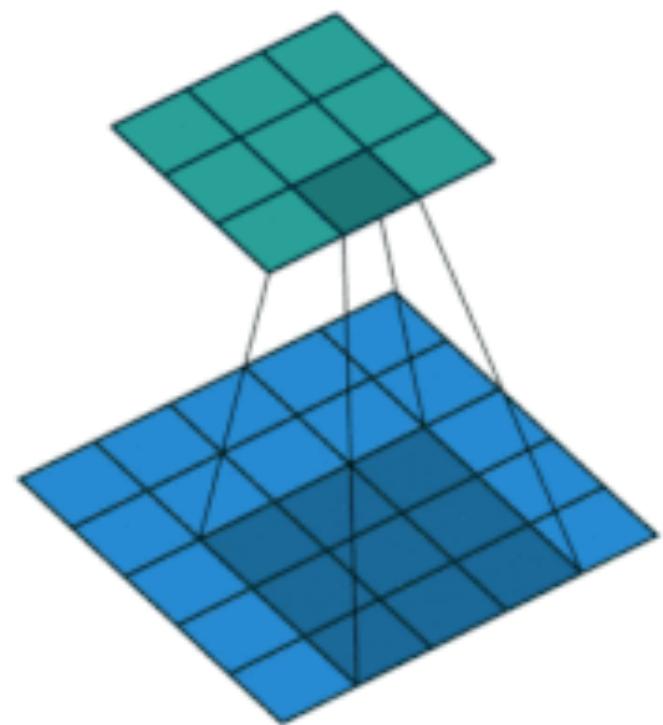
12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0



Convolutional Layer

3	3	2	1	0
0	0	1	3	1
3	1 ₀	2 ₁	2 ₂	3
2	0 ₂	0 ₂	2 ₀	2
2	0 ₀	0 ₁	0 ₂	1

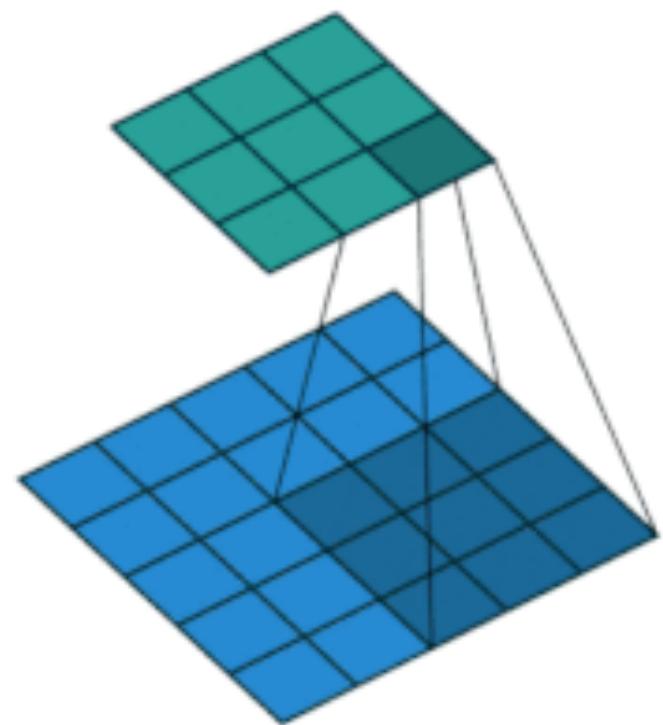
12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0



Convolutional Layer

3	3	2	1	0
0	0	1	3	1
3	1	2 ₀	2 ₁	3 ₂
2	0	0 ₂	2 ₂	2 ₀
2	0	0 ₀	0 ₁	1 ₂

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0



Convolutional Layer

0	0	0	0	0	0	0	...
0	156	155	156	158	158	...	
0	153	154	157	159	159	...	
0	149	151	155	158	159	...	
0	146	146	149	153	158	...	
0	145	143	143	148	158	...	
...

Input Channel #1 (Red)

0	0	0	0	0	0	0	...
0	167	166	167	169	169	...	
0	164	165	168	170	170	...	
0	160	162	166	169	170	...	
0	156	156	159	163	168	...	
0	155	153	153	158	168	...	
...

Input Channel #2 (Green)

0	0	0	0	0	0	0	...
0	163	162	163	165	165	...	
0	160	161	164	166	166	...	
0	156	158	162	165	166	...	
0	155	155	158	162	167	...	
0	154	152	152	157	167	...	
...

Input Channel #3 (Blue)

-1	-1	1
0	1	-1
0	1	1

Kernel Channel #1

1	0	0
1	-1	-1
1	0	-1

Kernel Channel #2

0	1	1
0	1	0
1	-1	1

Kernel Channel #3

308

+

-498

164

+ 1 = -25

Bias = 1

-25					...
					...
					...
					...
...



Convolutional Layer

padding : ensure output image has the same size

0	0	0	0	0	0	...
0	156	155	156	158	158	...
0	153	154	157	159	159	...
0	149	151	155	158	159	...
0	146	146	149	153	158	...
0	145	143	143	148	158	...
...

Input Channel #1 (Red)

0	0	0	0	0	0	...
0	167	166	167	169	169	...
0	164	165	168	170	170	...
0	160	162	166	169	170	...
0	156	156	159	163	168	...
0	155	153	153	158	168	...
...

Input Channel #2 (Green)

0	0	0	0	0	0	...
0	163	162	163	165	165	...
0	160	161	164	166	166	...
0	156	158	162	165	166	...
0	155	155	158	162	167	...
0	154	152	152	157	167	...
...

Input Channel #3 (Blue)

-1	-1	1
0	1	-1
0	1	1

Kernel Channel #1



308

+

1	0	0
1	-1	-1
1	0	-1

Kernel Channel #2



-498

0	1	1
0	1	0
1	-1	1

Kernel Channel #3



164 + 1 = -25

Bias = 1

Output

-25					...
					...
					...
					...
...

Convolutional Layer

0	0	0	0	0	0	...
0	156	155	156	158	158	...
0	153	154	157	159	159	...
0	149	151	155	158	159	...
0	146	146	149	153	158	...
0	145	143	143	148	158	...
...

Input Channel #1 (Red)

0	0	0	0	0	0	...
0	167	166	167	169	169	...
0	164	165	168	170	170	...
0	160	162	166	169	170	...
0	156	156	159	163	168	...
0	155	153	153	158	168	...
...

Input Channel #2 (Green)

0	0	0	0	0	0	...
0	163	162	163	165	165	...
0	160	161	164	166	166	...
0	156	158	162	165	166	...
0	155	155	158	162	167	...
0	154	152	152	157	167	...
...

Input Channel #3 (Blue)

-1	-1	1
0	1	-1
0	1	1

Kernel Channel #1

1	0	0
1	-1	-1
1	0	-1

Kernel Channel #2

0	1	1
0	1	0
1	-1	1

Kernel Channel #3

↓

310

+

↓

-170

+

325

+ 1 = 466

0,17 2 3,1
-17

weights learned,

Output

-25	466			...
				...
				...
				...
...

↑
Bias = 1

Convolutional Layer

0	0	0	0	0	0	0	...
0	156	155	156	158	158	158	...
0	153	154	157	159	159	159	...
0	149	151	155	158	159	159	...
0	146	146	149	153	158	158	...
0	145	143	143	148	158	158	...
...

Input Channel #1 (Red)

0	0	0	0	0	0	0	...
0	167	166	167	169	169	169	...
0	164	165	168	170	170	170	...
0	160	162	166	169	170	170	...
0	156	156	159	163	168	168	...
0	155	153	153	158	168	168	...
...

Input Channel #2 (Green)

0	0	0	0	0	0	0	...
0	163	162	163	165	165	165	...
0	160	161	164	166	166	166	...
0	156	158	162	165	166	166	...
0	155	155	158	162	167	167	...
0	154	152	152	157	167	167	...
...

Input Channel #3 (Blue)

-1	-1	1
0	1	-1
0	1	1

Kernel Channel #1



314

+

1	0	0
1	-1	-1
1	0	-1

Kernel Channel #2



-175

0	1	1
0	1	0
1	-1	1

Kernel Channel #3



+

326

+ 1 = 466

Bias = 1

Output

-25	466	466	...
...
...
...
...

Convolutional Layer

0	0	0	0	0	0	0	...
0	156	155	156	158	158	158	...
0	153	154	157	159	159	159	...
0	149	151	155	158	159	159	...
0	146	146	149	153	158	158	...
0	145	143	143	148	158	158	...
...

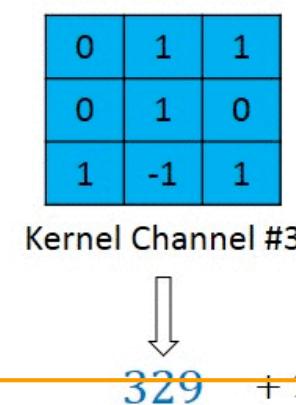
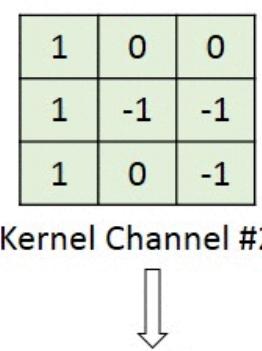
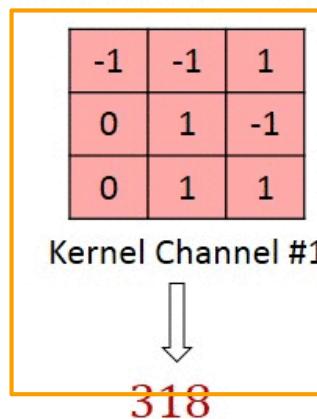
Input Channel #1 (Red)

0	0	0	0	0	0	0	...
0	167	166	167	169	169	169	...
0	164	165	168	170	170	170	...
0	160	162	166	169	170	170	...
0	156	156	159	163	168	168	...
0	155	153	153	158	168	168	...
...

Input Channel #2 (Green)

0	0	0	0	0	0	0	...
0	163	162	163	165	165	165	...
0	160	161	164	166	166	166	...
0	156	158	162	165	166	166	...
0	155	155	158	162	167	167	...
0	154	152	152	157	167	167	...
...

Input Channel #3 (Blue)



Learne d

Output

-25	466	466	475	...
...
...
...

Bias = 1

Convolutional Layer

0	0	0	0	0	0	...
0	156	155	156	158	158	...
0	153	154	157	159	159	...
0	149	151	155	158	159	...
0	146	146	149	153	158	...
0	145	143	143	148	158	...
...

Input Channel #1 (Red)

0	0	0	0	0	0	...
0	167	166	167	169	169	...
0	164	165	168	170	170	...
0	160	162	166	169	170	...
0	156	156	159	163	168	...
0	155	153	153	158	168	...
...

Input Channel #2 (Green)

0	0	0	0	0	0	...
0	163	162	163	165	165	...
0	160	161	164	166	166	...
0	156	158	162	165	166	...
0	155	155	158	162	167	...
0	154	152	152	157	167	...
...

Input Channel #3 (Blue)

-1	-1	1
0	1	-1
0	1	1

Kernel Channel #1



298

+

1	0	0
1	-1	-1
1	0	-1

Kernel Channel #2



-491

0	1	1
0	1	0
1	-1	1

Kernel Channel #3



487

+

+ 1 = 295

Bias = 1
↑

-25	466	466	475	...
295				...
				...
				...
...

Output

Convolutional Layer

0	0	0	0	0	0	...
0	156	155	156	158	158	...
0	153	154	157	159	159	...
0	149	151	155	158	159	...
0	146	146	149	153	158	...
0	145	143	143	148	158	...
...

Input Channel #1 (Red)

0	0	0	0	0	0	...
0	167	166	167	169	169	...
0	164	165	168	170	170	...
0	160	162	166	169	170	...
0	156	156	159	163	168	...
0	155	153	153	158	168	...
...

Input Channel #2 (Green)

0	0	0	0	0	0	...
0	163	162	163	165	165	...
0	160	161	164	166	166	...
0	156	158	162	165	166	...
0	155	155	158	162	167	...
0	154	152	152	157	167	...
...

Input Channel #3 (Blue)

-1	-1	1
0	1	-1
0	1	1

Kernel Channel #1



148

+

1	0	0
1	-1	-1
1	0	-1

Kernel Channel #2



-8

0	1	1
0	1	0
1	-1	1

Kernel Channel #3



646

+

+ 1 = 787

Bias = 1
↑

-25	466	466	475	...
295	787			...
				...
				...
...

Summary

CNN:

- allow to take spacial information into account
- reduce the number of parameters per layer. We can have deeper networks!

Summary

CNN:

- allow to take spacial information into account
- reduce the number of parameters per layer. We can have deeper networks!

LeNet 1998:

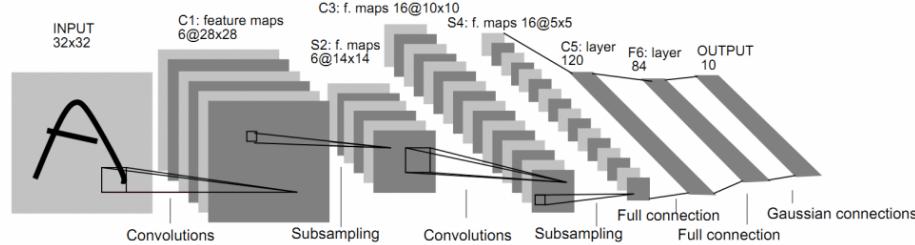
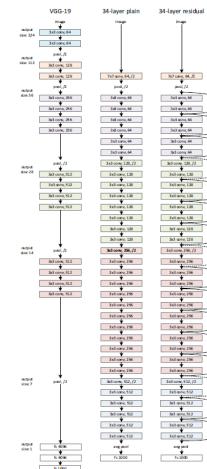


Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

Resnet 2016



8 layers, 60k parameters.

~ 40 Layers, 140M parameters

Outline

1 Goals

2 The first Neural Network: the Perceptron

3 Convolutional Neural Networks

4 Applications

- Image classification
- Pose, action detection
- Object detection and segmentation
- Scene labeling - Semantic segmentation
- Object tracking - videos

Outline

1 Goals

2 The first Neural Network: the Perceptron

3 Convolutional Neural Networks

4 Applications

- **Image classification**
- Pose, action detection
- Object detection and segmentation
- Scene labeling - Semantic segmentation
- Object tracking - videos

Image classification - CNN Tree

Category	Confusion Set					
tench	gar	sturgeon	coho	eel	barracouta	
indigo bunting	European gallinule	jacamar	peacock	coucal	macaw	jay
red-breasted merganser	albatross	pelican	oystercatcher	drake	redshank	goose
echidna	porcupine	beaver	armadillo	mongoose		American coot
shopping basket	bucket	shopping cart	packet	mailbag	hamper	grocery store

Figure: Confusion set outputs by AlexNet softmax prediction on validation set of ILSVRC 2015.

Image classification - CNN Tree

Category	Example Validation Images					
barracouta						
church						
spaghetti squash						
espresso						
trolleybus						

Figure: Top label is given by basic AlexNet CNN while bottom one is given by CNNTree (green color corresponds to a correct prediction)

Outline

1 Goals

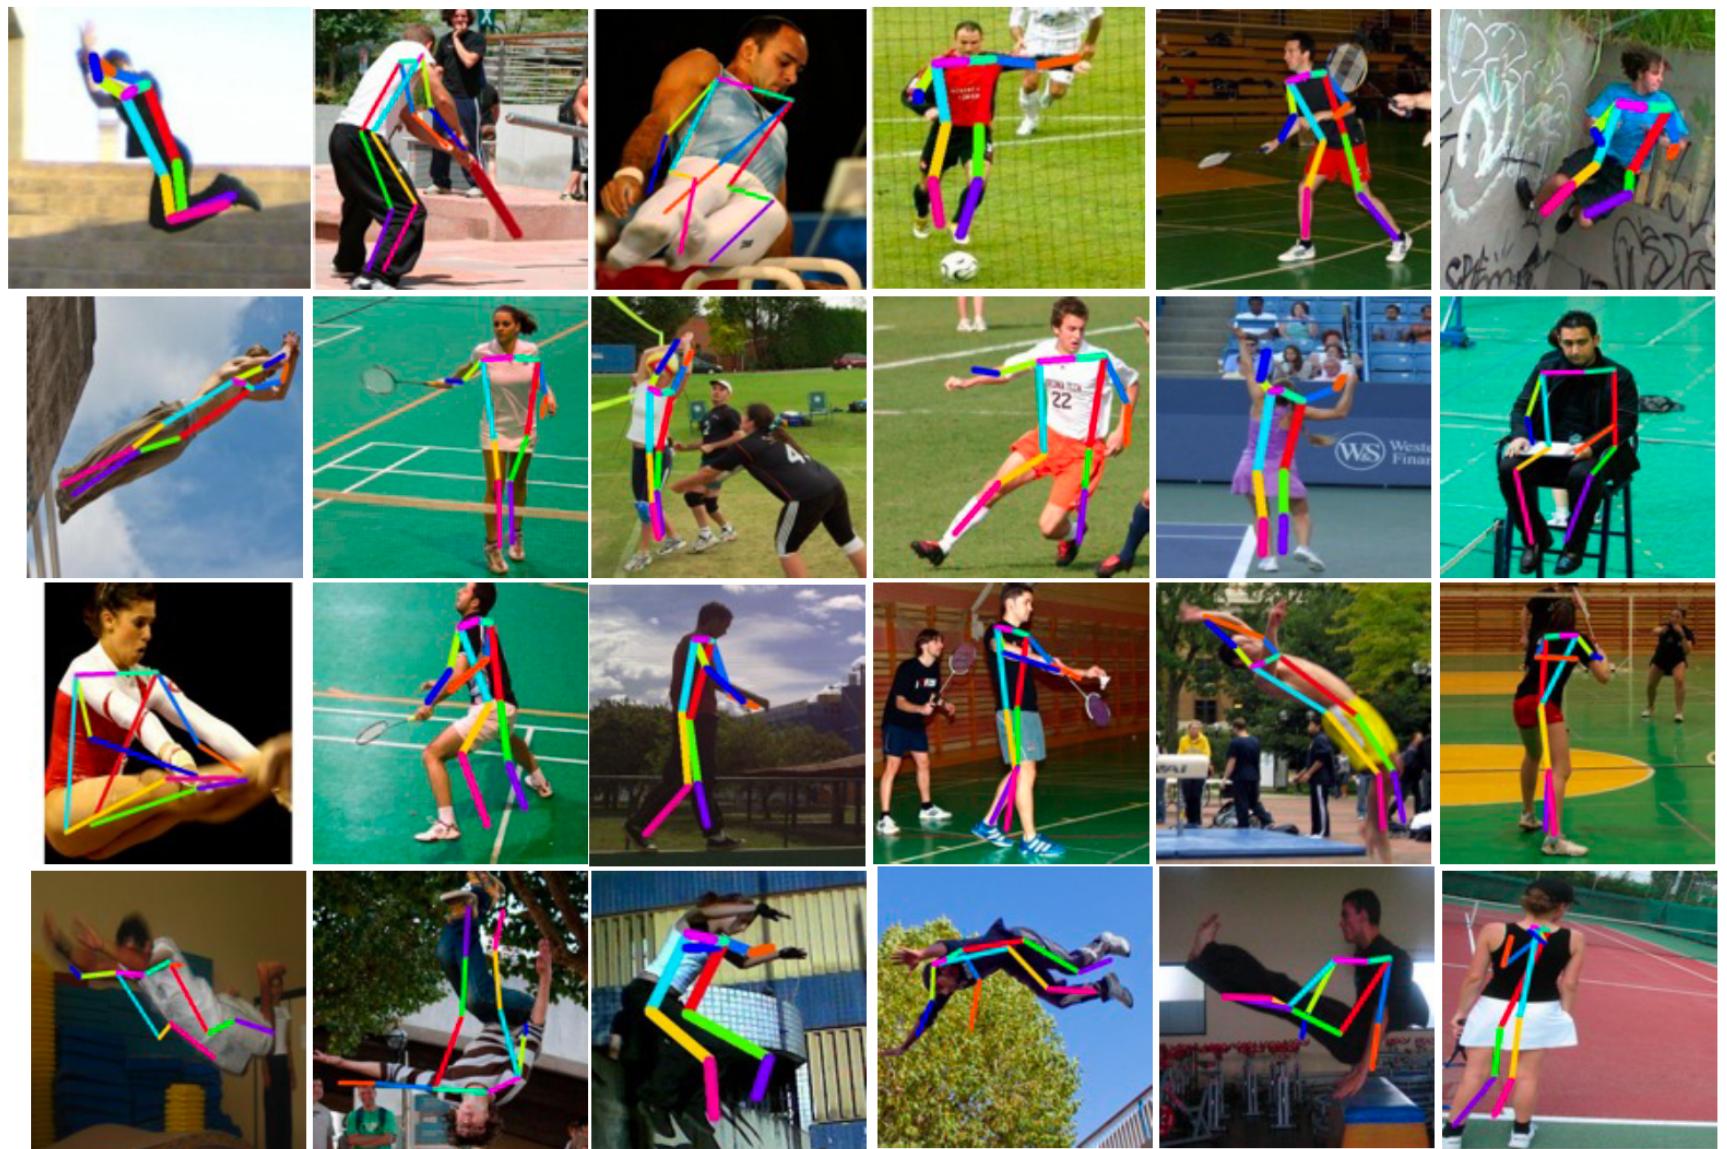
2 The first Neural Network: the Perceptron

3 Convolutional Neural Networks

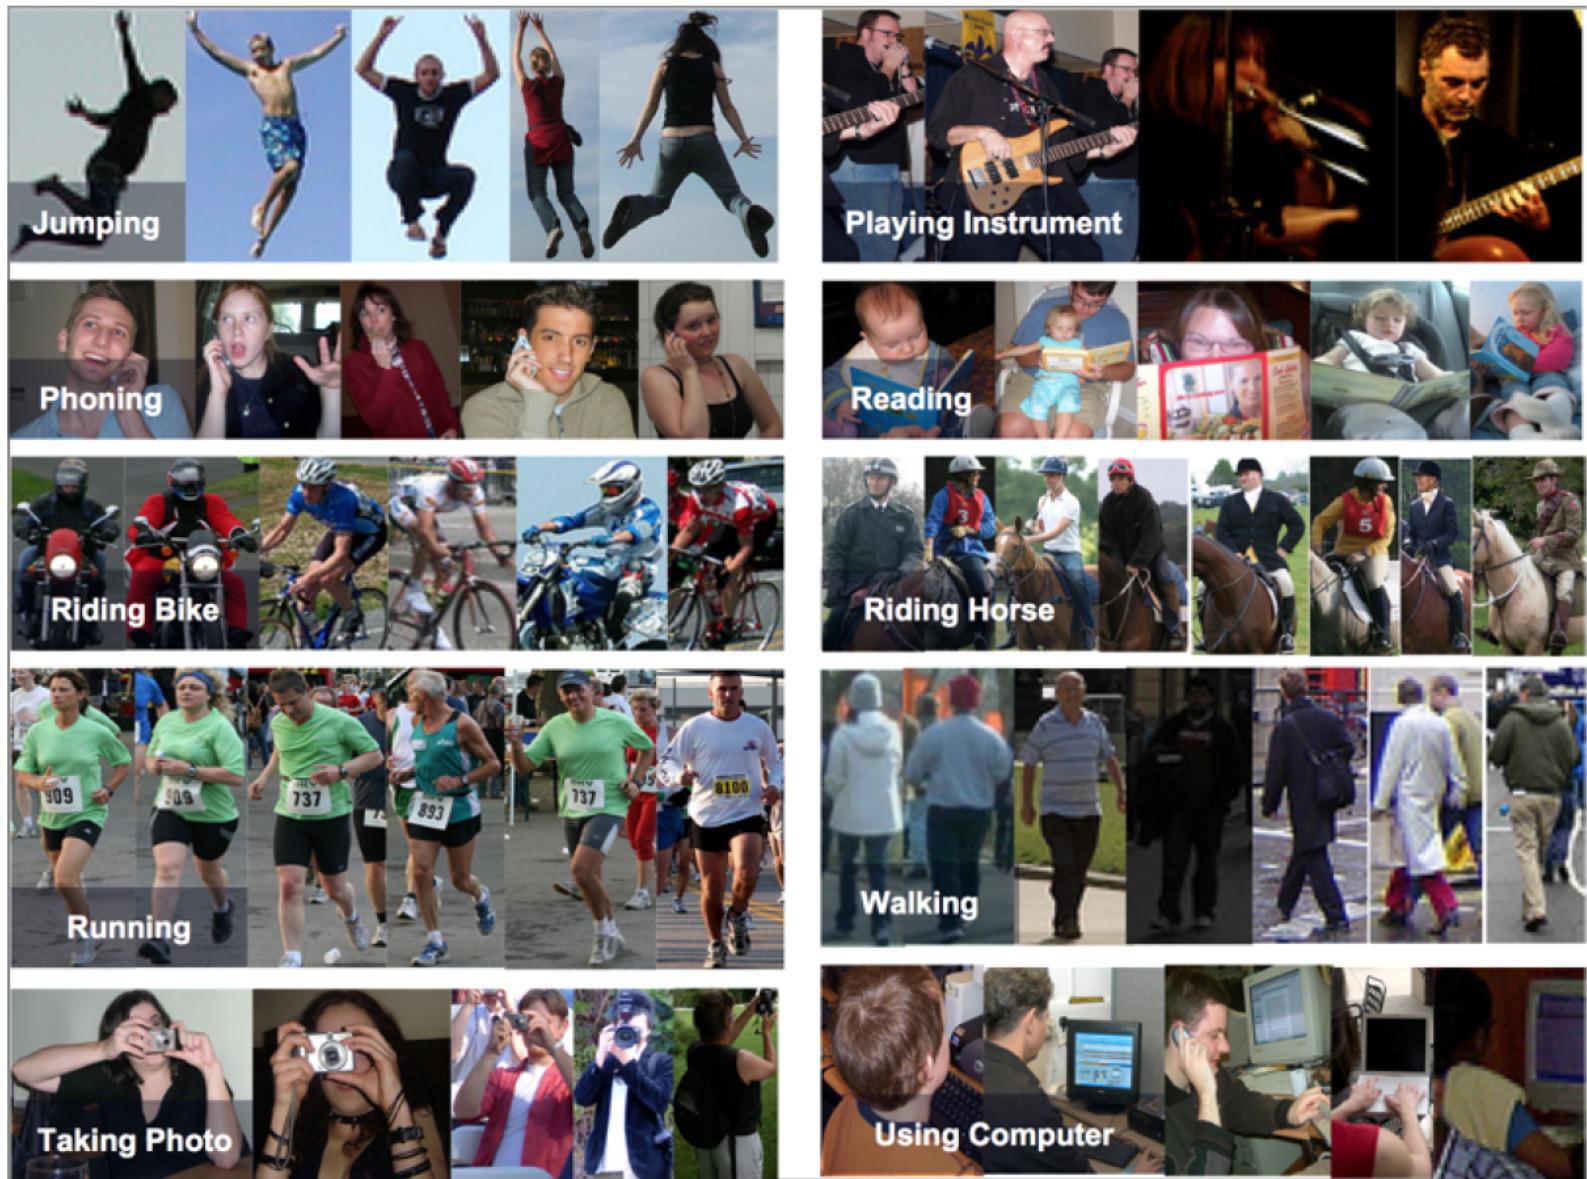
4 Applications

- Image classification
- Pose, action detection
- Object detection and segmentation
- Scene labeling - Semantic segmentation
- Object tracking - videos

Pose estimation - Deeppose



Action recognition



Outline

1 Goals

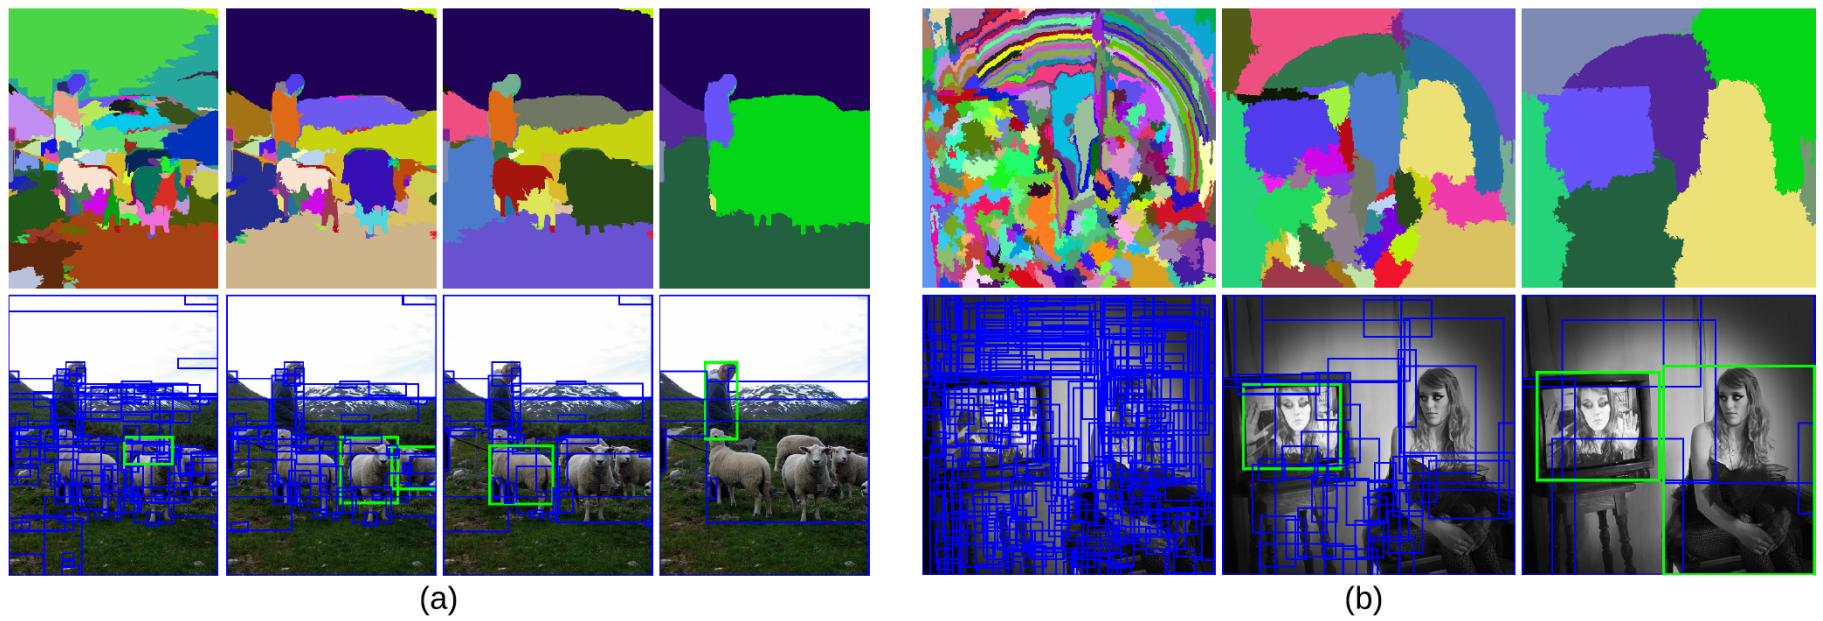
2 The first Neural Network: the Perceptron

3 Convolutional Neural Networks

4 Applications

- Image classification
- Pose, action detection
- **Object detection and segmentation**
- Scene labeling - Semantic segmentation
- Object tracking - videos

Object detection and segmentation



Object detection - YOLO, SDD

More recently, YOLO (You Only Look Once) and SSD (Single Shot Detector) allow single pipeline detection that directly predicts class labels.

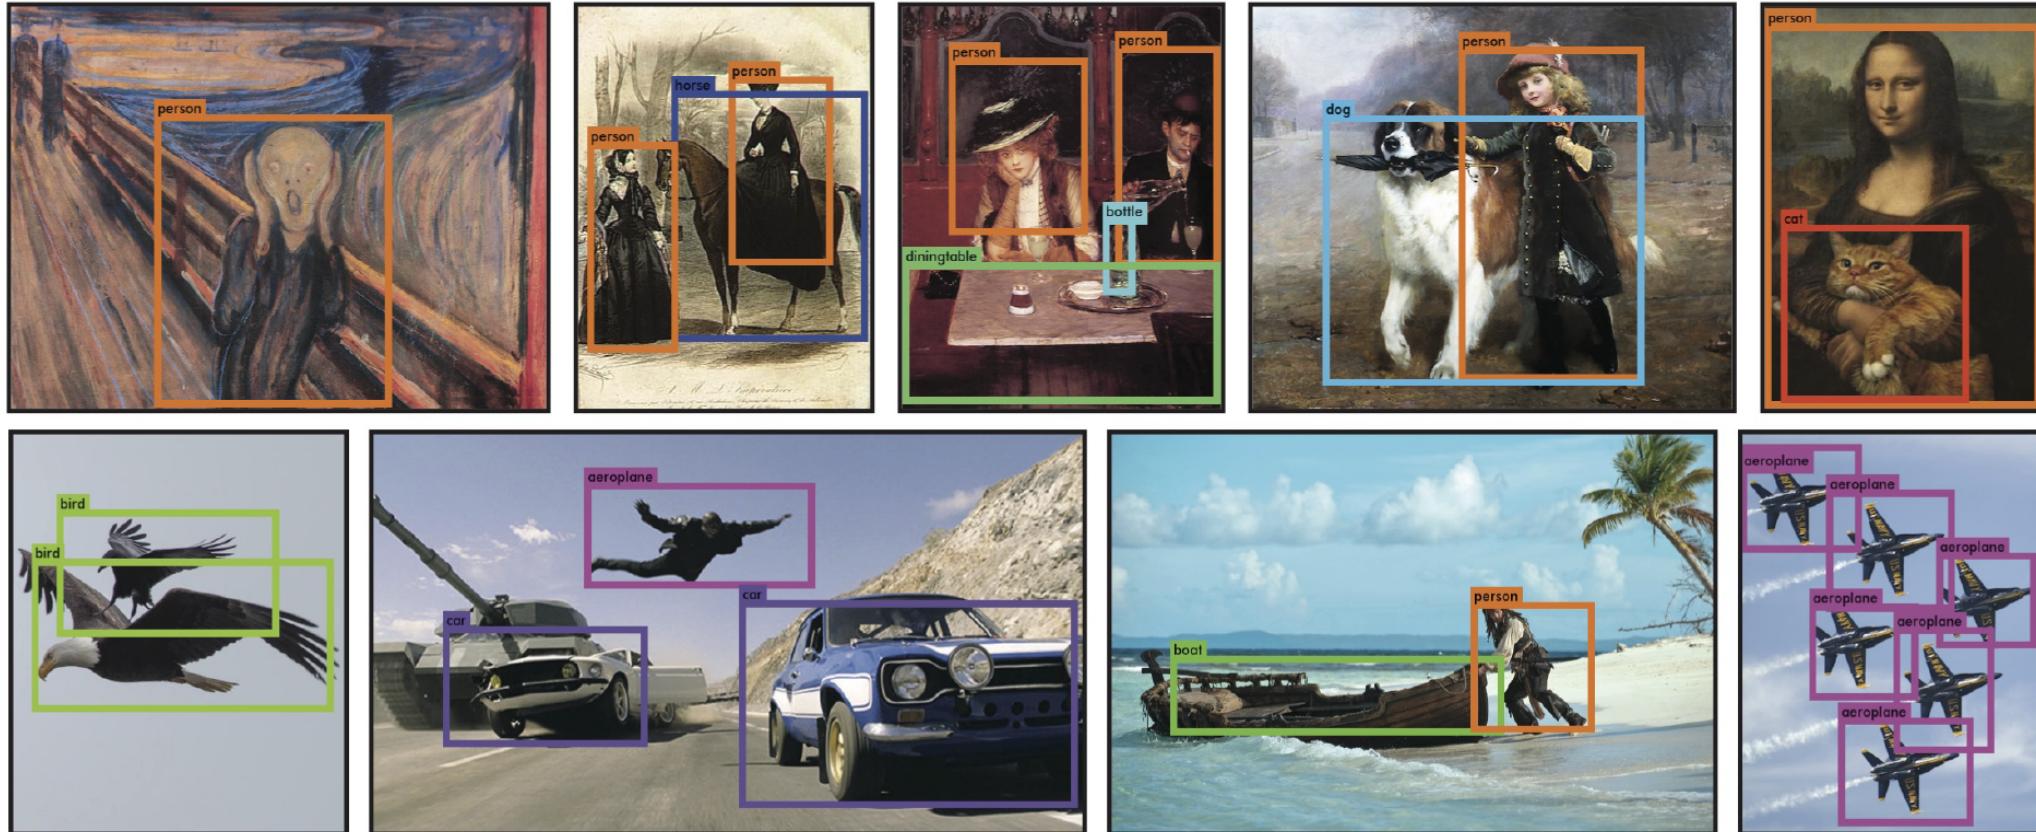
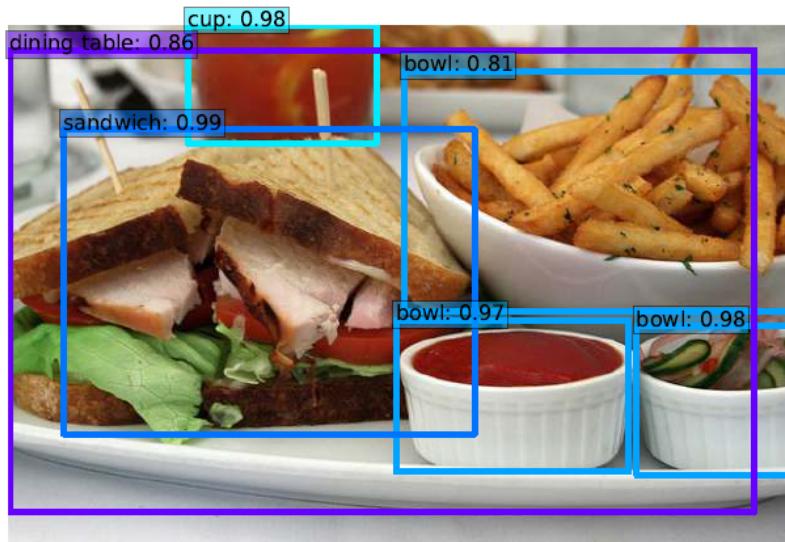
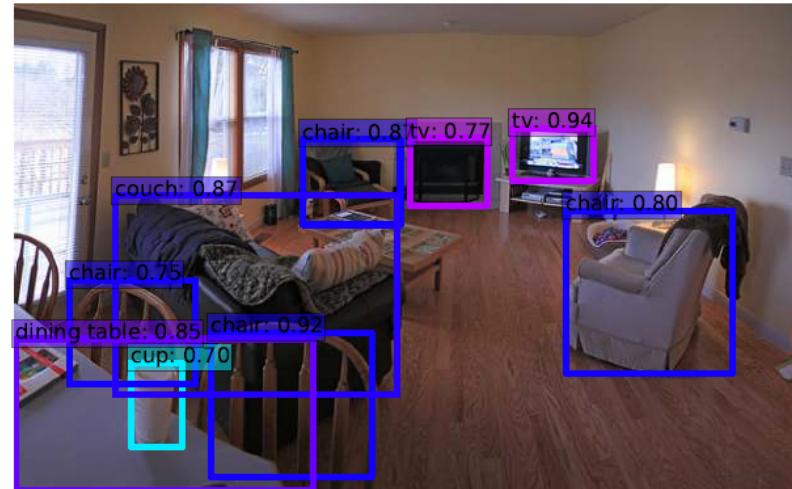
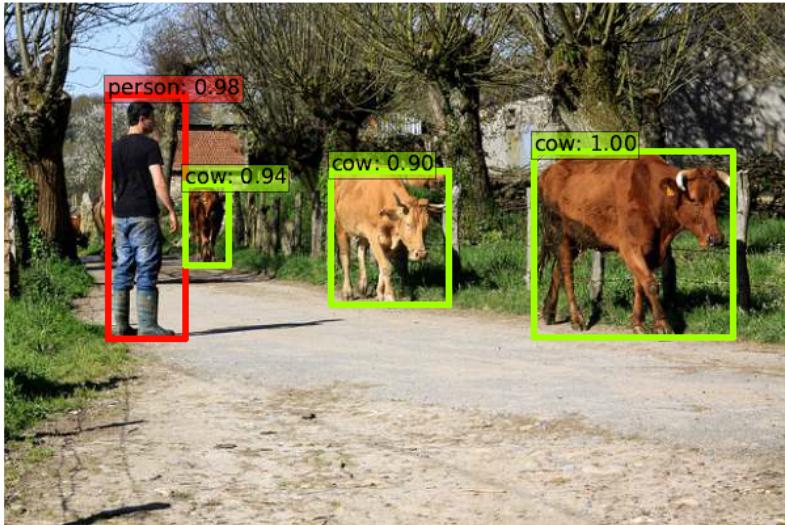


Figure: YOLO results



Outline

1 Goals

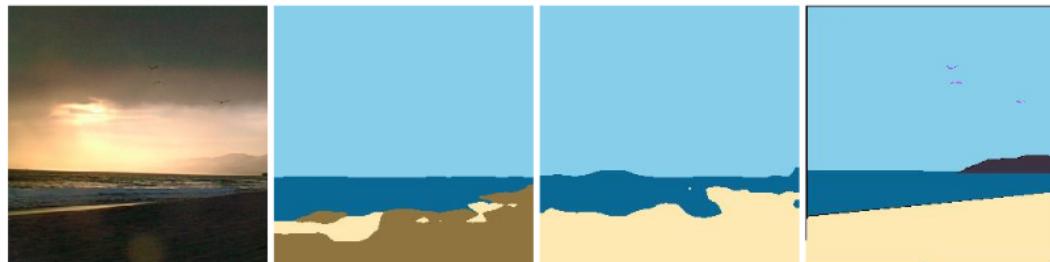
2 The first Neural Network: the Perceptron

3 Convolutional Neural Networks

4 Applications

- Image classification
- Pose, action detection
- Object detection and segmentation
- **Scene labeling - Semantic segmentation**
- Object tracking - videos

Scene labeling - DAG-RNN

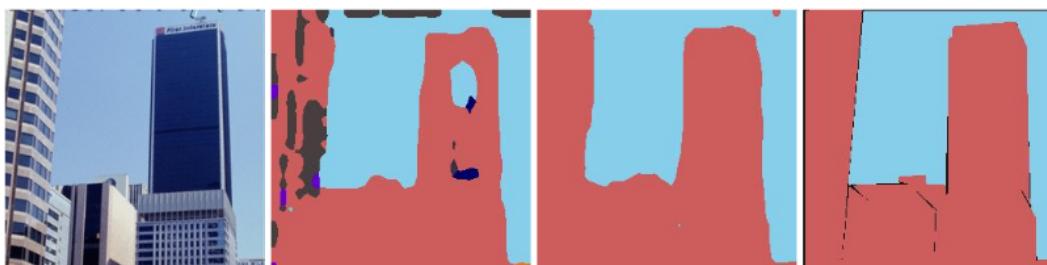


Input Image

CNN

DAG-RNN

Ground Truth



Input Image

CNN

DAG-RNN

Ground Truth

Outline

1 Goals

2 The first Neural Network: the Perceptron

3 Convolutional Neural Networks

4 Applications

- Image classification
- Pose, action detection
- Object detection and segmentation
- Scene labeling - Semantic segmentation
- Object tracking - videos

Object tracking

Object Detection in a video:

- YOLO <https://pjreddie.com/darknet/yolo/>
- YOLO - James Bond <https://www.youtube.com/watch?v=V0C3huqHrss>

Deep Learning - Many many other applications !

- ① Image and video (super-resolution, 3D, Image captioning), medical applications...
- ② Self Driving cars (scene segmentation, etc.)
- ③ In NLP (language, text), automatic translation, voice recognition, text generation, next word completion ... (Virtual assistants)
- ④ Recommendation systems...
- ⑤ For time series, complex datasets too...

But also beyond the supervised settings: using adversarial networks (GANS) for generation of images, faces, voice, etc.

Click on the person who is real.



<http://www.whichfaceisreal.com/>

Limits and problems

Limits of Deep Learning:

Limits and problems

Limits of Deep Learning:

- Interpretability

Limits and problems

Limits of Deep Learning:

- Interpretability
- Size of the models: Applicability, Ethical, Energy & Ecological concerns (Ai has both a positive and a huge negative carbon footprint)!

Limits and problems

Limits of Deep Learning:

- Interpretability
- Size of the models: Applicability, Ethical, Energy & Ecological concerns (Ai has both a positive and a huge negative carbon footprint)!
- Theory !

Limits and problems

Limits of Deep Learning:

- Interpretability
- Size of the models: Applicability, Ethical, Energy & Ecological concerns (Ai has both a positive and a huge negative carbon footprint)!
- Theory !
- Ethical and practical concerns !

Deep Learning: failures

CNN are not always “robust” to adversarial attacks !

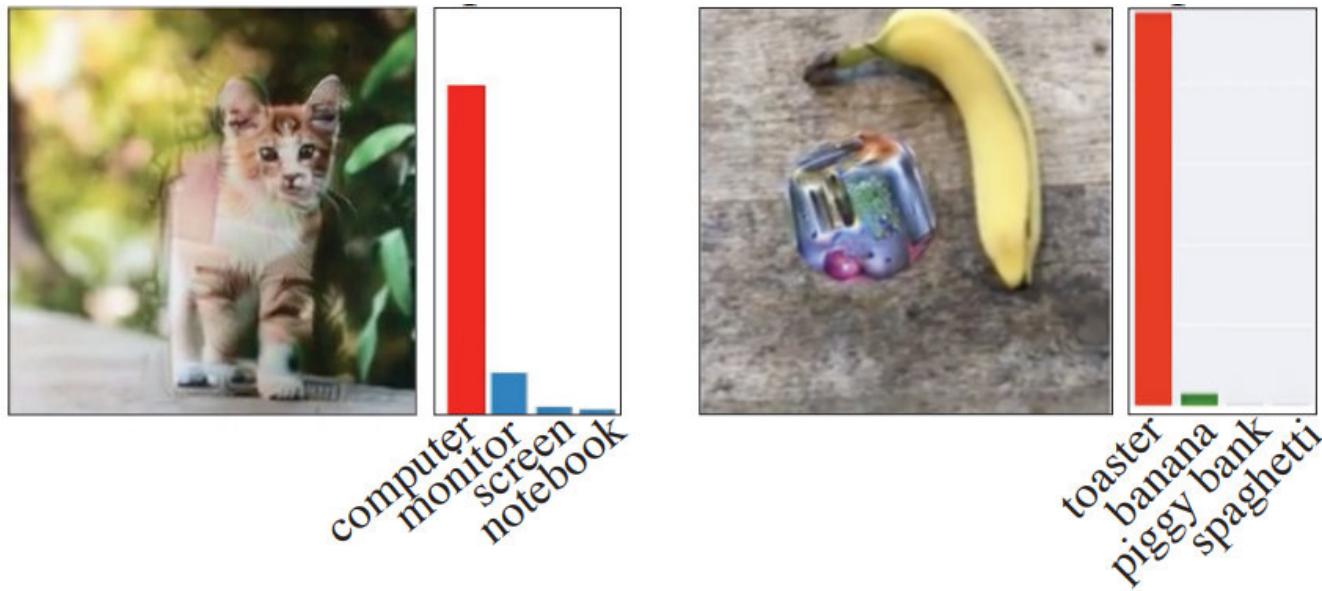
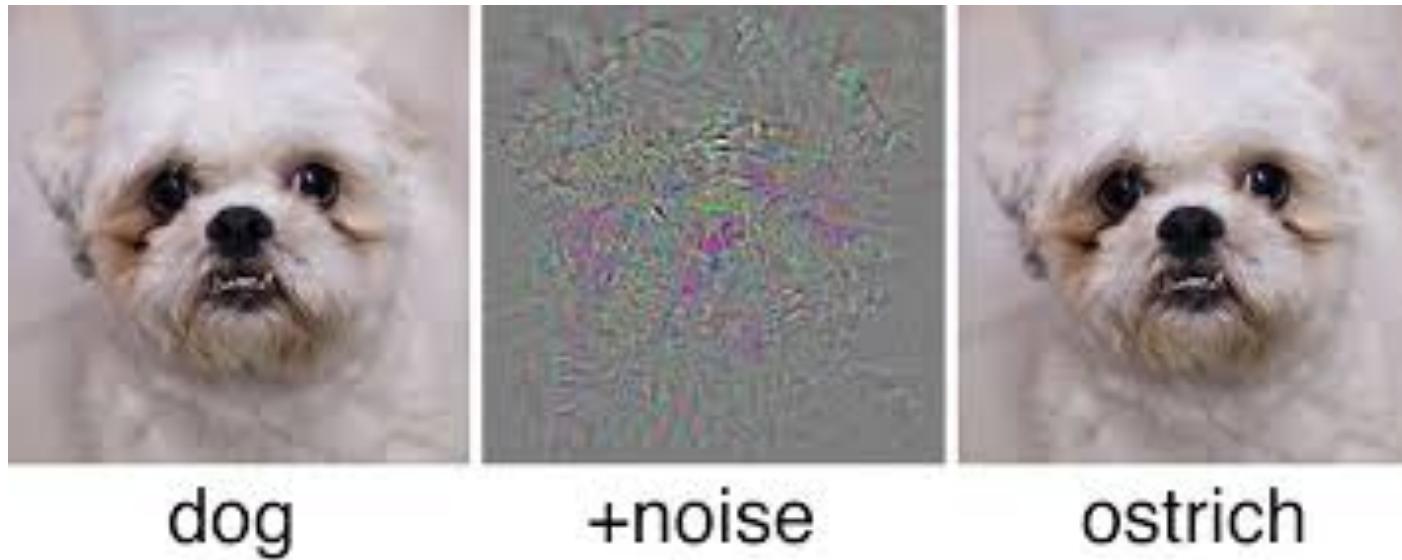


Figure: What happened here ??

Deep Learning: failures

CNN are not always “robust” to adversarial attacks !



Adding a small well chosen noise can completely fool a CNN !

↳ Can we trust deep networks on cars, medical applications, planes...?

Deep Learning: failures

There can be biases in Learning sets:



Deep Learning: failures

There can be biases in Learning sets, or AI can be manipulated...

- ① Chatbot becomes racist
- ② Google apologises for Photos app's racist blunder

Deep Learning: questionable applications

Are some applications just bad (they all already exist :() ?

- Deep Fakes: insert someone in a video.
- Underground exploration: finding new fossil resources.
- Voice imitation: what if I cannot check who is calling me.
- Military applications + automatic weapons.
- Mass surveillance.

Deep Learning: questionable applications

Are some applications just bad (they all already exist :() ?

- Deep Fakes: insert someone in a video.
- Underground exploration: finding new fossil resources.
- Voice imitation: what if I cannot check who is calling me.
- Military applications + automatic weapons.
- Mass surveillance.

Deep Learning: questionable applications

A challenging problem: chain of responsibility. The same tools are used for positive and negative applications, e.g.:

- mass surveillance
- tumor recognition

are both based on image recognition...

Good news ! the research community is very aware of the situation.

Some references:

- Asilomar AI Principles <https://www.oecd.org/going-digital/ai-intelligent-machines-smart-policies/conference-agenda/ai-intelligent-machines-smart-policies-oheigearthaigh.pdf>
<https://futureoflife.org/ai-principles/>
- European guidelines <https://ec.europa.eu/digital-single-market/en/news/ethics-guidelines-trustworthy-ai>
- AI for Good <https://ai4good.org/>

Conclusion: Deep Learning in one slide

- **How does it work:**

- ▶ Automatically learn representations of observations
- ▶ Learn highly non-linear models.

- **What does it require:**

- ▶ Large datasets with structure
- ▶ Computational power

- **Why now:**

- ▶ Combination of the 2 points above
- ▶ investment !

- **Some Applications**

- ▶ Image classification; object / face recognition
- ▶ Self driving cars
- ▶ Automatic Translation, Information extraction
- ▶ Caption Generation
- ▶ Ads, recommendation systems, etc.

Goals: Understanding core concepts of Deep Learning.

- When and how it works on paper (data, models, architecture)
- How to implement a simple neural network with Python.
- Overview of some of the main applications and challenges.

