

# Deep Learning

Aymeric DIEULEVEUT

May 2024

## 1 Goals

## 2 The first Neural Network: the Perceptron

## 3 Convolutional Neural Networks

## 4 Applications

- Image classification
- Pose, action detection
- Object detection and segmentation
- Scene labeling - Semantic segmentation
- Object tracking - videos

# Outline

## 1 Goals

## 2 The first Neural Network: the Perceptron

## 3 Convolutional Neural Networks

## 4 Applications

- Image classification
- Pose, action detection
- Object detection and segmentation
- Scene labeling - Semantic segmentation
- Object tracking - videos

# Deep Learning in one slide

- **How does it work:**

- ▶ Automatically learn representations of observations
- ▶ Learn highly non-linear models.

- **What does it require:**

- ▶ Large datasets with structure
- ▶ Computational power

- **Why now:**

- ▶ Combination of the 2 points above
- ▶ investment !

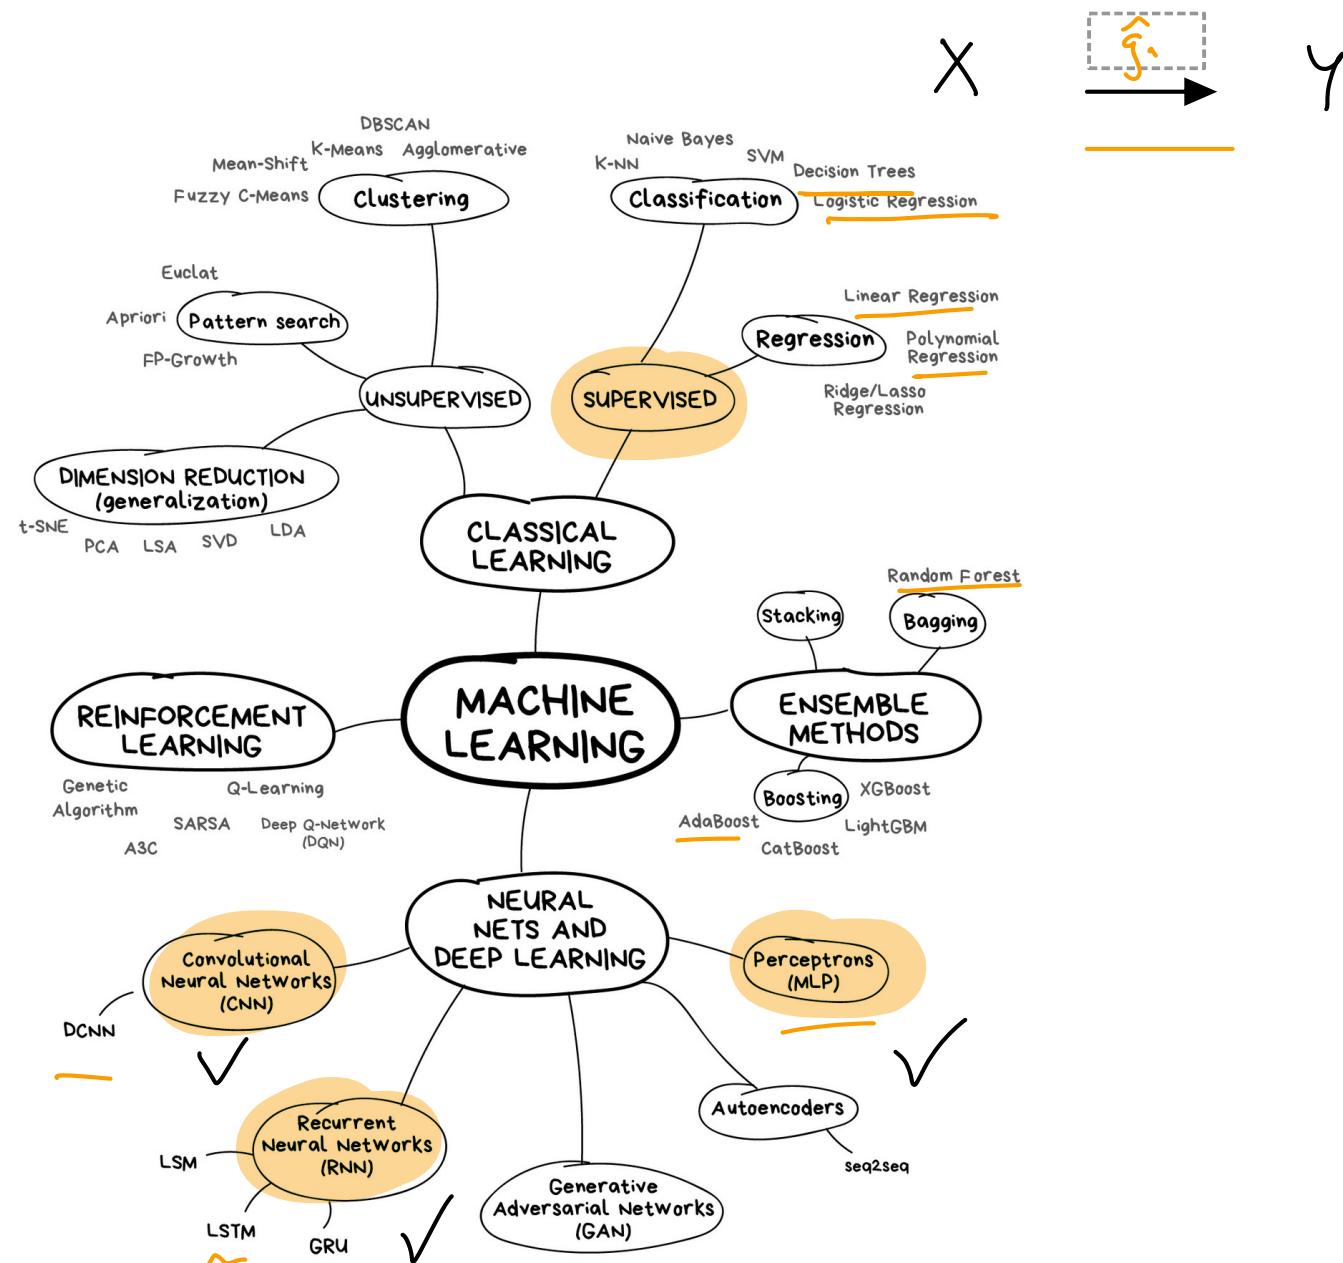
- **Some Applications**

- ▶ Image classification; object / face recognition
- ▶ Self driving cars
- ▶ Automatic Translation, Information extraction
- ▶ Caption Generation
- ▶ Ads, recommendation systems, etc.

**Goals:** Understanding core concepts of Deep Learning.

- When and how it works on paper (data, models, architecture)
- How to implement a simple neural network with Python.
- Overview of some of the main applications and challenges.

# Machine Learning

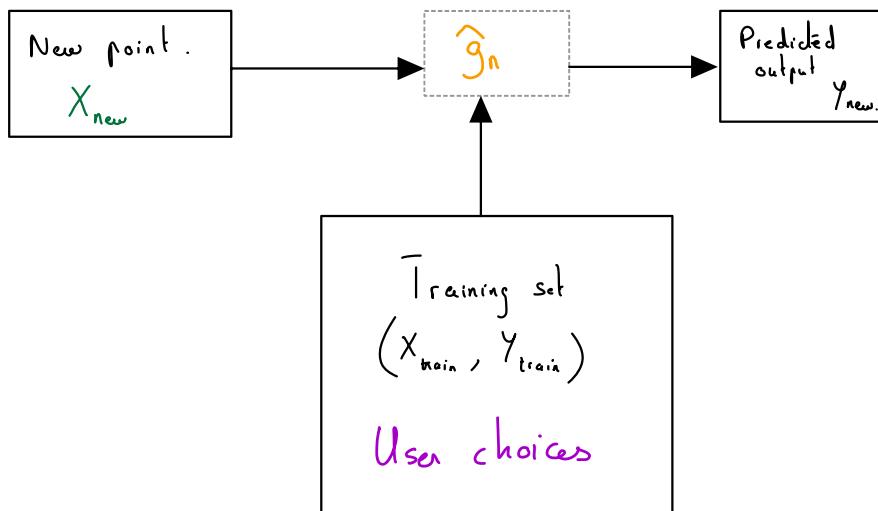


# Summary of previous lectures on Machine Learning

Take home message (in 2 words) : generalization + overfitting

## ① General framework: loss, risk, generalization risk vs training risk

$$R(\hat{g}_n) = \mathbb{E}[\ell(\hat{g}_n(x_{\text{new}}), y_{\text{new}})]$$



Loss function, Generalization Risk and data-dependent predictors

### Loss Function

- Loss function:  $\ell(y, g(x))$  quantifies the quality of the prediction  $g(x)$  of  $y$ , e.g.:
  - ▶ 0-1 Loss (classification):  $\ell(y, g(x)) = 1_{y \neq g(x)}$ ,  $y \in \mathcal{Y} = \{-1, 1\}$
  - ▶ Quadratic Loss (regression):  $\ell(y, g(x)) = \|y - g(x)\|^2$ ,  $y \in \mathbb{R}^d$

### Risk of a Decision Rule = average loss

- Risk:  $R(g) = \mathbb{E}[\ell(Y, g(x))] = \int \ell(y, g(x)) d\mathbb{P}(x, y)$ , e.g.:
    - ▶ 0-1 Risk (classification):  $R(g) = \mathbb{P}(Y \neq g(x))$
    - ▶ Quadratic Risk (regression):  $R(g) = \mathbb{E}\|Y - g(x)\|^2$
- ↪  $R$  is also referred to as *Generalization risk*, or *True risk*.

⚠ The distribution  $\mathbb{P}$  is unknown.

### Data-dependent predictors

- Training set:  $\mathcal{D}_n = \{(X_1, Y_1), \dots, (X_n, Y_n)\}$  (*i.i.d.~* $\sim \mathbb{P}$ )
- Goal: build a **data dependent predictor / classifier**  $\hat{g}_n$  based on the training data
- That has a **minimal generalization risk**  $R(\hat{g}_n)$ .

$R(\hat{g}_n)$  = average loss on a "new point"  $(X, Y) \sim \mathbb{P}$ , independent from  $D_n$

↪ Our goal is to predict well on inputs/outputs pairs that we have not seen in the dataset, i.e. *generalize well*.

THM 1.

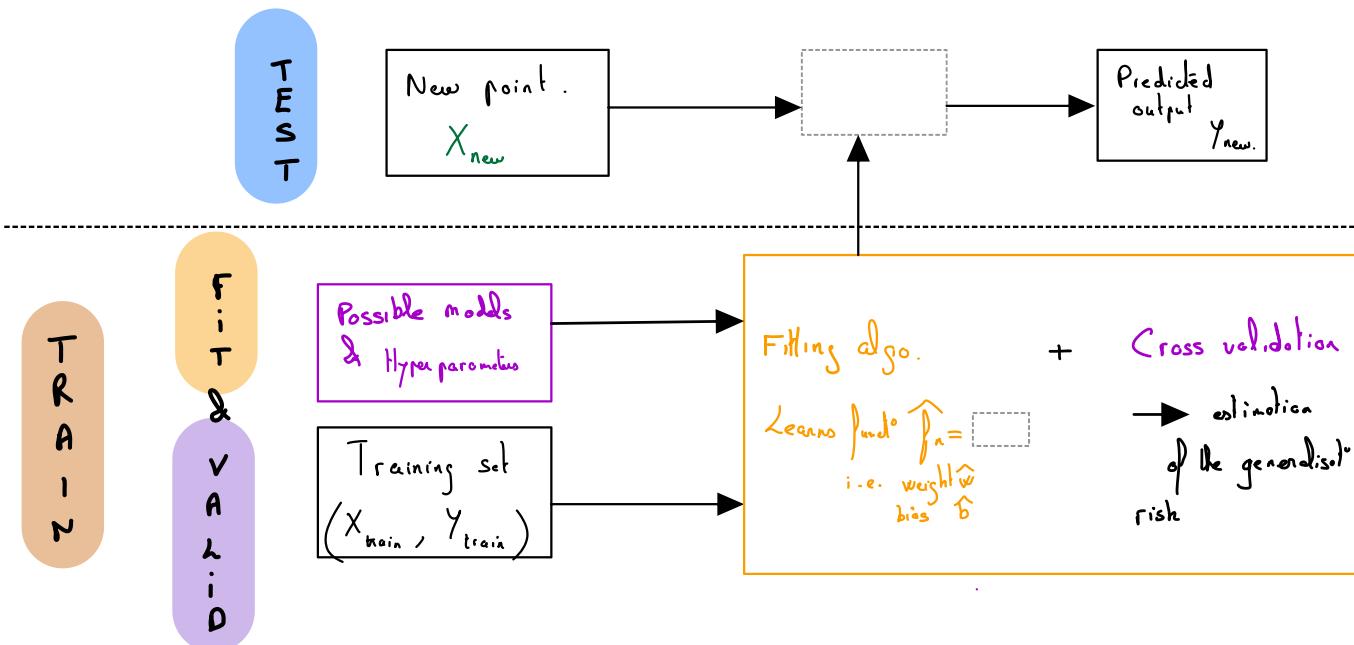
# Summary of previous lectures on Machine Learning

THM : Cross-Validation

① General framework: loss, risk, generalization risk vs training risk

② SML Pipeline

- ▶ Train (Fit + validation) / Test / Deployment
- ▶ Evaluation of validation scores through cross validation to pick the best method/hyperparams



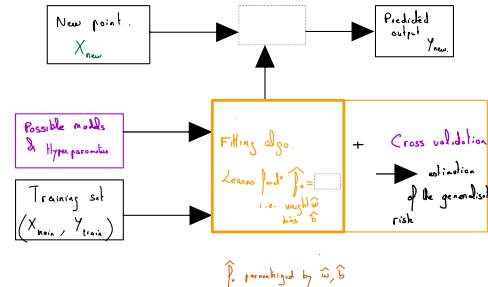
# Summary of previous lectures on Machine Learning

① General framework: loss, risk, generalization risk vs training risk

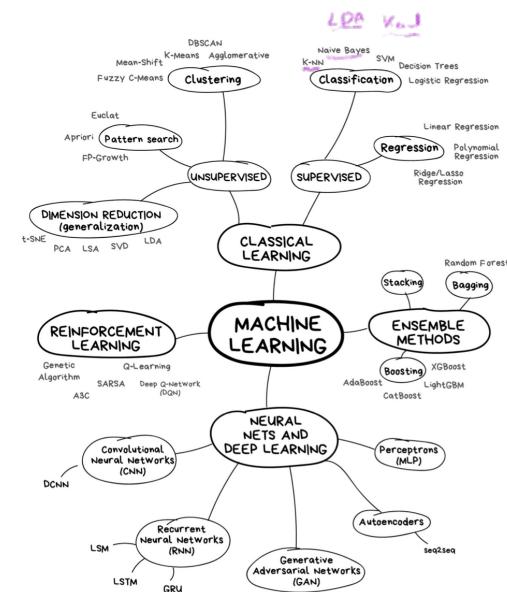
② SML Pipeline

- ▶ Train (Fit + validation) / Test / Deployment
- ▶ Evaluation of validation scores through cross validation to pick the best method/hyperparams

③ Methods: Linear, Logistic regression, Decision Trees, Random Forests, Boosting



Fitting algo.  
Learn  $f$  s.t.  $\hat{Y}_n = \boxed{\quad}$   
i.e. weight  $\hat{w}$  bias  $\hat{b}$



# Summary of previous lectures on Machine Learning

- ① General framework: loss, risk, generalization risk vs training risk
- ② SML Pipeline
  - ▶ Train (Fit + validation) / Test / Deployment
  - ▶ Evaluation of validation scores through cross validation to pick the best method/hyperparams
- ③ Methods: Linear, Logistic regression, Decision Trees, Random Forests, Boosting
- ④ Opening the black box: for each method
  - ▶ How to predict (↔ interpretability)
  - ▶ How fit is made (↔ speed, complexity) [Computer learning parameters]
  - ▶ Important Hyper-parameters [User chosen]
  - ▶ Drawbacks and Strengths
  - ▶ Underfitting, overfitting risks (↔ Performance)  
→ regularization techniques (avoiding overfitting by modifying the method)

$\hat{y}_n =$ [ ]	: Lin. Reg	Polynomial Linear reg	Logistic Reg	Dec Trees	RF/ Boosting
How to predict	✓ → ~		✓	✓	✓
How to fit	explicit →		algo. (opt)	✓	✓
Key hyperparameters	none (regulariz.) simple stable interpretation		not many "	✓	✓
⊕			"	✓	✓
⊖	underfitting risk		"	✓	✓
Overfitting / Underfitting → regularization techniq	mostly under depends on d vs n if d too large, overfit Lasso ✗ Ridge ✗	much more overfit	"	✓	✓
Approach	ERM	ERM	ERM	not ERM	not ERM

✓ = seen in Lecture.

\* Advanced

## ERM: empirical risk minimisation.

\* Define class of function  $\{f\}_{n \text{ possible}} \subseteq \mathcal{C}$

Find  $\min_{f \in \mathcal{C}} \left\{ \widehat{R}_n(f) := \frac{1}{n} \sum_{i=1}^n l(p_k^i, y_i^i) \right\}$   
on the train set

choice 1 ① define loss  $\rightarrow$  general risk

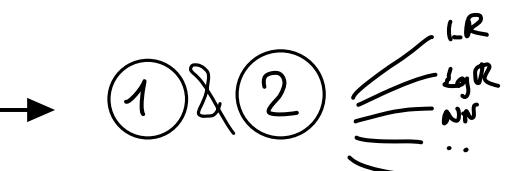
empirical risk  $E_R$   $\widehat{R}_n$

choice 2 ② choose a class  $\mathcal{C}$

ERM = ③  $\rightarrow$  find the best  $f \in \mathcal{C}$  for the  $E_R$   
ie on train set

Verif) ④ Check overfitting (risk is high if  $C$  large)

ERM encapsses:



seen	seen
* Linear Reg ✓	* Poly Log Reg ~
* Logistic Reg ✓	* Poly Lin Reg ~
* NN ☺	* Kernels ✗

# Outline

1 Goals

2 The first Neural Network: the Perceptron

3 Convolutional Neural Networks

4 Applications

- Image classification
- Pose, action detection
- Object detection and segmentation
- Scene labeling - Semantic segmentation
- Object tracking - videos

# From logistic regression to Multi Layer Perceptron 1

Supervised ML problem:

- ① Data  $D_n = (X_i, Y_i)_{i=1,\dots,n}, X_i \in \mathbb{R}^d, Y_i \in \{0, \dots, K\}$
- ② Goal: Predict a the label for a new point.

2 approaches:

## Generative approach

- Define a model on your data (e.g., logit model).
- Estimate parameter (likelihood maximization).

→ LDA, QDA .

→ statist. based.

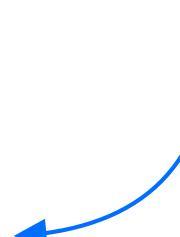
## Discriminative approach

- No statistical model !
- Define a loss function  $\ell$
- Goal:

$$\min_{f \in \mathcal{F}} \mathbb{E}[\ell(f(X), Y)]$$

- Approach:

$$\min_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n [\ell(f(X_i), Y_i)]$$



# From logistic regression to Multi Layer Perceptron 1

Supervised ML problem:

- ① Data  $D_n = (X_i, Y_i)_{i=1,\dots,n}, X_i \in \mathbb{R}^d, Y_i \in \{0, \dots, K\}$
- ② Goal: Predict a the label for a new point.

2 approaches:

## Generative approach

- Define a model on your data (e.g., logit model).
- Estimate parameter (likelihood maximization).



## Discriminative approach - ERN

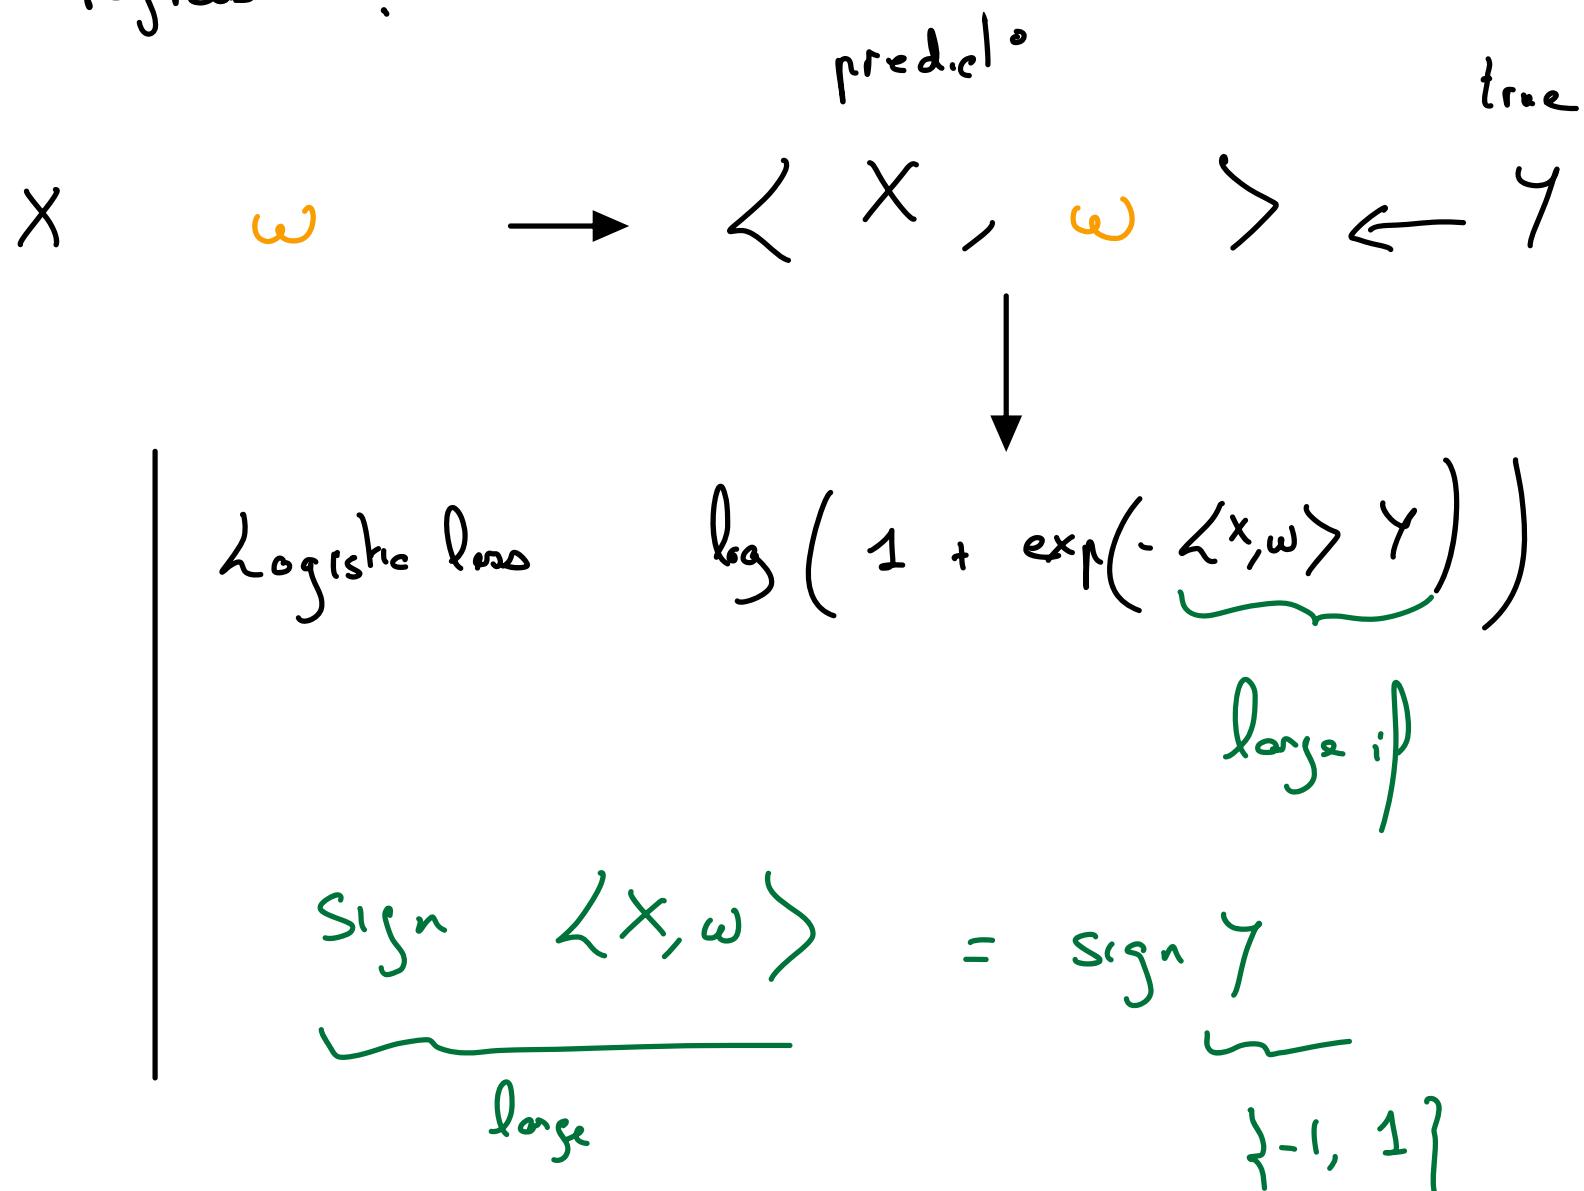
- No statistical model ! ✓
  - Define a loss function  $\ell$  ✓
  - Goal:
- $$\min_{f \in \mathcal{F}} \mathbb{E}[\ell(f(X), Y)] \rightarrow \text{gen risk}$$
- Approach: ↑ overfitting ! ⚠
- $$\min_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n [\ell(f(X_i), Y_i)] \quad \checkmark$$

Logistic regression can be seen both ways !

$$\min_{w \in \mathbb{R}^d} \frac{1}{n} \sum_{i=1}^n \log(1 + \exp(-Y_i w^T X_i))$$

# Let's NN playground: [Link](#)

Logistic regression.



# Let's NN playground: Link

FEATURES

Which properties do you want to feed in?

blue is  $>0$

orange is  $<0$

$X_1$

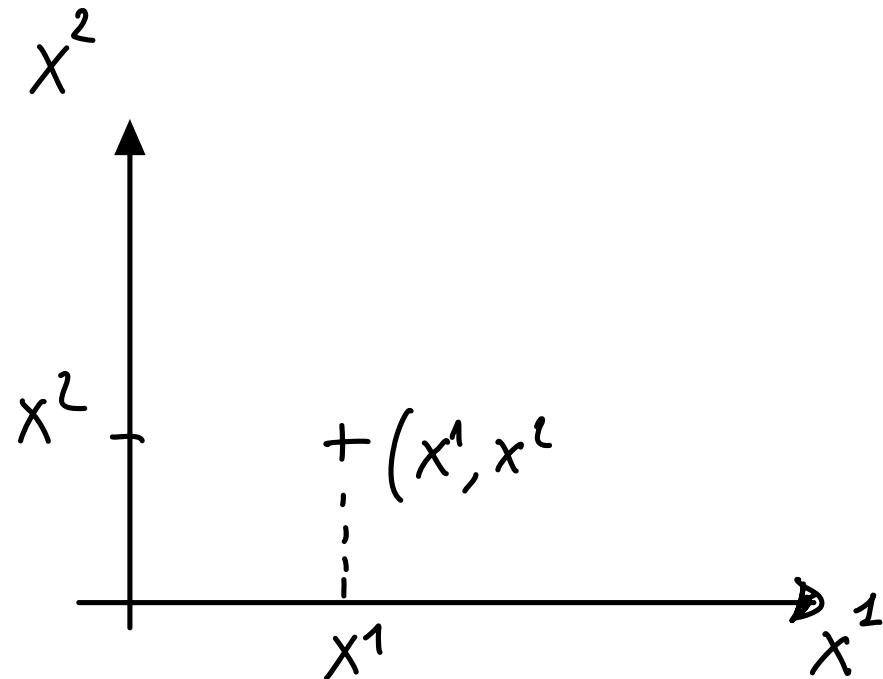
$X_2$

$X_{12}$   $\rightarrow (X^1)^2$

$X_{22}$   $\rightarrow (X^2)^2$

$X^1 \times X^2$

$\rightarrow 0 \rightarrow X^1 \times X^2 < 0 \text{ or } 0$



understanding the NNPG:

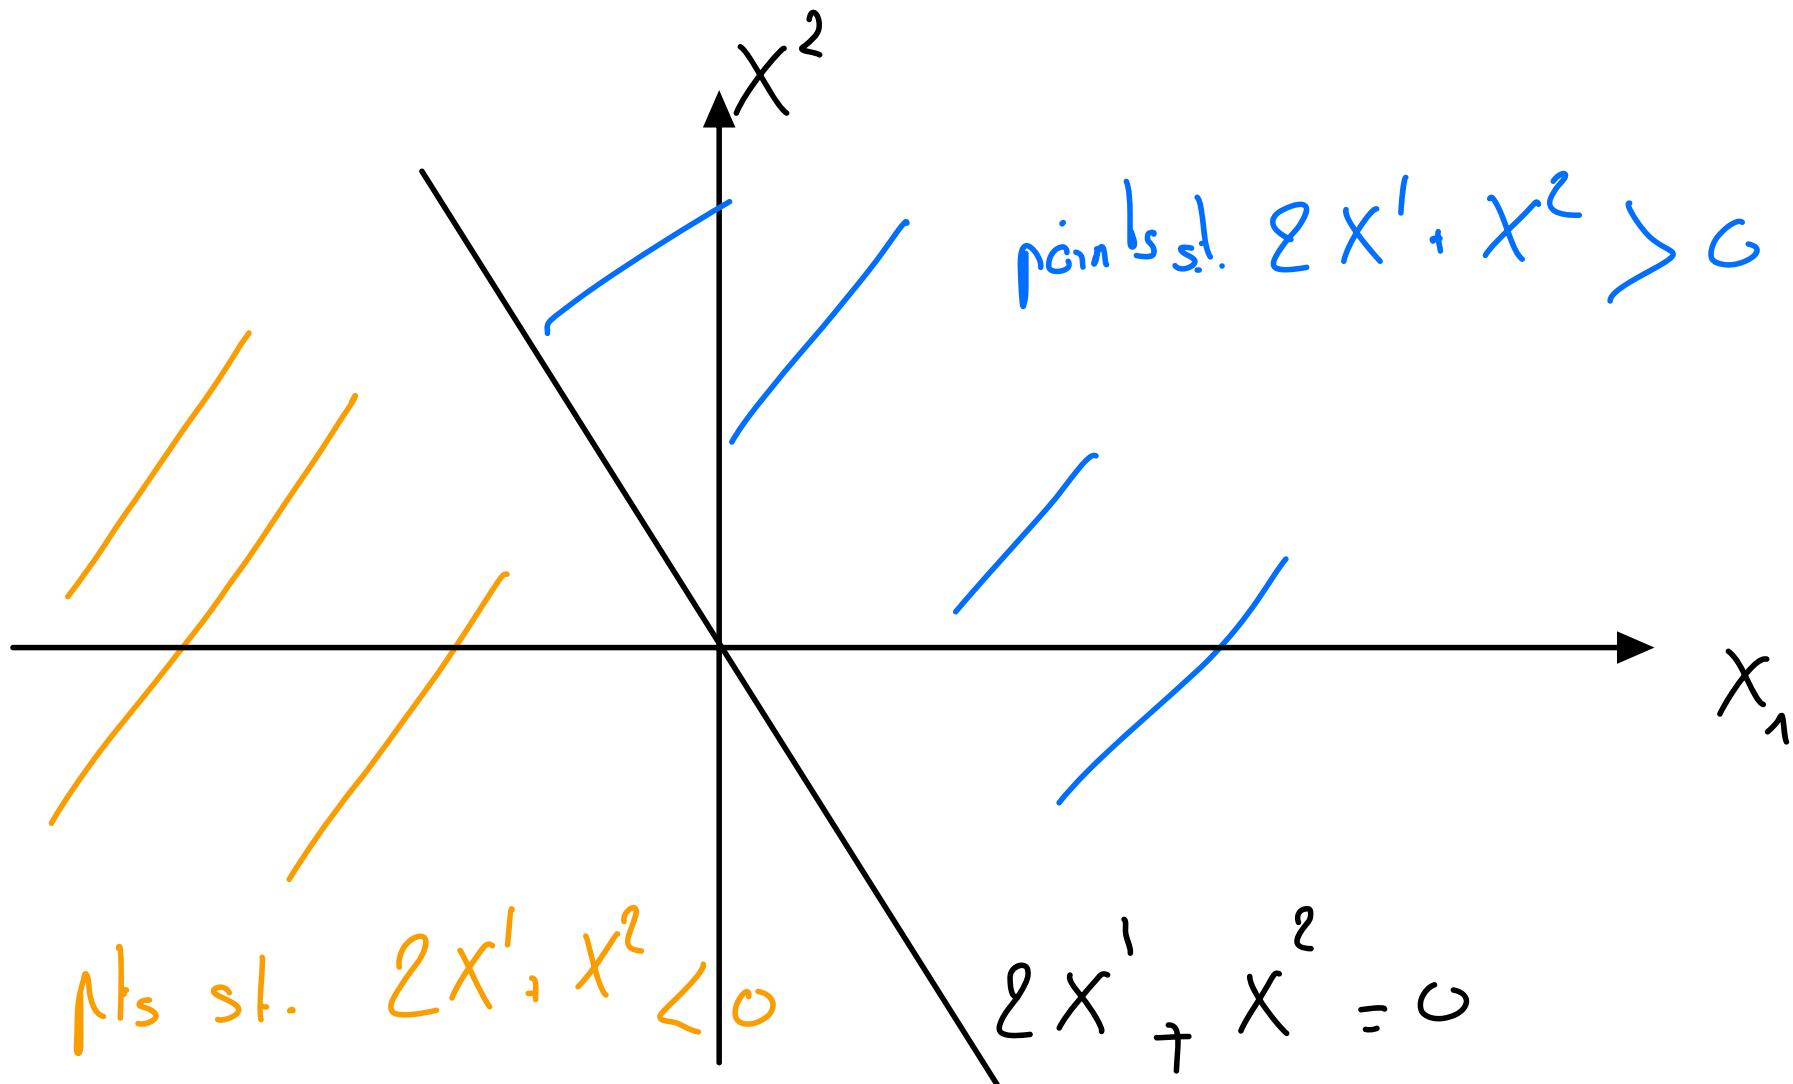
- \* lectures
- \* colors.

Let's NN playground: [Link](#)

$\otimes$  hand picked weights

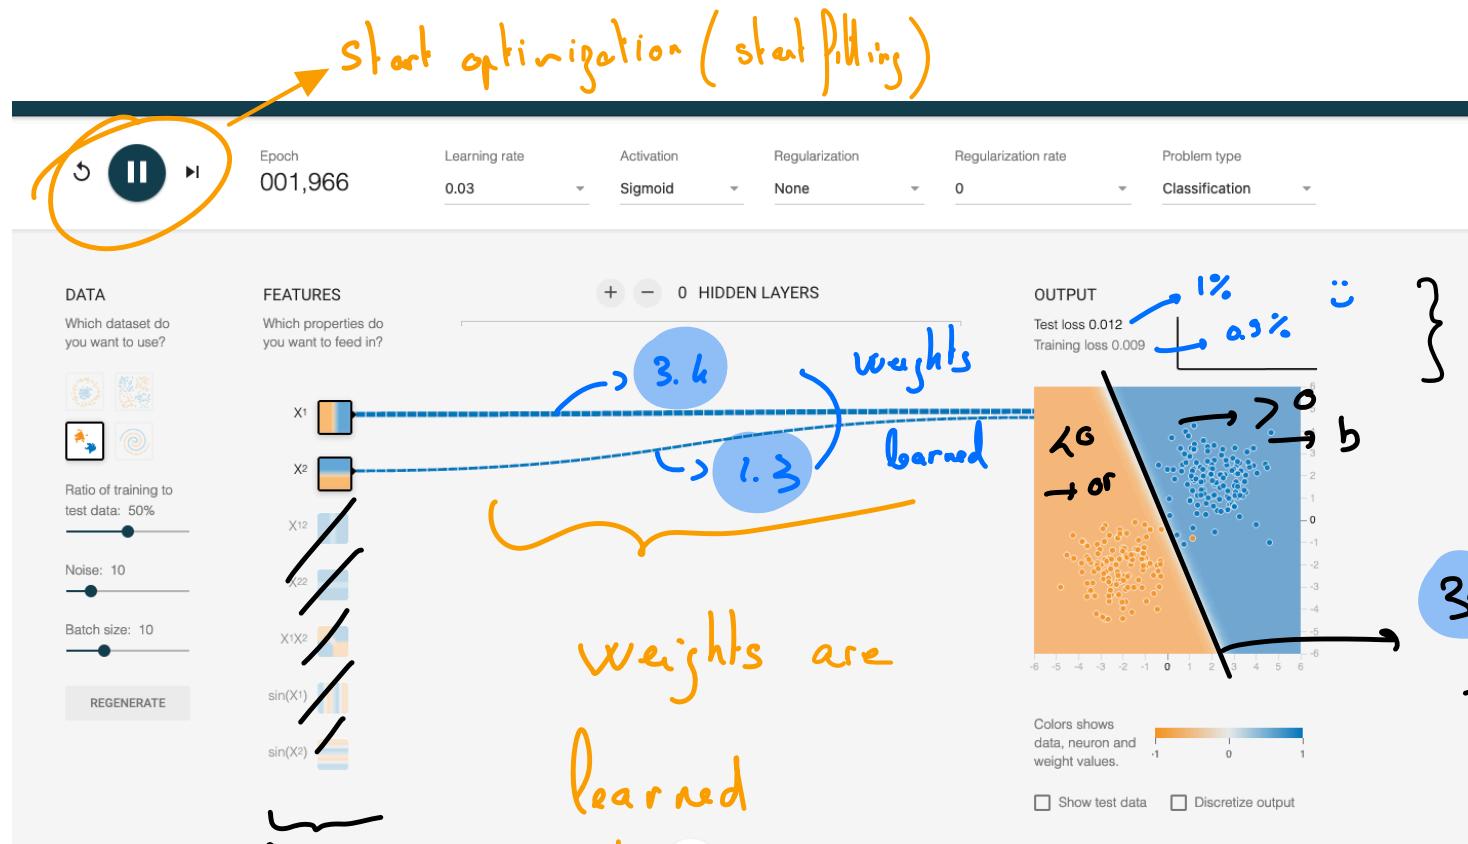
$$2x^1, 1x^2$$

$$\omega_1 \quad \omega_2$$



# Let's NN playground: Link

- \* Let us run the optim process, with 0 hid layer
- \* this is equivalent to logistic regression

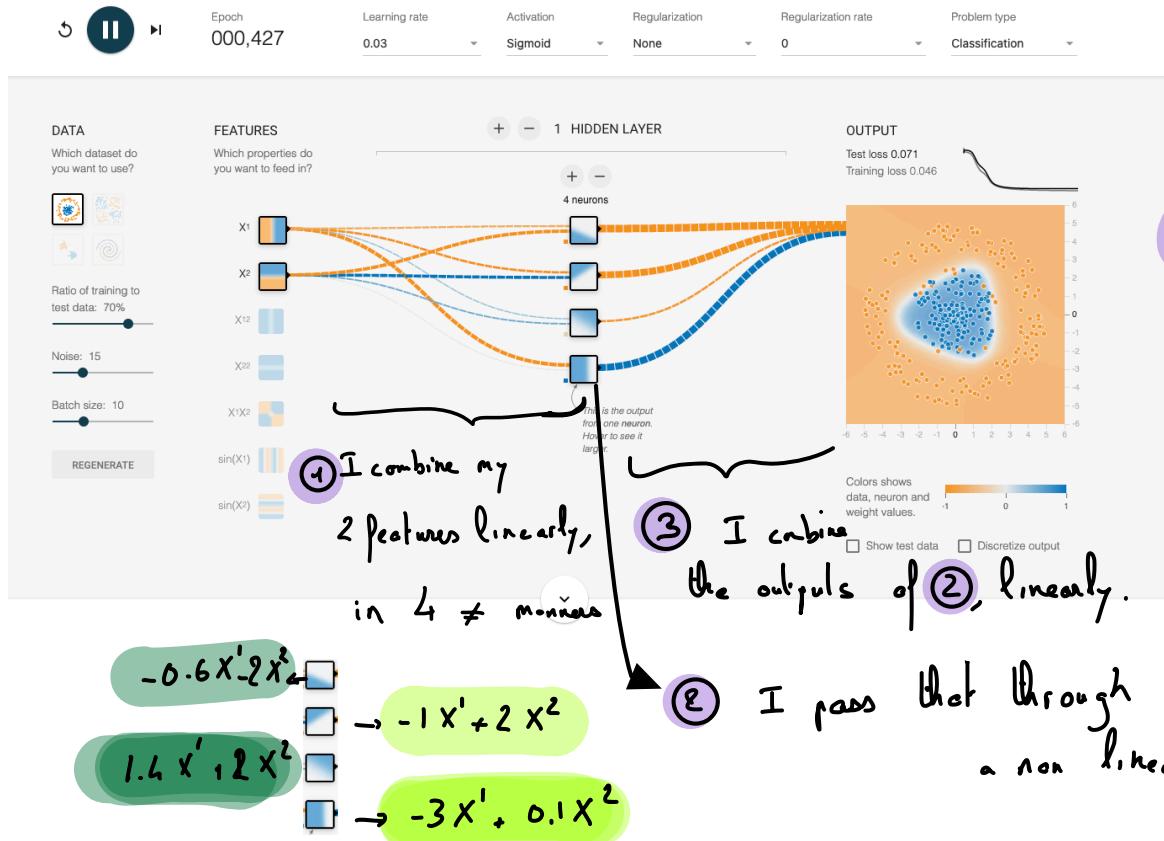


th:

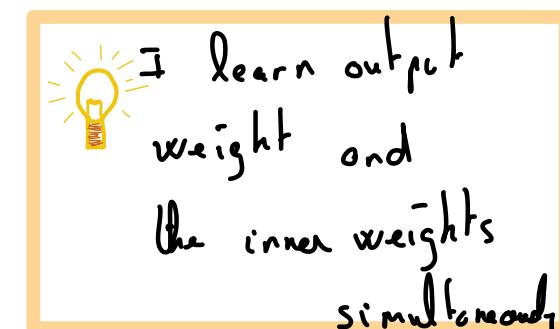
Logistic regression → stable : 2 ≠ init give the same solut° !!

# Let's NN playground: [Link](#)

- \* Let us move to non linearly sep data → Log reg fails
- \* Let us add more "neurons". → So does single neuron hidden layer



(4) Success!!



Take home msg → I have learned a non linear fit!! → I have learned 4 good features!!  $(\sigma(\quad), \sigma(\quad), \sigma(\quad), \sigma(\quad))$

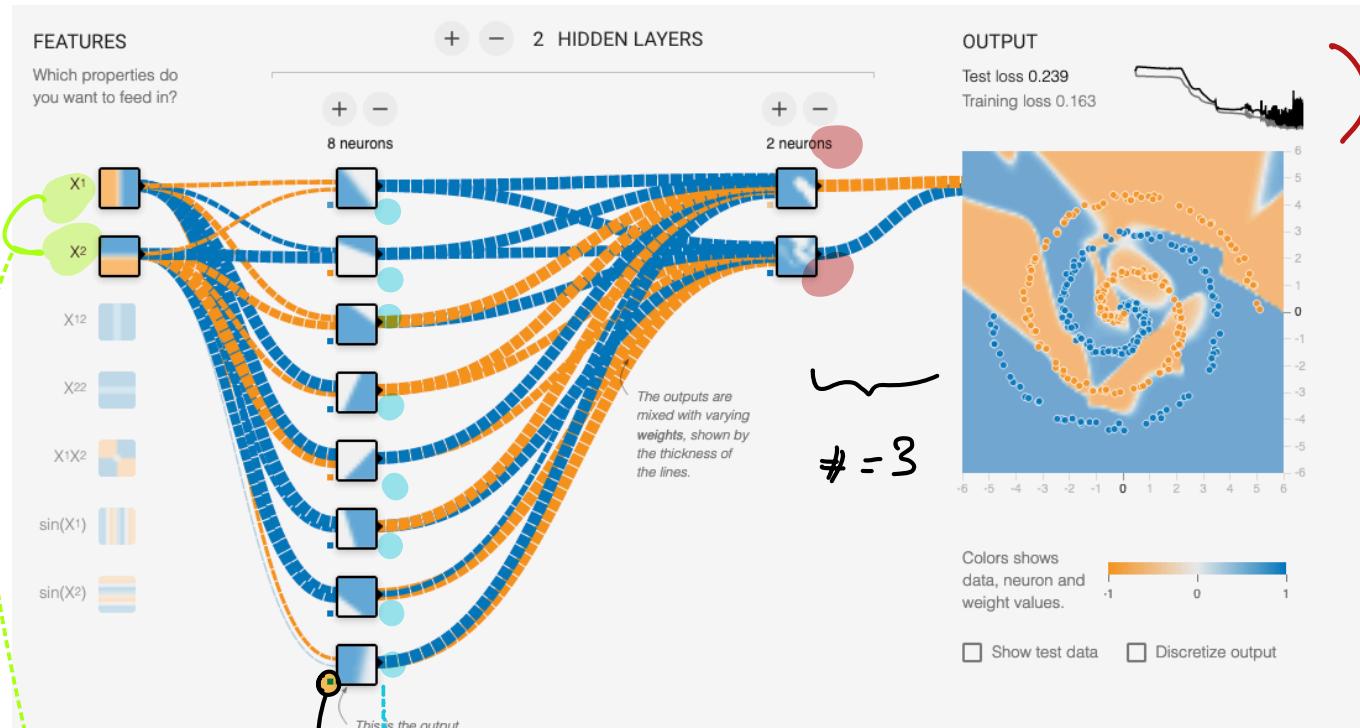
→ That can be linearly combined to obtain the output!!

# Let's NN playground: Link

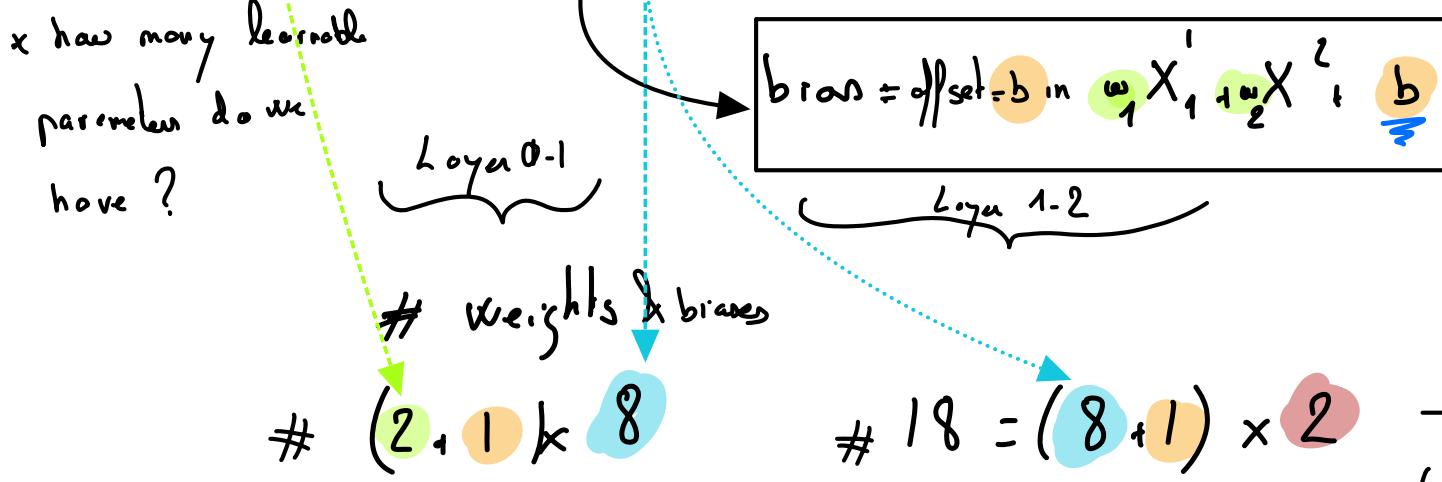
\* Let us move to more advanced data



\* Let us try with 8 neurons ?  
1 L  
x 2 L  
8N ?



x how many learnable parameters do we have ?



→ 45 parameters overall  
(vs GPT 4 :  $10^{10}$  +)

# Let's NN playground: Link Impact of activation:

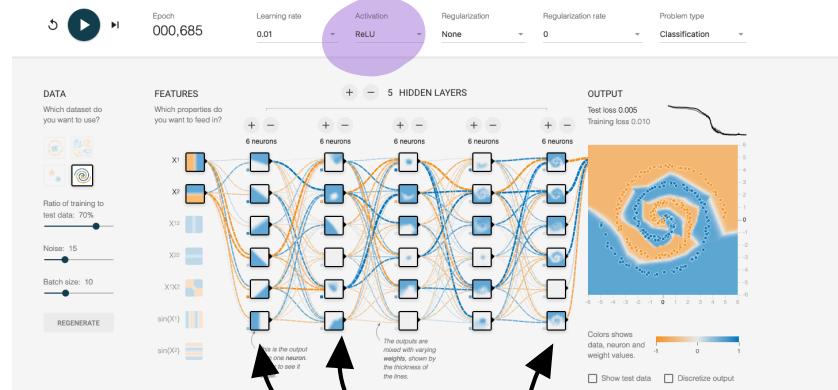
① Linear activat° → no learning of non linear behavior

② Sigmoid → gradient killer → alg does not progress ::

③ Tanh → improves on Sigmoid  
→ overfitting

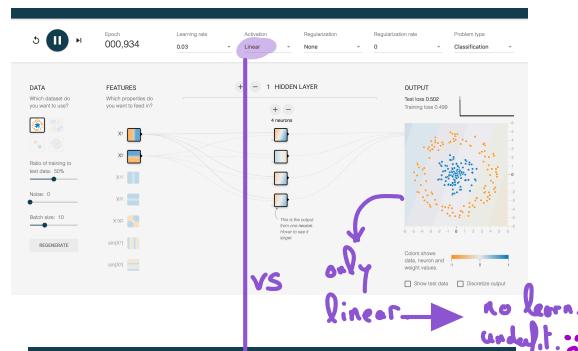
④ ReLU → work amazingly well here

④



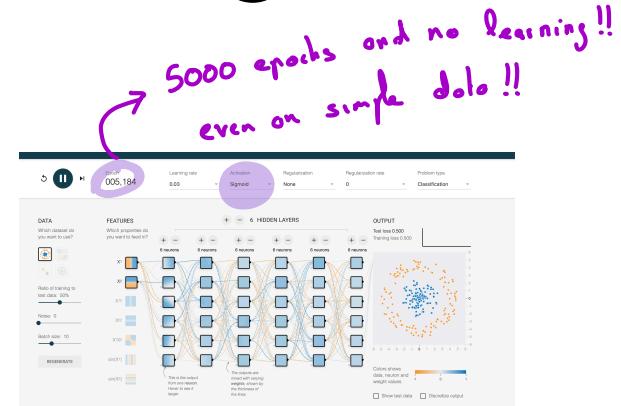
note the evolution of features learned on each layer!

①



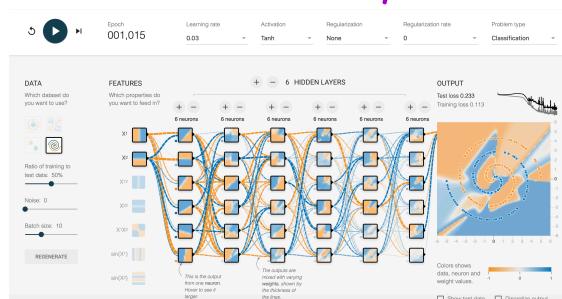
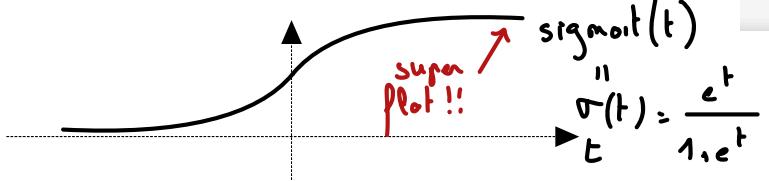
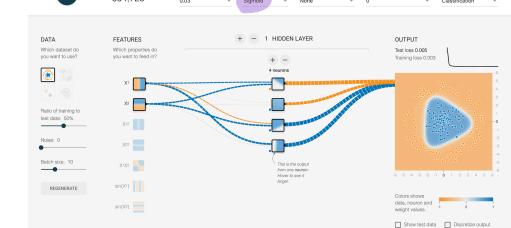
vs  
only linear → no learn. underfit...  
no learn. underfit...  
no learn. underfit...

②



③

now, learning happens  
→ but overfit



# Let's NN playground: Link

Optimization: \* iterative

↳ "million \$" \* batch: nb of points "looked at" at each iteration  
question \*

→ very complicated \* opt: global for logistic (convex function, unique minimum)  
unstable otherwise → non convex, many minima!

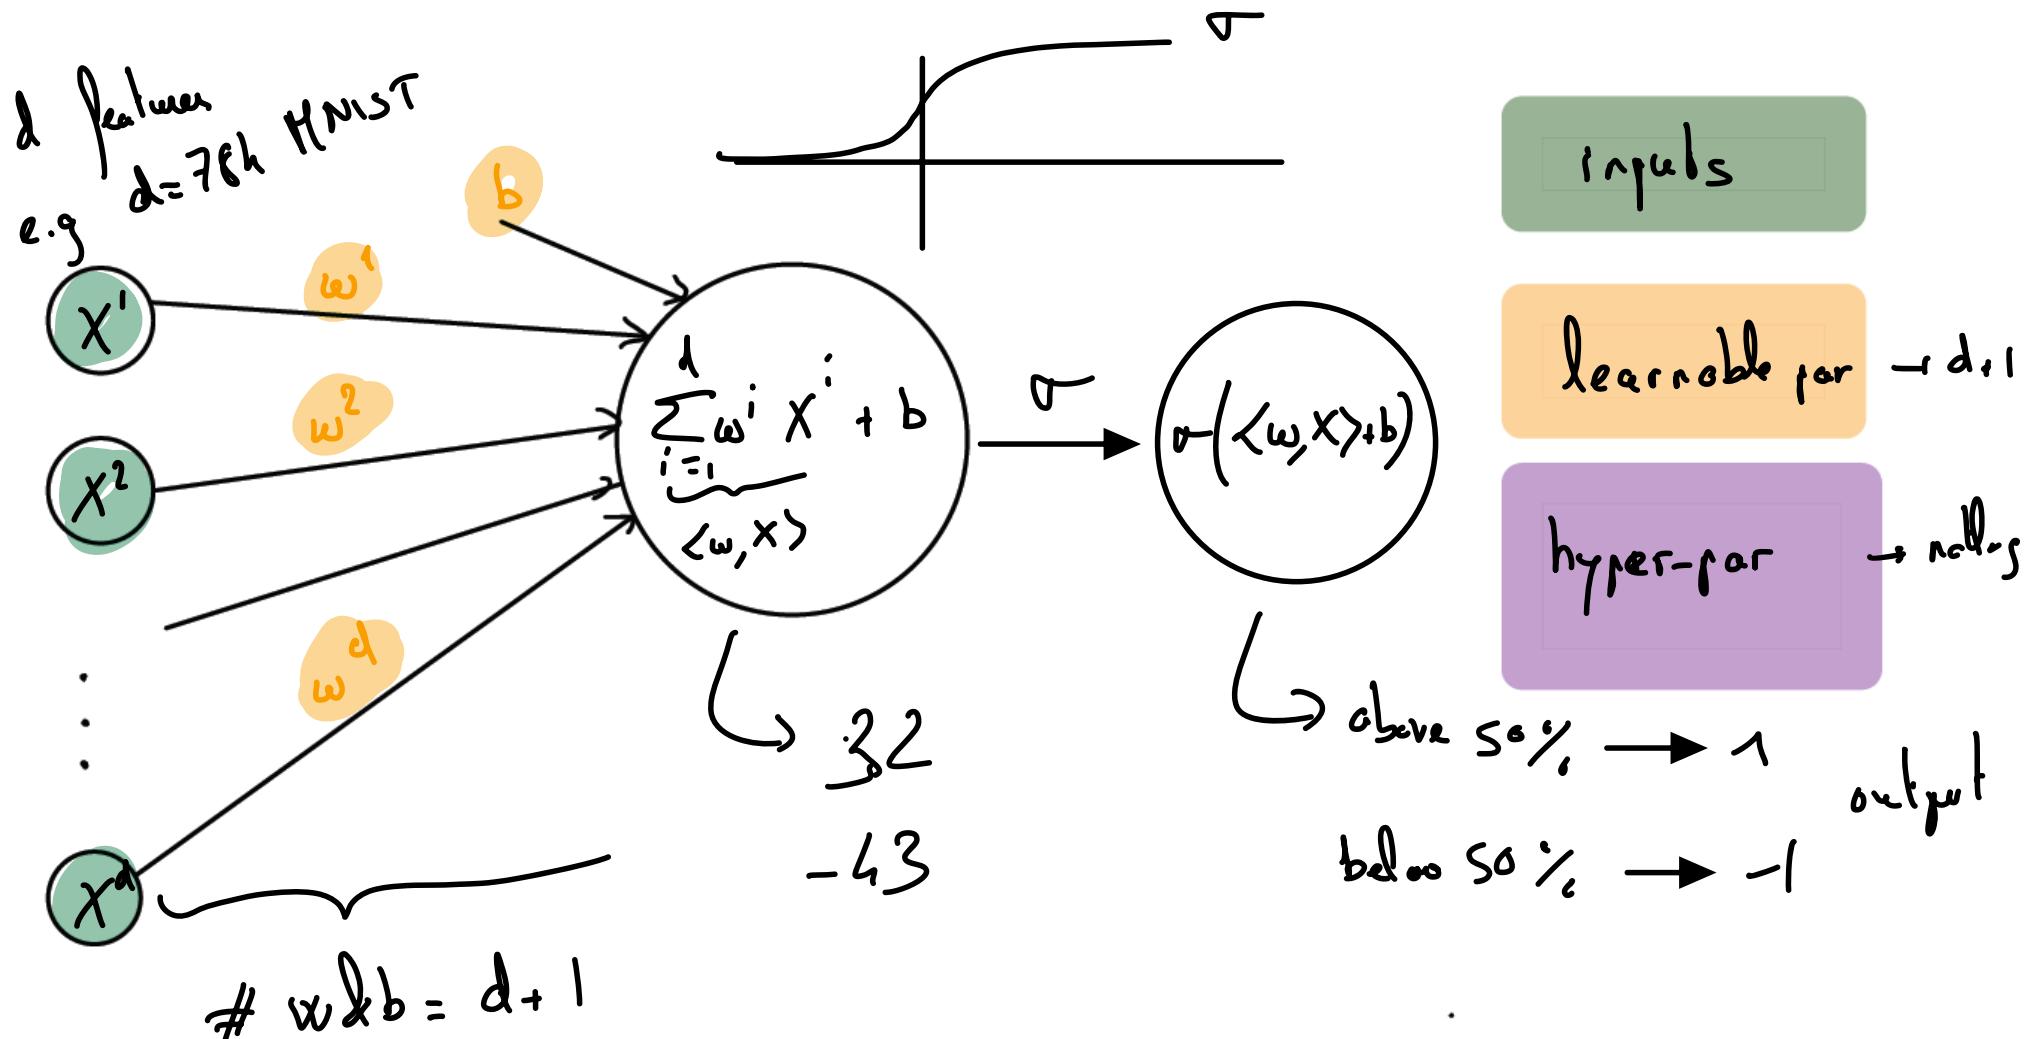
\* very long  
for architecture choice  
CV not possible  
→ pick student architecture

## Let's NN playground: [Link](#)

- \* Architect : "Fully-connected" here
  - \* We have visualized & discussed :
    - Neuron, layer, bias, weights, activation, optimization, learned features, etc
- Pro tip: try again the playground in one year  
add a task in your calendar not to forget

# From logistic regression to Multi Layer Perceptron 3

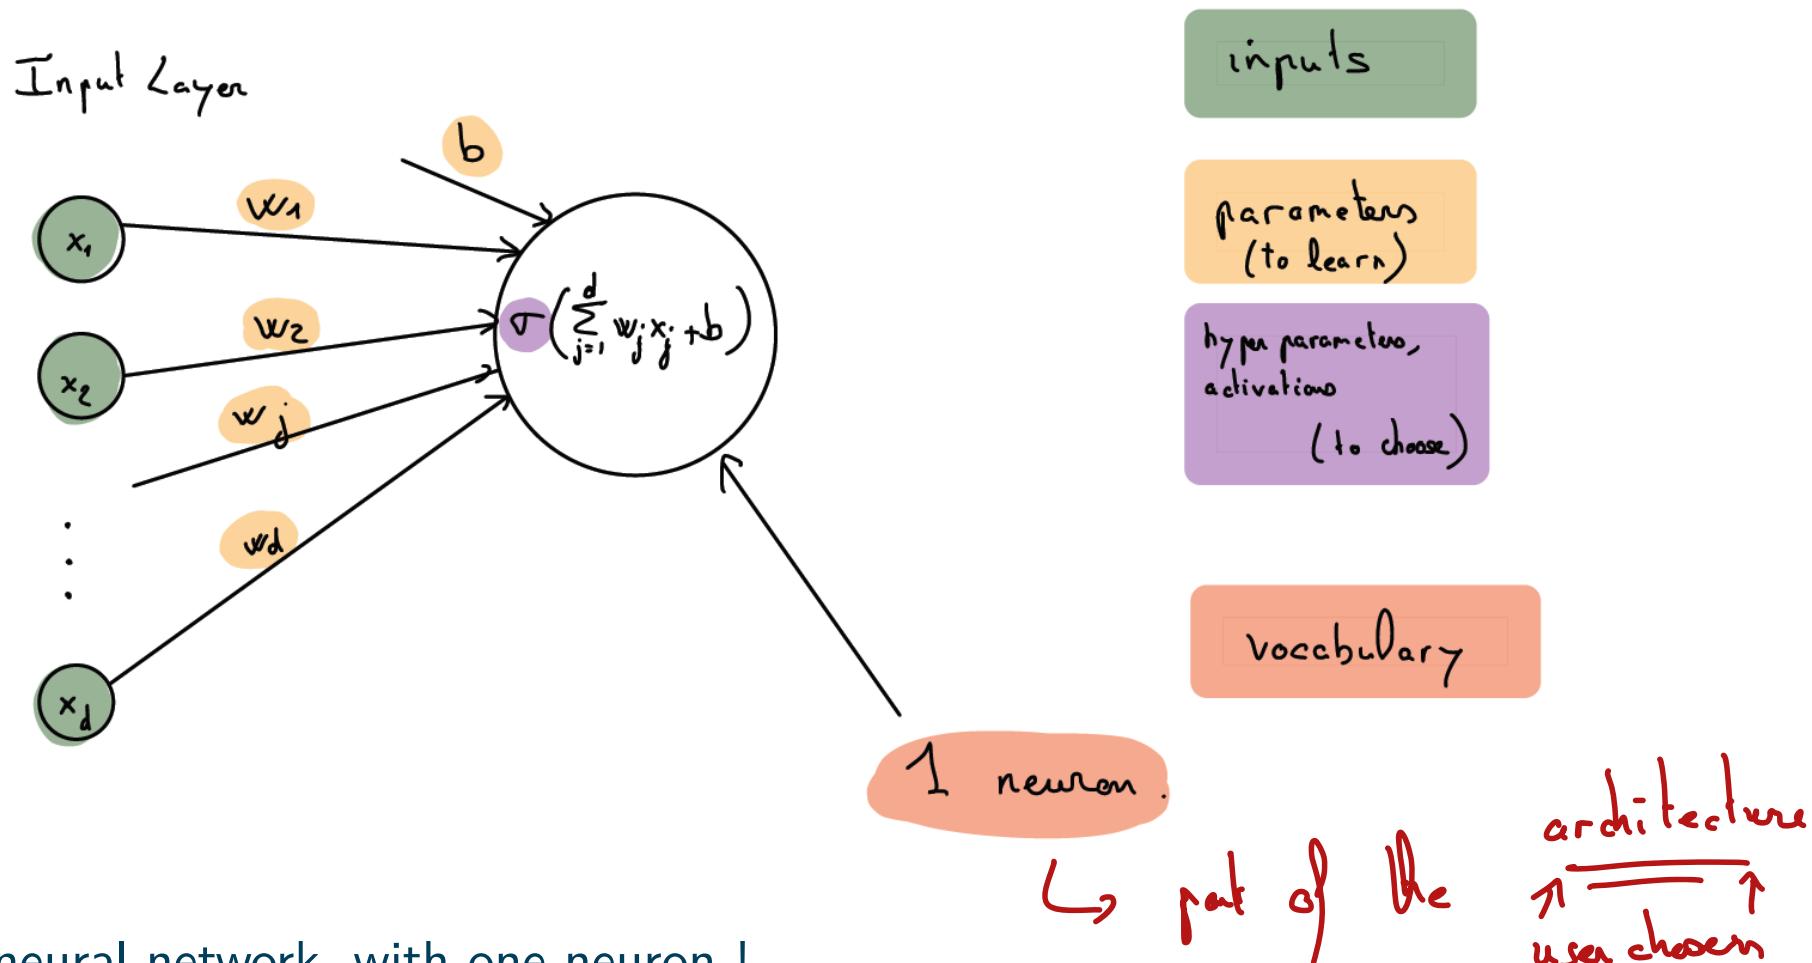
- Class of functions for logistic regression:  $\mathcal{F} = \{X \mapsto w^T X, w \in \mathbb{R}^{d+1}\}$ .
- Prediction for a new point is  $1_{\sigma(X_i) > 1/2}$
- Can be seen as predicting a probability  $\sigma(X_i)$  that the output is 1



This is a neural network, with one neuron !

# From logistic regression to Multi Layer Perceptron 3

- Class of functions for logistic regression:  $\mathcal{F} = \{X \mapsto w^T X, w \in \mathbb{R}^d\}$ .
- Prediction for a new point is  $1_{\sigma(X_i) > 1/2}$
- Can be seen as predicting a probability  $\sigma(X_i)$  that the output is 1

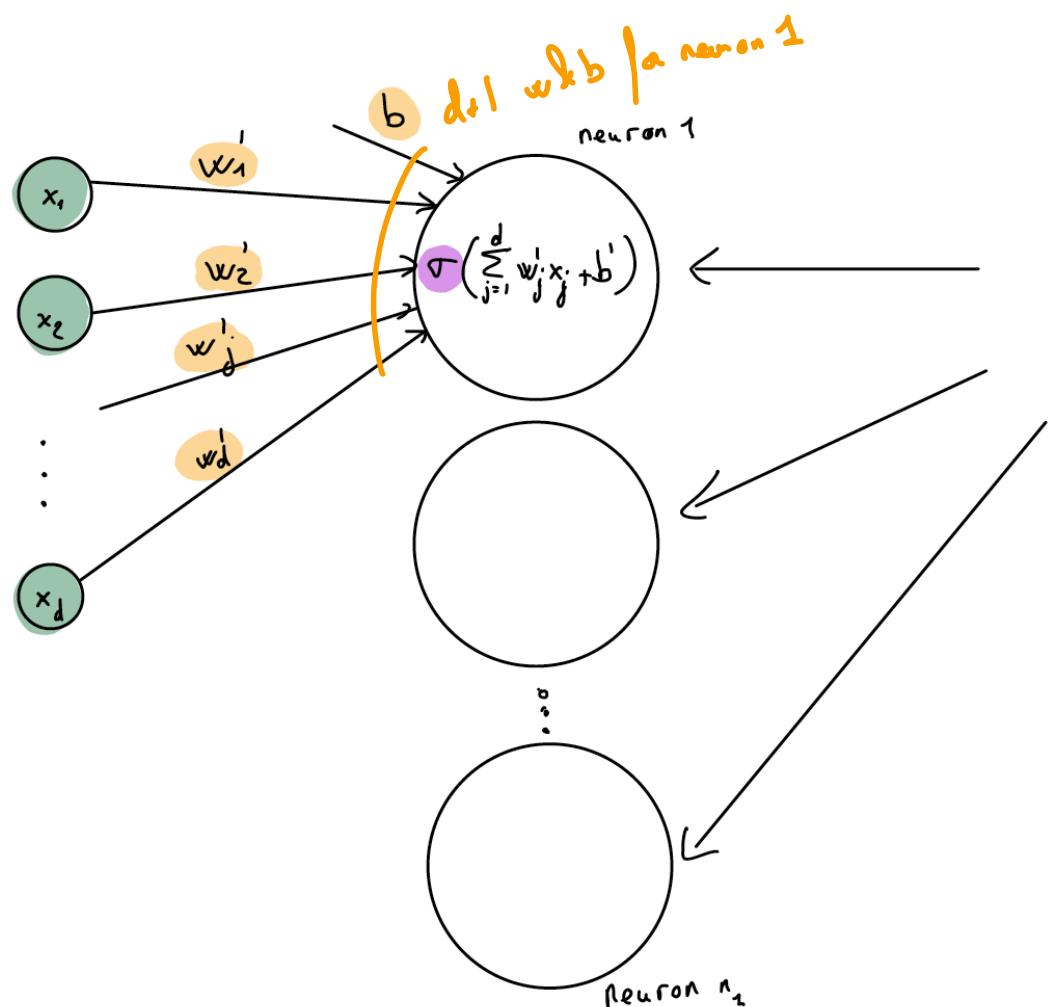


This is a neural network, with one neuron !

# From logistic regression to Multi Layer Perceptron 4

Extend to

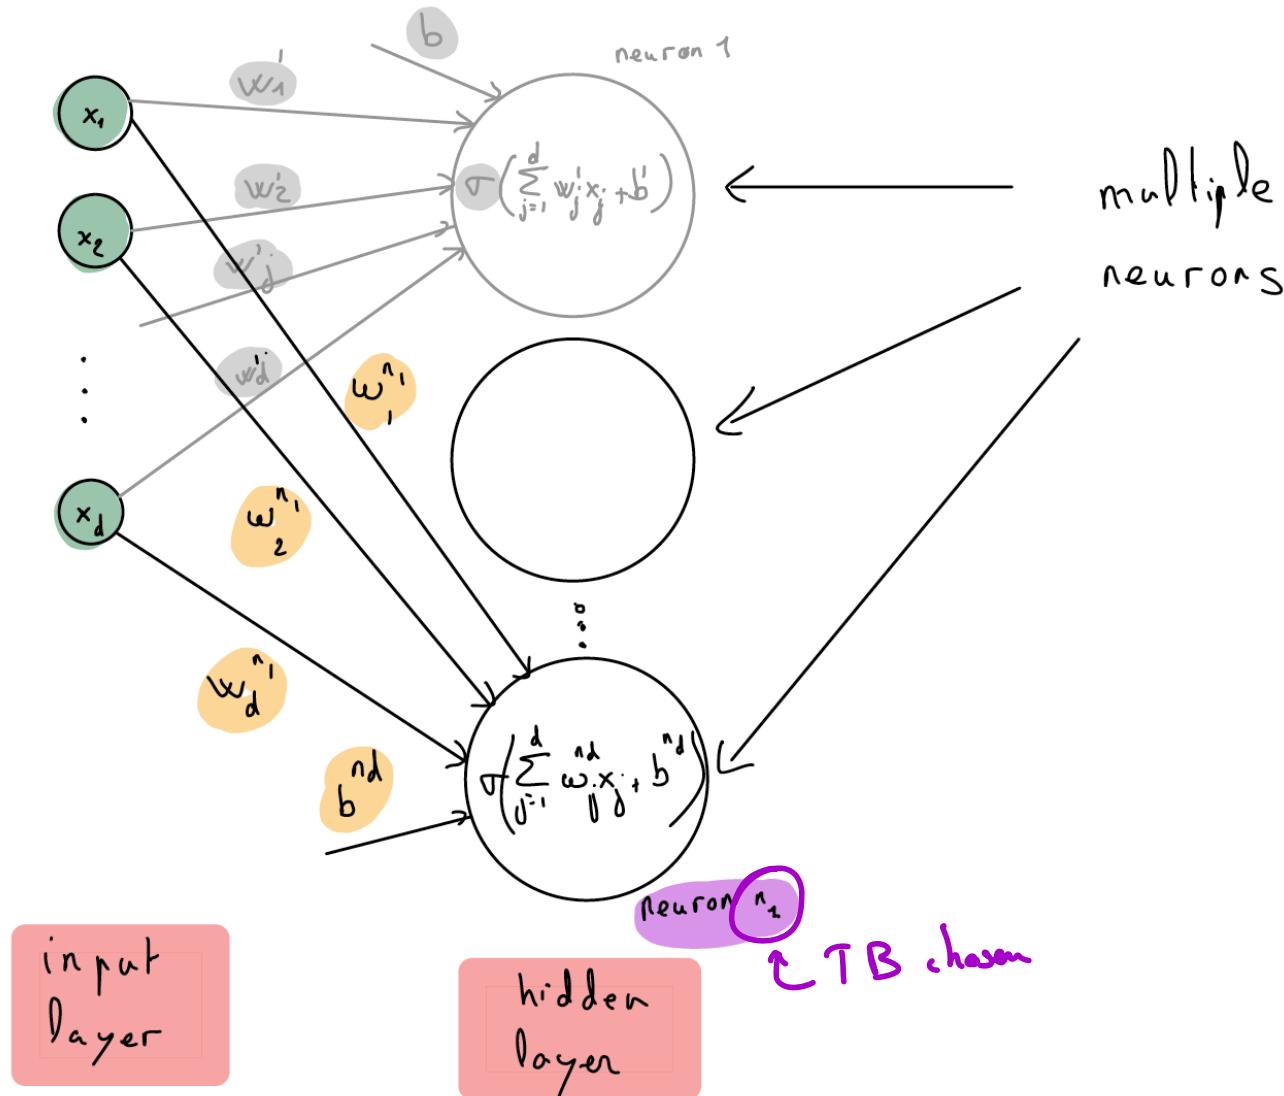
- Multiple Neurons ✓
- Multiple Layers



# From logistic regression to Multi Layer Perceptron 4

Extend to

- Multiple Neurons
- Multiple Layers

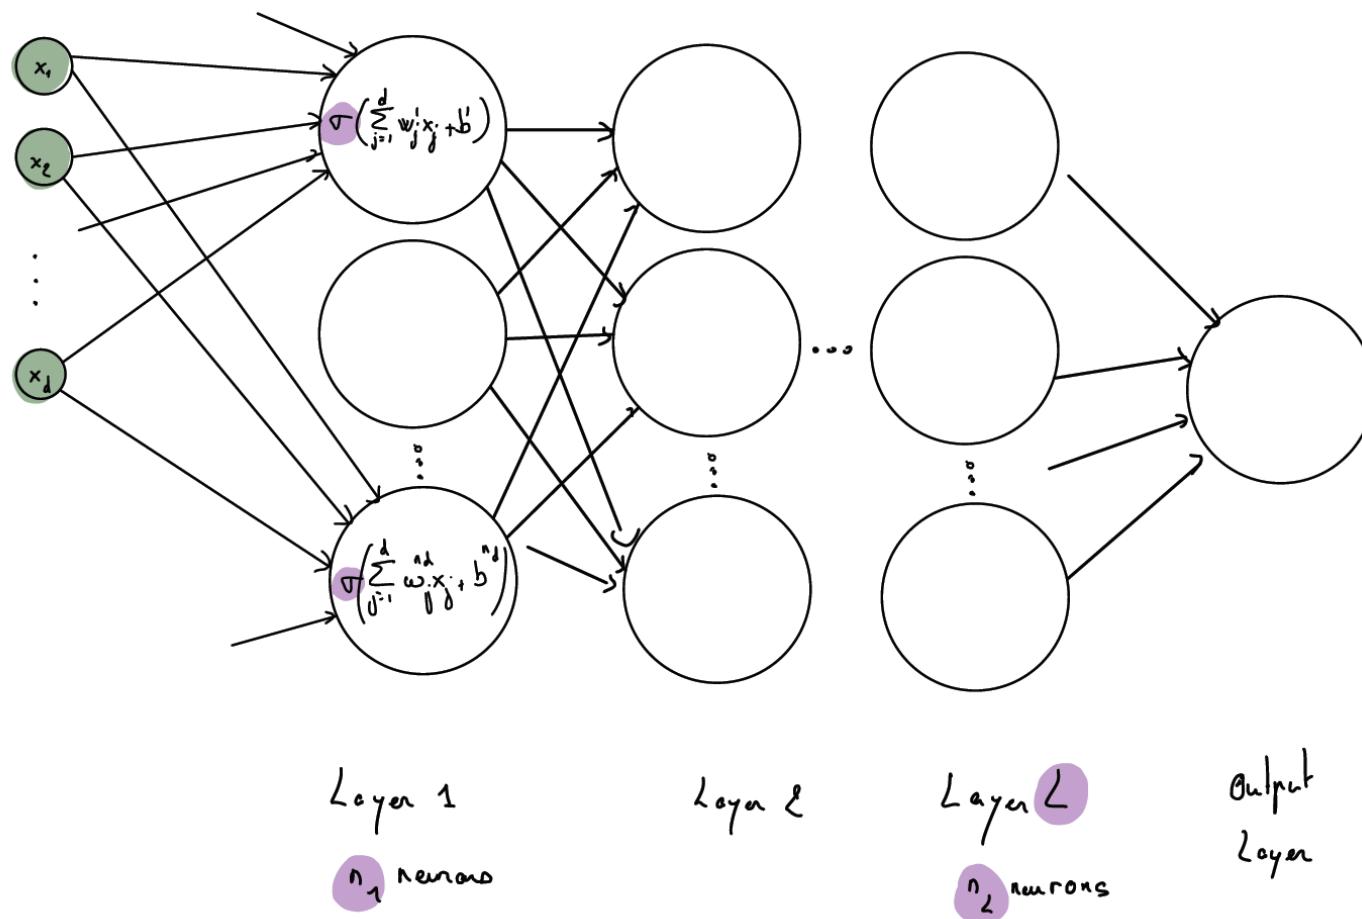


# From logistic regression to Multi Layer Perceptron 5

Extend to

- Multiple Neurons
- Multiple Layers

values after non linearity act as  
inputs for the next layer. and again . .



Again:   
learn weights  
that create features  
that can be combined  
in features (that can  
be combined in features)  
that ... can be combined  
(in output)

# From logistic regression to Multi Layer Perceptron 5

Neural network:

obj

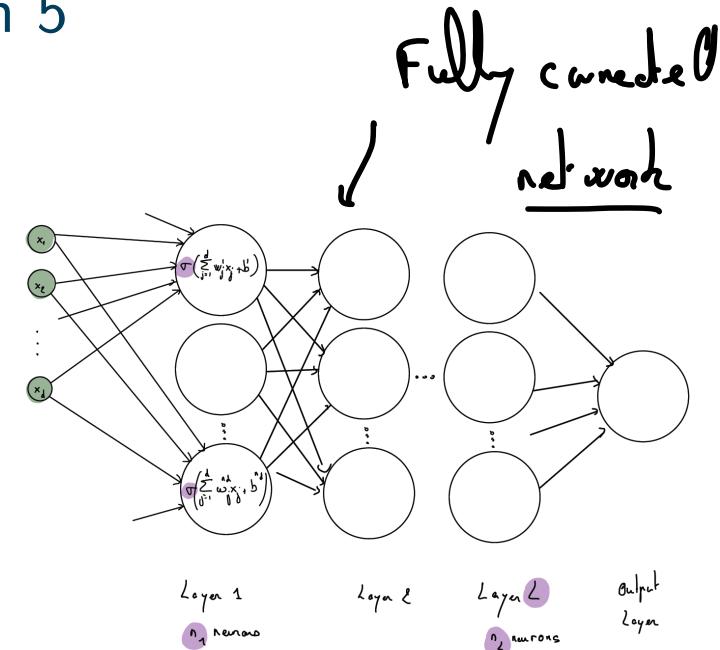
- Goal  $\min_{f \in \mathcal{F}_{MLP}} \frac{1}{n} \sum_{i=1}^n [\ell(f(X_i), Y_i)]$  ERR !

- Class of non linear functions.

$$\mathcal{F}_{MLP} = \{f(X; W, b) = \sigma(b_n + W_n \sigma(\dots(b_1 + W_1 x))),$$

$$W \in \mathbb{R}^{n \times n}, b \in \mathbb{R}^n\}$$

matrix of  
weights



Summary: NN generalize regression by looking for candidates in a much larger class of functions than linear ones.

What needs to be learned

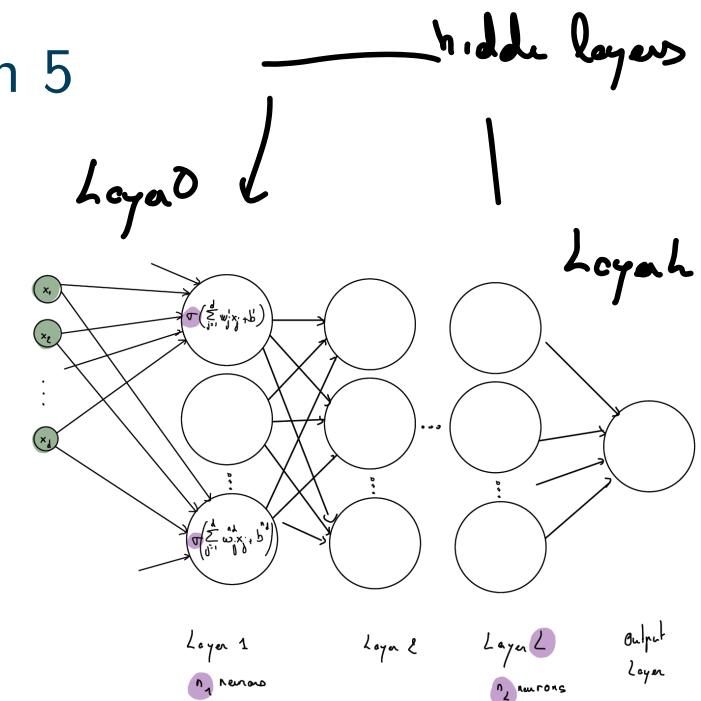
What needs to be chosen

Vocabulary

# From logistic regression to Multi Layer Perceptron 5

Neural network:

- Goal  $\min_{f \in \mathcal{F}_{\text{MLP}}} \frac{1}{n} \sum_{i=1}^n [\ell(f(X_i), Y_i)]$
  - Class of non linear functions.
- $$\mathcal{F}_{\text{MLP}} = \{f(X; W, b) = \sigma(b_n + W_n \sigma(\dots(b_1 + W_1 x))),$$
- $$W \in \mathbb{R}^{d \times n}, b \in \mathbb{R}^{n+1}\}$$



Summary: NN generalize regression by looking for candidates in a much larger class of functions than linear ones.

What needs to be learned

Parameters  $W, b$

# learnable weights on layer  
 $L$   
 $n_0 \times (n_i + 1)$  on first bld  
 $(d+1) \times n_1$

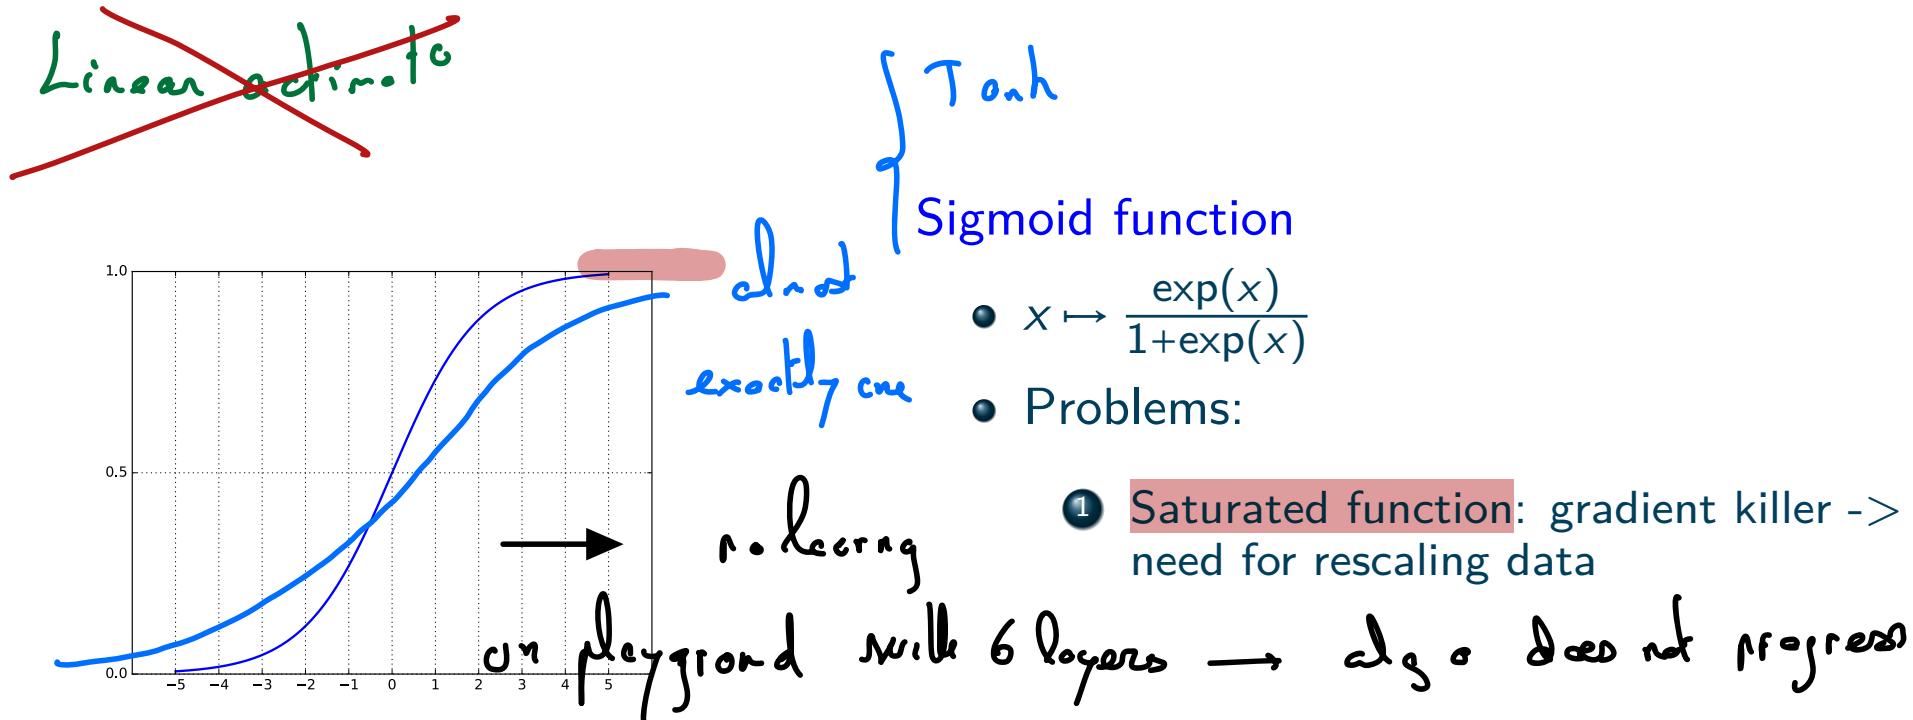
What needs to be chosen

- Activation  $\sigma$
  - Number of layers  $L$
  - Number of neurons per hidden layer
- $$n_i, i = 1, \dots, L.$$

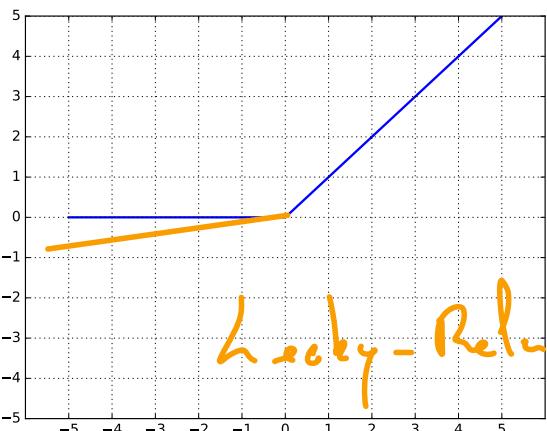
Vocabulary

- Neuron
- Activation
- Hidden Layer
- Width, Depth
- Fully connected layer

## 2 possible activations: Sigmoid and Rectified Linear Unit



### Rectified Linear Unit (ReLU)



- $x \mapsto \max(0, x)$
- Pros & cons:
  - ① Not a saturated function. Kills negative values.
  - ② Empirically, convergence is faster than sigmoid/tanh.
  - ③ Plus: biologically plausible

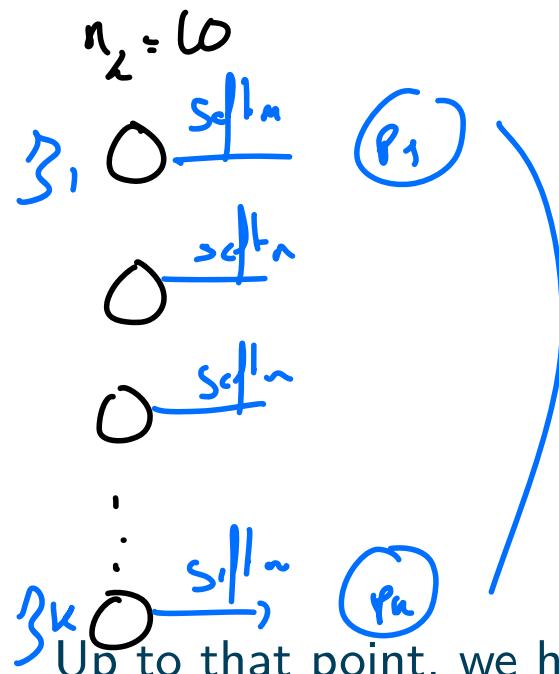
# How to deal with multi-class output: softmax activation

$$n_L = k$$

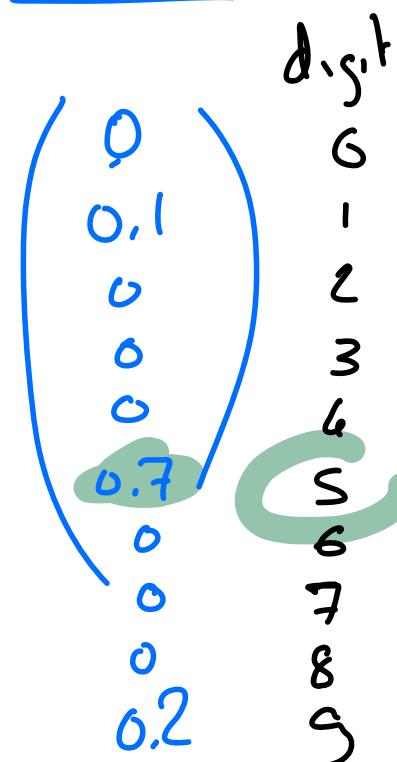
When we do multi-class classification, i.e.,  $Y_i \in 1, \dots, K$ , we output  $K$  different values:

- Each of them corresponds to the probability of belonging to the class  $j$ ,  $1 \leq j \leq K$
- To obtain probabilities (adding up to 1):
  - ① We have  $K$  neurons on the last layer
  - ② And use a Softmax activation to renormalize.

Softmax output unit, used to predict  $\{1, \dots, K\}$ :



$$\text{softmax}(z)_i = \frac{e^{z_i}}{\sum_{k=1}^K e^{z_k}}$$



Last layer:  
in regression, activation  
=  $\text{Id}$

no activation

in binary classif:

sigmoid.

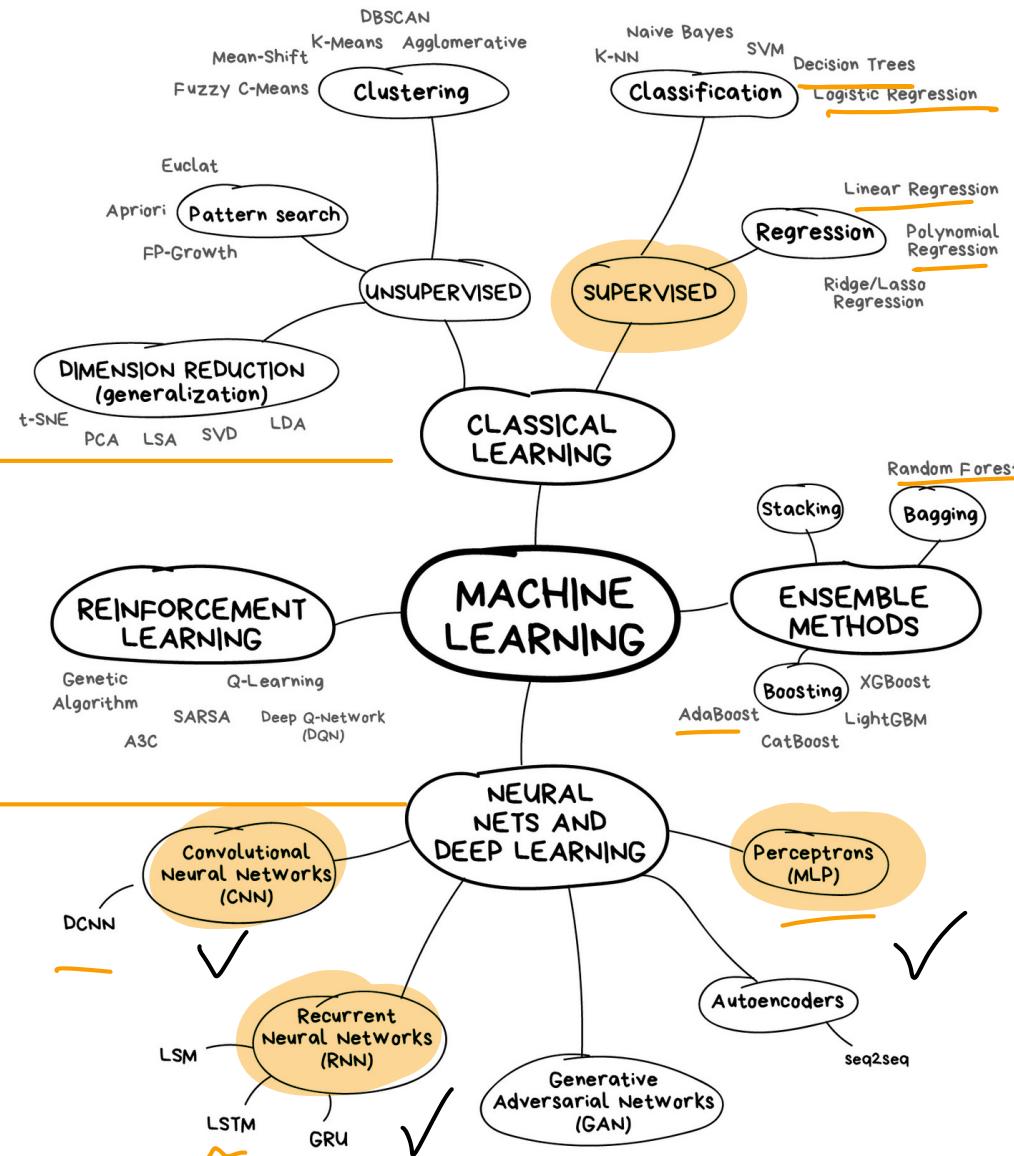
→ giving you  
a probability

Up to that point, we have seen :

- What a Neural Network was
- Why it is expected to learn better.

Next Question: how to implement it ?

# Machine Learning



# Python - Keras



## 1. Hardware:

- CPU
- GPU
- TPU

$\times 10$  times faster



## 2. Software (Python packages):

- Pytorch/Tensorflow
- → Keras: TF high level API. Ideal for applications.



TensorFlow



PyTorch



Keras

# Keras - A few lines !<sup>1</sup>

## What do we need to specify?

- \* architect ( $\#$  layer,  $\#$  neurons per h)
- \* activate
- \* algo for optimizat<sup>o</sup>, loss

## What do we need to do next?

11

predict!

```

1 # import tensorflow and keras (tf.keras not "vanilla" Keras)
2 import tensorflow as tf
3 from tensorflow import keras
4 # get data
5 (train_images, train_labels), (test_images,
6 | test_labels) = keras.datasets.mnist.load_data()
7
8 # setup model
9 model = keras.Sequential([
10 | keras.layers.Flatten(input_shape=(28, 28)),
11 | keras.layers.Dense(128, activation=tf.nn.relu),
12 | keras.layers.Dense(10, activation=tf.nn.softmax)
13 ])
14
15 model.compile(optimizer=tf.train.AdamOptimizer(),
16 | loss='sparse_categorical_crossentropy',
17 | metrics=['accuracy'])
18 | not used ↗ standard one
19 | for multiclass. ↗ 8
20 # train model in batches
21 model.fit(train_images, train_labels, epochs=5)
22
23 # evaluate
24 test_loss, test_acc = model.evaluate(test_images, test_labels)
25 print('test accuracy:', test_acc)
26
27 # make predictions
28 predictions = model.predict(test_images)

```

*Y* *X* *label*

*Tr* *Y*

enables to define a

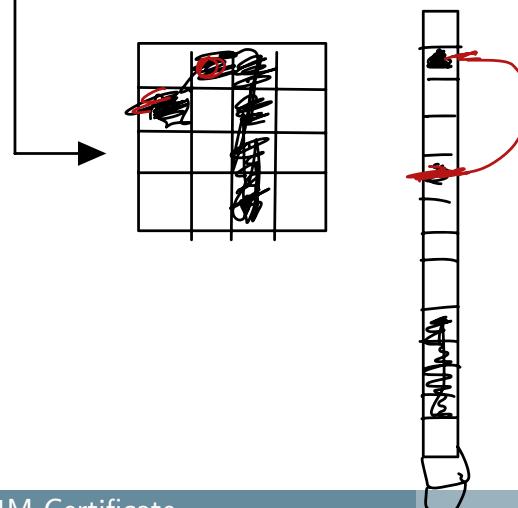
Layer 1 NN layer  
Layer 2 NN layer  
Layer 3 NN layer

per layer

30% on.

784  $n_0 = 10$

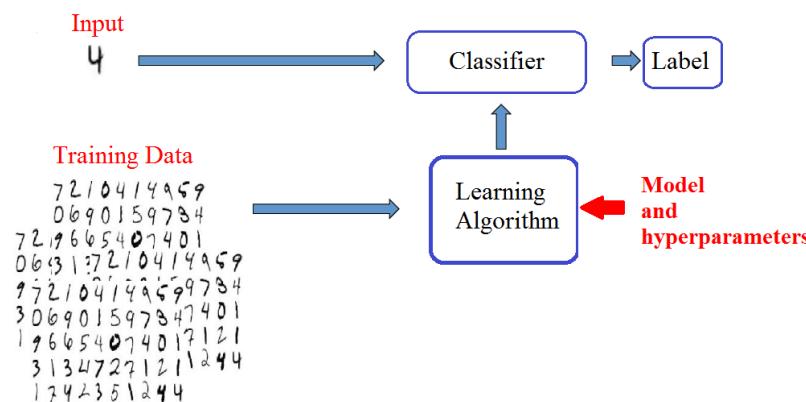
$n_1 = 128$



# Keras - A few lines !<sup>1</sup>

## What do we need to specify?

# What do we need to do next?



```

# import tensorflow and keras (tf.keras not "vanilla" Keras)
import tensorflow as tf
from tensorflow import keras
# get data
(train_images, train_labels), (test_images,
                               test_labels) = keras.datasets.mnist.load_data()

# setup model
model = keras.Sequential([
    keras.layers.Flatten(input_shape=(28, 28)),
    keras.layers.Dense(128, activation=tf.nn.relu),
    keras.layers.Dense(64, activation= tf.nn.relu),
    keras.layers.Dense(10, activation=tf.nn.softmax)
])

model.compile(optimizer=tf.optimizers.Adam(),
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# train models
model.fit(train_images, train_labels, epochs=5)

# evaluate
test_loss, test_acc = model.evaluate(test_images, test_labels)
print('test accuracy:', test_acc)

# make predictions
predictions = model.predict(test_images)

```

```

Epoch 1/5
1875/1875 [=====] - 6s 3ms/step - loss: 1.5450 - accuracy: 0.8669
Epoch 2/5
1875/1875 [=====] - 6s 3ms/step - loss: 0.2851 - accuracy: 0.9290
Epoch 3/5
1875/1875 [=====] - 5s 3ms/step - loss: 0.1913 - accuracy: 0.9474
Epoch 4/5
1875/1875 [=====] - 6s 3ms/step - loss: 0.1578 - accuracy: 0.9548
Epoch 5/5
1875/1875 [=====] - 5s 3ms/step - loss: 0.1373 - accuracy: 0.9607
313/313 [=====] - 1s 2ms/step - loss: 0.1607 - accuracy: 0.9586
test accuracy: 0.9585999846458435
313/313 [=====] - 1s 1ms/step

```

Randomness in training  
due to adj.

→ Final accuracy

784  
8      0

```

model.summary()
Model: "sequential"
Layer (type)      Output Shape       Param #
=====
flatten (Flatten) (None, 784)          0
dense (Dense)     (None, 128)         128 * 784 = 98,560
dense_1 (Dense)   (None, 64)          64 * (128+1) = 8,192
dense_2 (Dense)   (None, 10)           10 * (64+1) = 650
=====
Total params: 109,386
Trainable params: 109,386
Non-trainable params: 0

```

785\*128

100480 ✓

# Intuition

Up to that point, we have seen :

- What a Neural Network was.
- How to implement it in Python.

Next Question: why and when does it work?

# Intuition

Up to that point, we have seen :

- What a Neural Network was.
- How to implement it in Python.

Next Question: why and when does it work?

- ➊ Approximation theorem
- ➋ Optimization



You cannot vote anymore

Which of the following things need to be chosen (hyperparameters of the NN)?

1 neuron (13% | 5 people) 5 input size (46% | 18 people)

2 activation functions (95% | 37 checked) 6 number of layers (100% | 39 checked)

3 number of neurons per layer (92% | 36 checked) 7 optimization algorithm (38% | 15 checked)

4 weights (18% | 7 people) 8 biases (8% | 3 people)

**wooclap** Questions 1 / 10 Messages 🔒 100 % 🔍 Exit Tout afficher

california.ipynb

You cannot vote anymore

Which of the following things need to be learned ?

1 neuron (9% | 3 people) 5 input size (9% | 3 people)

2 activation functions (3% | 1 person) 6 number of layers (0% | 0 people)

3 number of neurons per layer (3% | 1 person) 7 optimization algorithm (9% | 3 people)

4 weights (97% | 32 checked) 8 biases (76% | 25 checked)

**wooclap** Questions 2 / 10 Messages 🔒 100 % 🔍 Exit Tout afficher

california.ipynb

You cannot vote anymore

What happens here?

1 Logistic Regression works (data is linearly separable) (58% | 15 checked)

2 I should use more neurons (0% | 0 people)

3 The network overfits (4% | 1 person)

4 I should add more layers (0% | 0 people)

5 I should change the activation function (4% | 1 person)

6 Logistic regression fails (0% | 0 people)

7 One hidden layer is sufficient (31% | 8 people)

8 I should change to a non linear activation (4% | 1 person)

1 Logistic Regression works (data is linearly separable) (13% | 3 people)

2 I should use more neurons (46% | 11 checked)

3 The network overfits (8% | 2 people)

4 I should add more layers (54% | 13 checked)

5 I should change the activation function (21% | 5 people)

6 Logistic regression fails (75% | 18 checked)

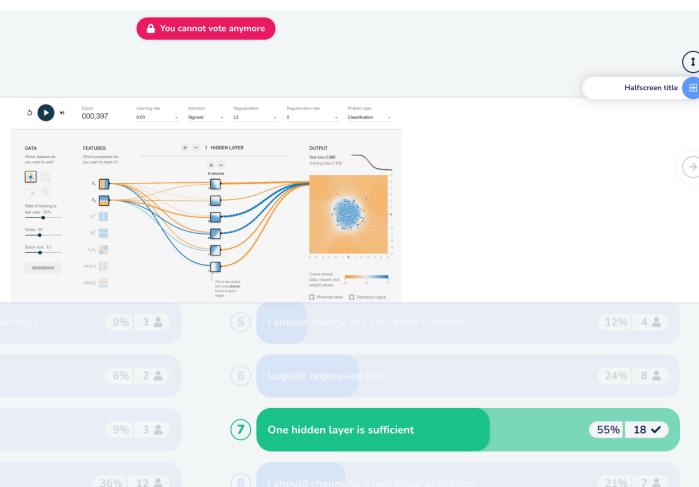
7 One hidden layer is sufficient (4% | 1 person)

8 I should change to a non linear activation (25% | 6 people)

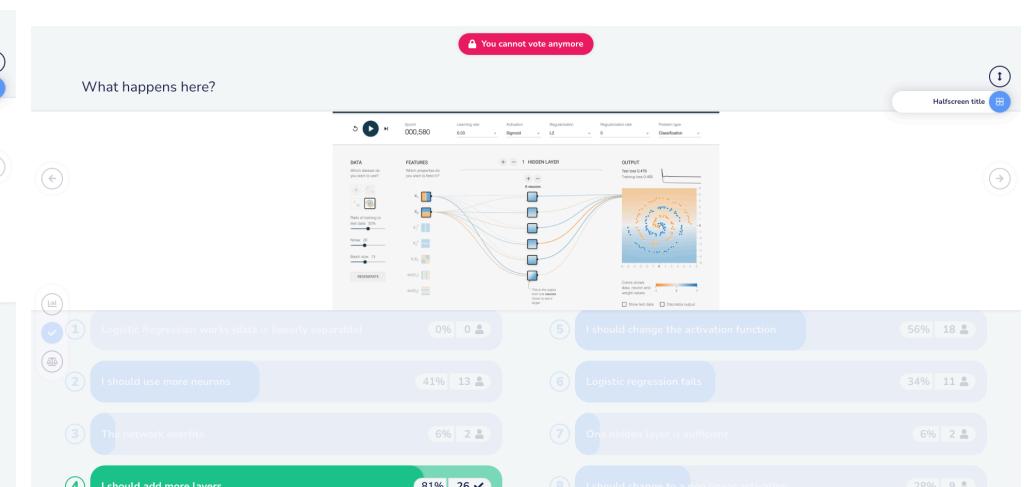
**wooclap** Questions 3 / 10 Messages 🔒 100 % 🔍 Exit Tout afficher

california.ipynb

What happens here?



What happens here?



The screenshot shows a Wooclap poll interface. At the top, a banner reads "You cannot vote anymore". Below it, the question "What happens here?" is displayed. The main content area features a neural network diagram with four hidden layers, each containing 4 neurons. The input layer has 4 neurons labeled  $x_1$ ,  $x_2$ ,  $x_3$ , and  $x_4$ . The output layer has 3 neurons labeled  $y_1$ ,  $y_2$ , and  $y_3$ . The diagram includes labels for "DATA", "FEATURES", "4 HIDDEN LAYERS", and "OUTPUT". Below the diagram, there are eight numbered options for users to vote on:

- 1 Logistic Regression works. (data is linearly separable) 0% 0
- 2 I should use more neurons. 12% 3
- 3 The network overfits. 8% 2
- 4 I should add more layers. 15% 4
- 5 I should change the activation function 85% 22 ✓
- 6 Logistic regression fails. 12% 3
- 7 One hidden layer is sufficient. 0% 0
- 8 I should change to a non-linear activation 0% 0

At the bottom, the Wooclap logo is visible along with navigation links for "Questions", "Messages", "Search", "Exit", and a progress bar indicating "73% correct" and "26 / 45".

What happens here?
You cannot vote anymore

(1)
Epoch 001.348
Learning rate 0.001
Activation ReLU
Regularization None
Dropout rate 0
Problem type Classifier

(1)
(2)
(3)
(4)

**DATA**  
SVM handles all linearly separable datasets  
For this example, we have 2 classes: Iris-setosa and Iris-versicolor.  
Number of training samples: 150  
Number of features: 4  
Number of hidden layers: 4

**FEATURES**  
SVM handles all linearly separable datasets  
For this example, we have 2 classes: Iris-setosa and Iris-versicolor.  
Number of training samples: 150  
Number of features: 4  
Number of hidden layers: 4

**OUTPUT**  
The output of the network is a probability distribution over the classes.  
Color shows data points are correctly classified.  
Color bar: 0.0 (blue) to 1.0 (orange)  
Show raw data (checkbox checked)

(1)
(2)
(3)
(4)

1 Logistic Regression works. Data is linearly separable.

0% 0

I should change the activation function.

4% 1

2 I should use more neurons.

4% 1

6 Logistic regression fails.

0% 0

3 The network overfits.

83% 19 ✓

7 One hidden layer is sufficient.

0% 0

4 I should add more layers.

0% 0

8 I should change to a non-linear activation.

9% 2

wooclap
Questions 9 / 10 +
Messages
100 %
Exit

california.ignb
83% correct
23 / 45
i

Tout afficher

The screenshot shows a Wooclap poll interface. At the top, a red banner displays the message "You cannot vote anymore". Below it, the question "What happens here?" is asked. The main content area features a neural network diagram with 6 hidden layers, input data (a handwritten digit 'G'), and output data (a digit 'G' with a color scale). On the left, a sidebar provides information about the data, features, and input. Below the diagram is a list of 8 numbered options, each with a progress bar indicating the percentage of votes:

- ① Logistic Regression works (data is linearly separable) - 0% (blue)
- ② I should change the activation function - 0% (blue)
- ③ I could use more neurons - 0% (blue)
- ④ I should change the learning rate - 0% (blue)
- ⑤ The network overfits - 0% (blue)
- ⑥ I could use a regression loss - 0% (blue)
- ⑦ It works! - 100% (green, highlighted)
- ⑧ I should add more layers - 0% (blue)

# Approximation Theorem

## Goal

$$\min_{f \in \mathcal{F}_{\text{MLP}}} \frac{1}{n} \sum_{i=1}^n [\ell(f(X_i), Y_i)].$$

As said before, a MLP can output non-linear functions... But how powerful is it?

# Approximation Theorem

Summary is

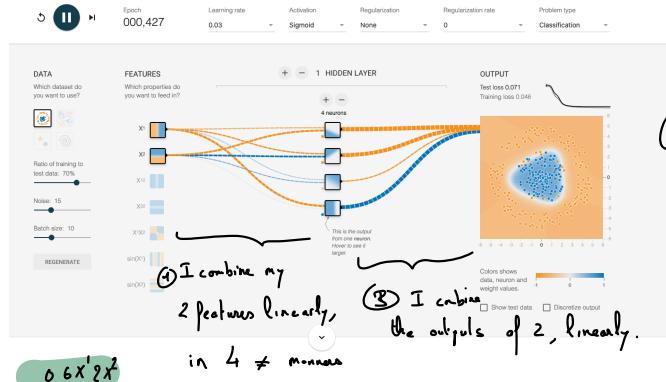


"is uni. Neural"

## Goal

$$\min_{f \in \mathcal{F}_{\text{MLP}}} \frac{1}{n} \sum_{i=1}^n [\ell(f(X_i), Y_i)].$$

As said before, a MLP can output non-linear functions... But how powerful is it?



0  
except

h neurons 1 Layer "approx-ater"

Continuous Neural Networks with one single hidden layer and any bounded and non constant activation function can approximate any function in  $L^p$ , provided a sufficient number of hidden units.

~ ~ ~ nb of neurons if  $L=1$  can be extremely large  
much smaller if  $L \gg 1$

→ Very powerful in terms of approximations.

Not very surprising: non linear function with millions of parameters!

# Optimization - Fitting the network

$$\widehat{R}_n(\omega)$$

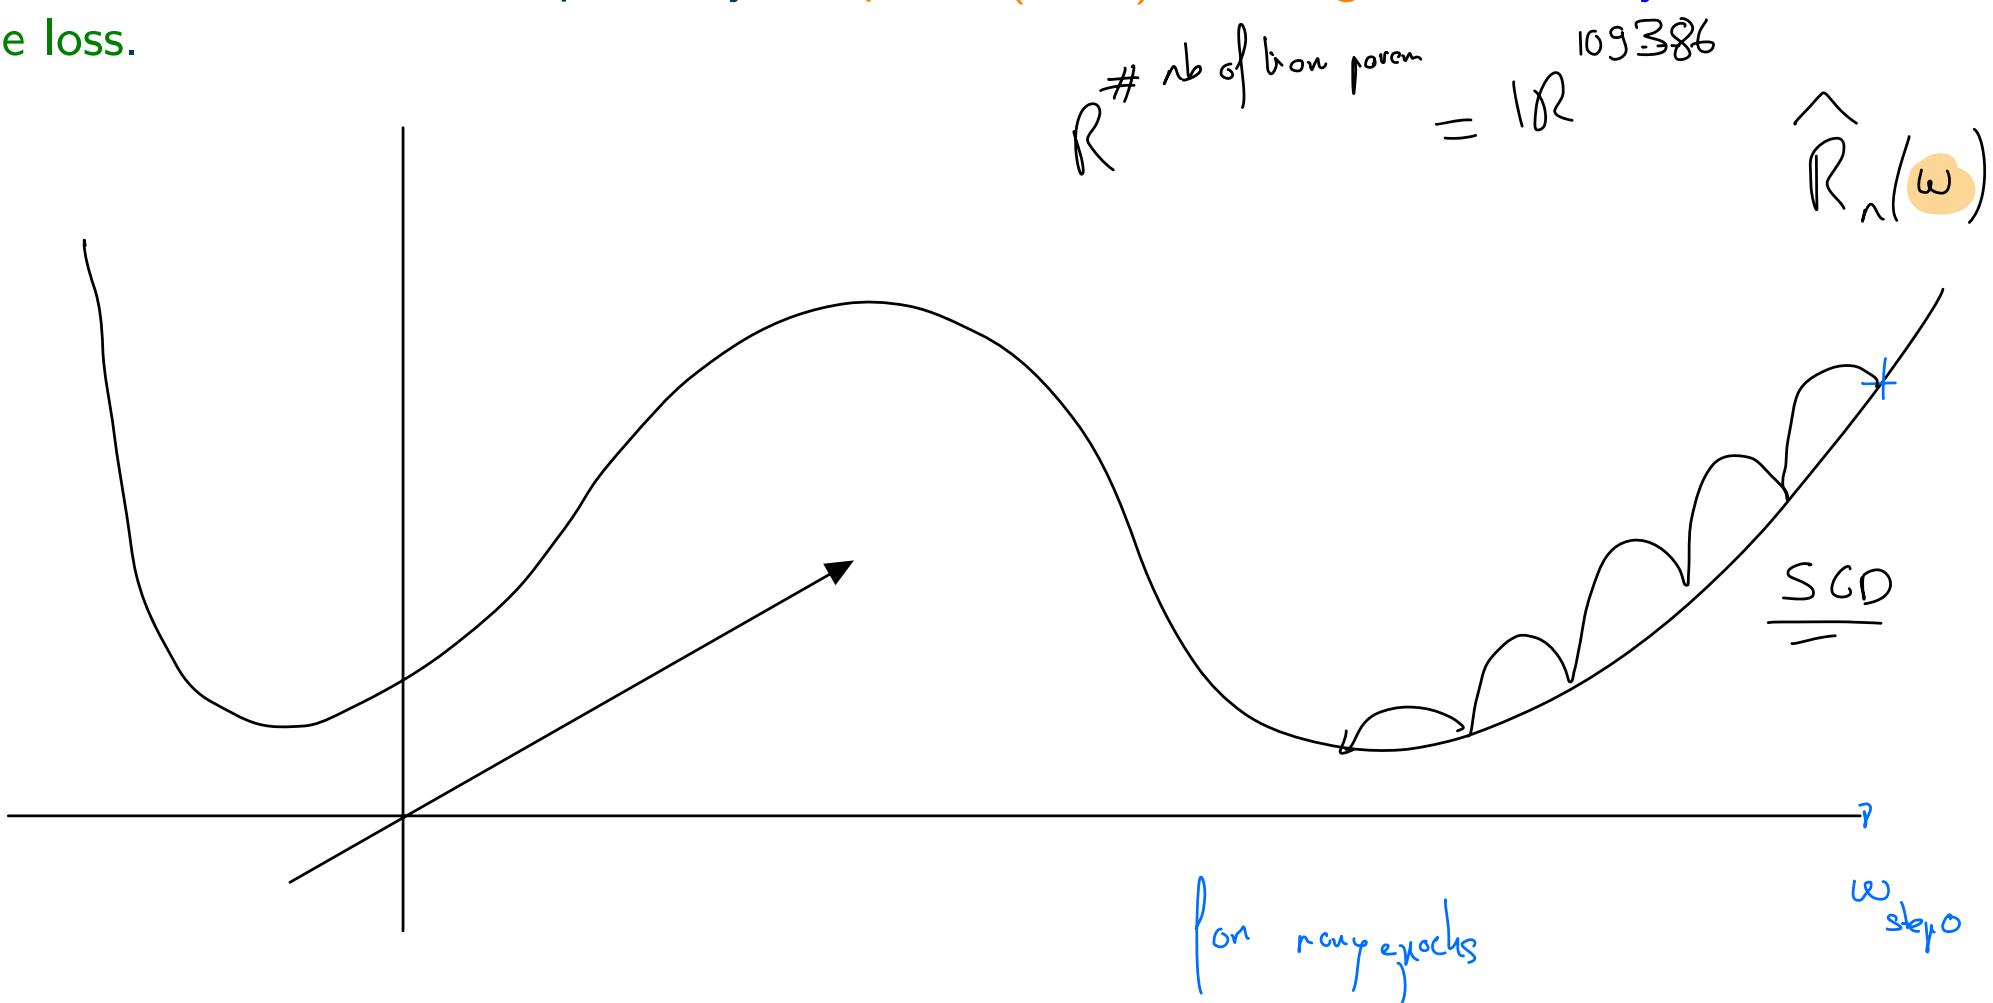
//

Goal

$$\min_{\omega \in \mathcal{F}_{MLP}} \frac{1}{n} \sum_{i=1}^n [\ell(f_\omega(X_i), Y_i)].$$

How to minimize a function?

Gradient Methods. = cheapest way to update (learn) the weights iteratively to minimize the loss.



# Optimization - Fitting the network

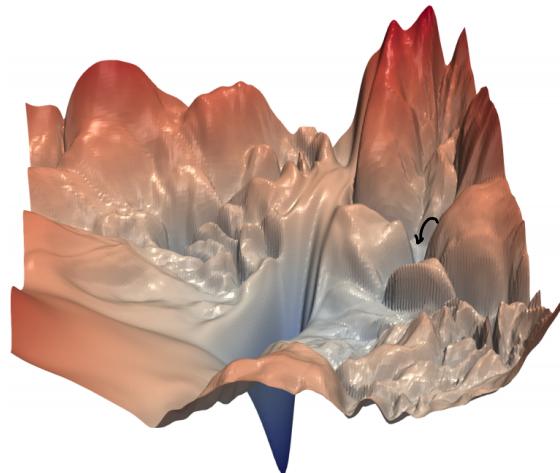
Goal

$$\min_{f \in \mathcal{F}_{\text{MLP}}} \frac{1}{n} \sum_{i=1}^n [\ell(f(X_i), Y_i)].$$

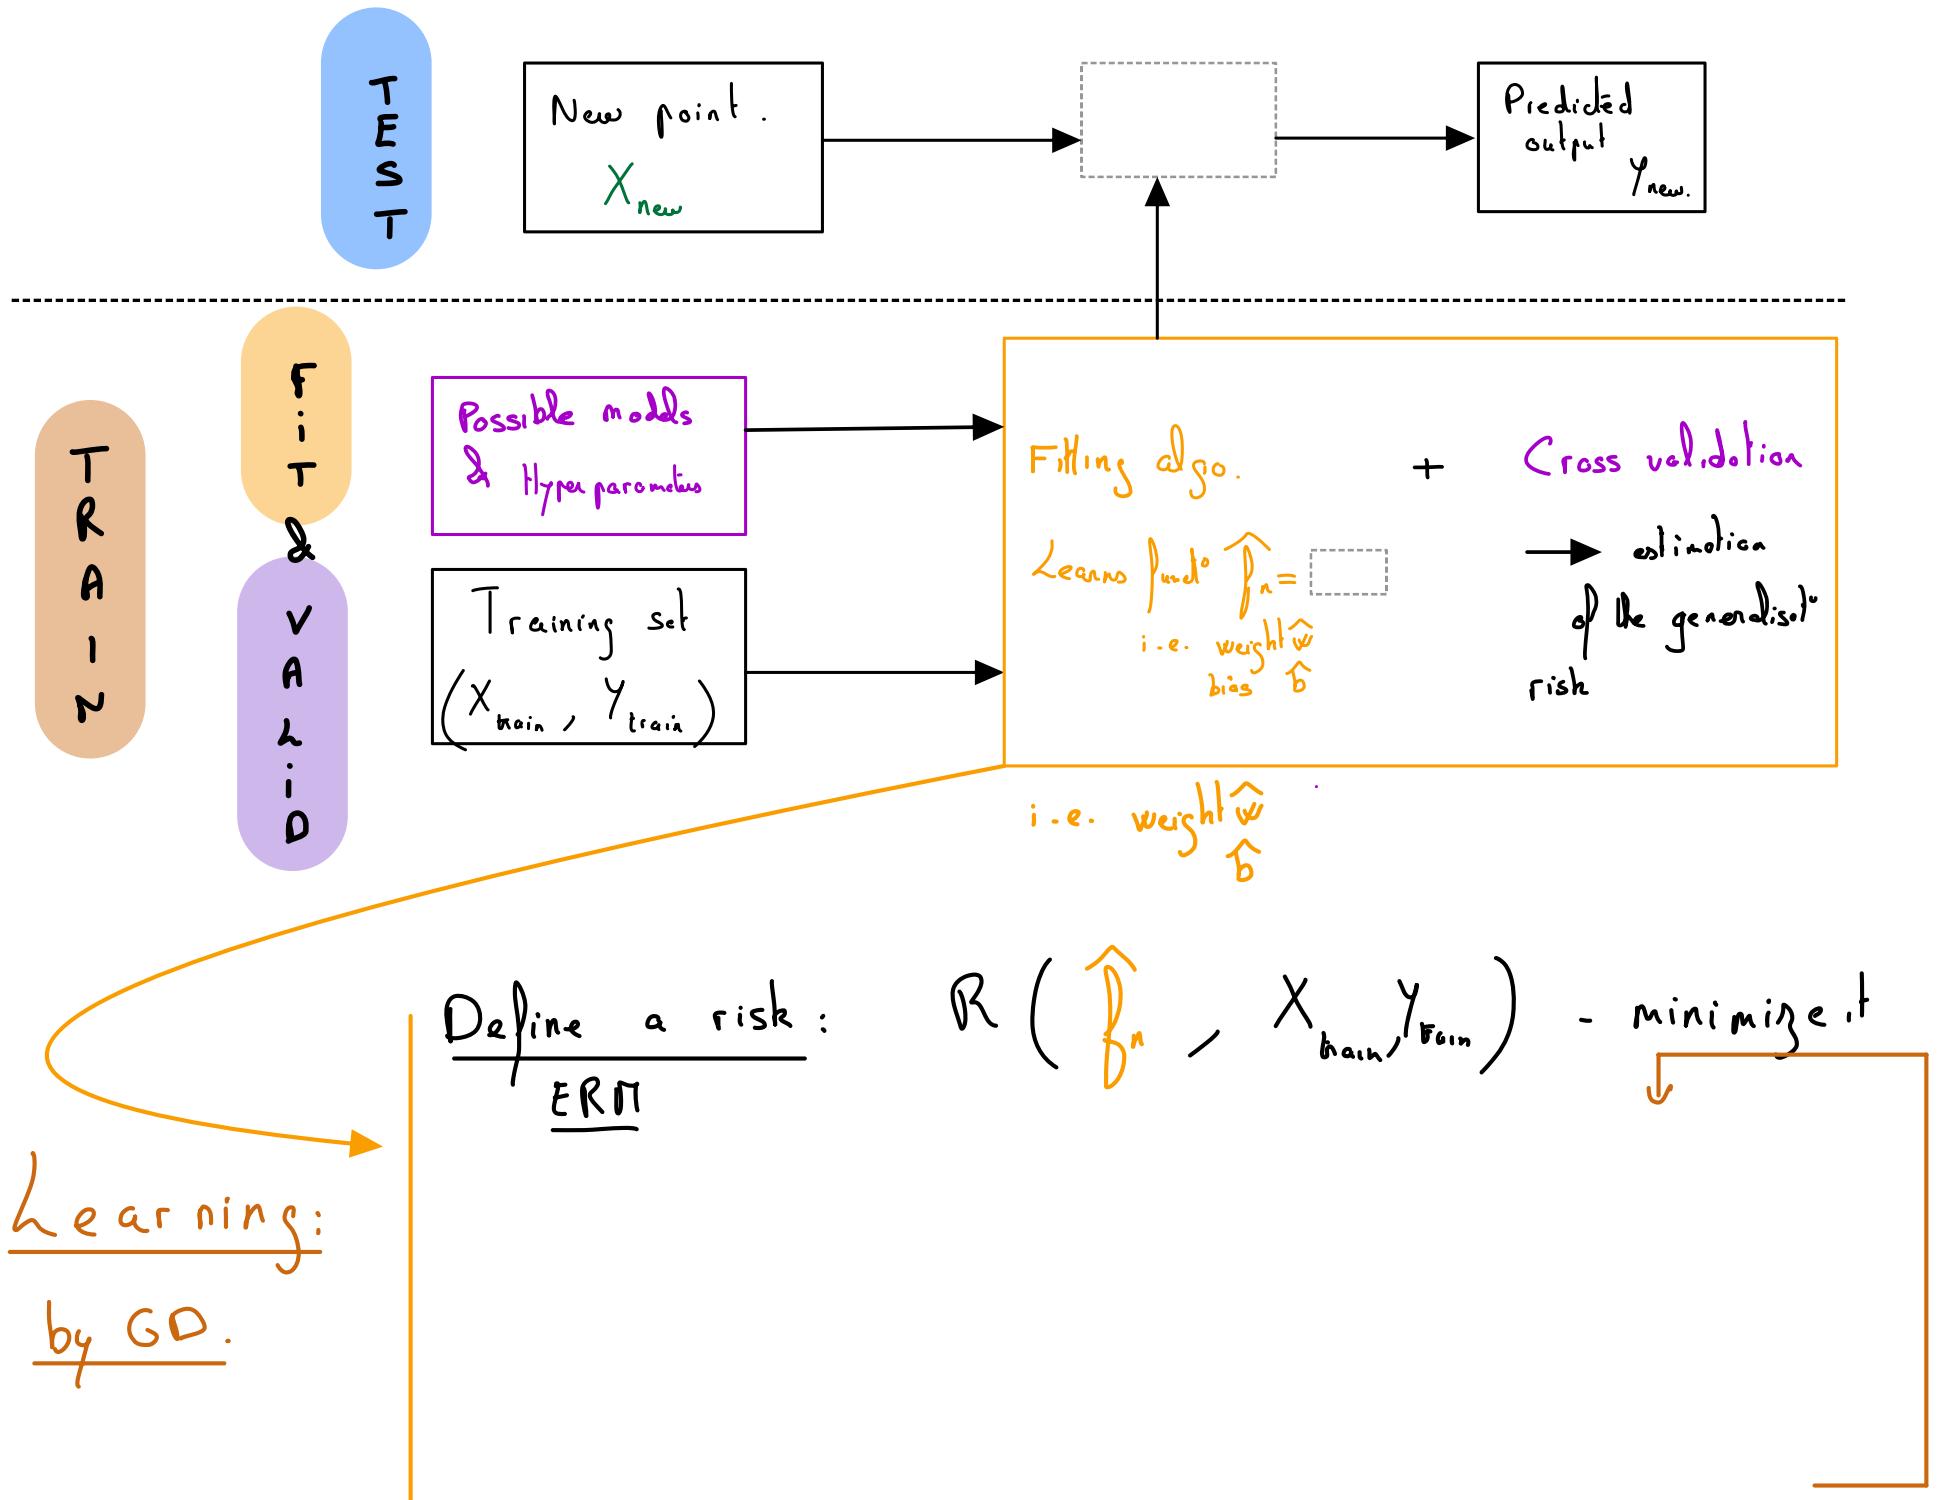
Gradient Methods. = cheapest way to update (learn) the weights iteratively to minimize the loss.

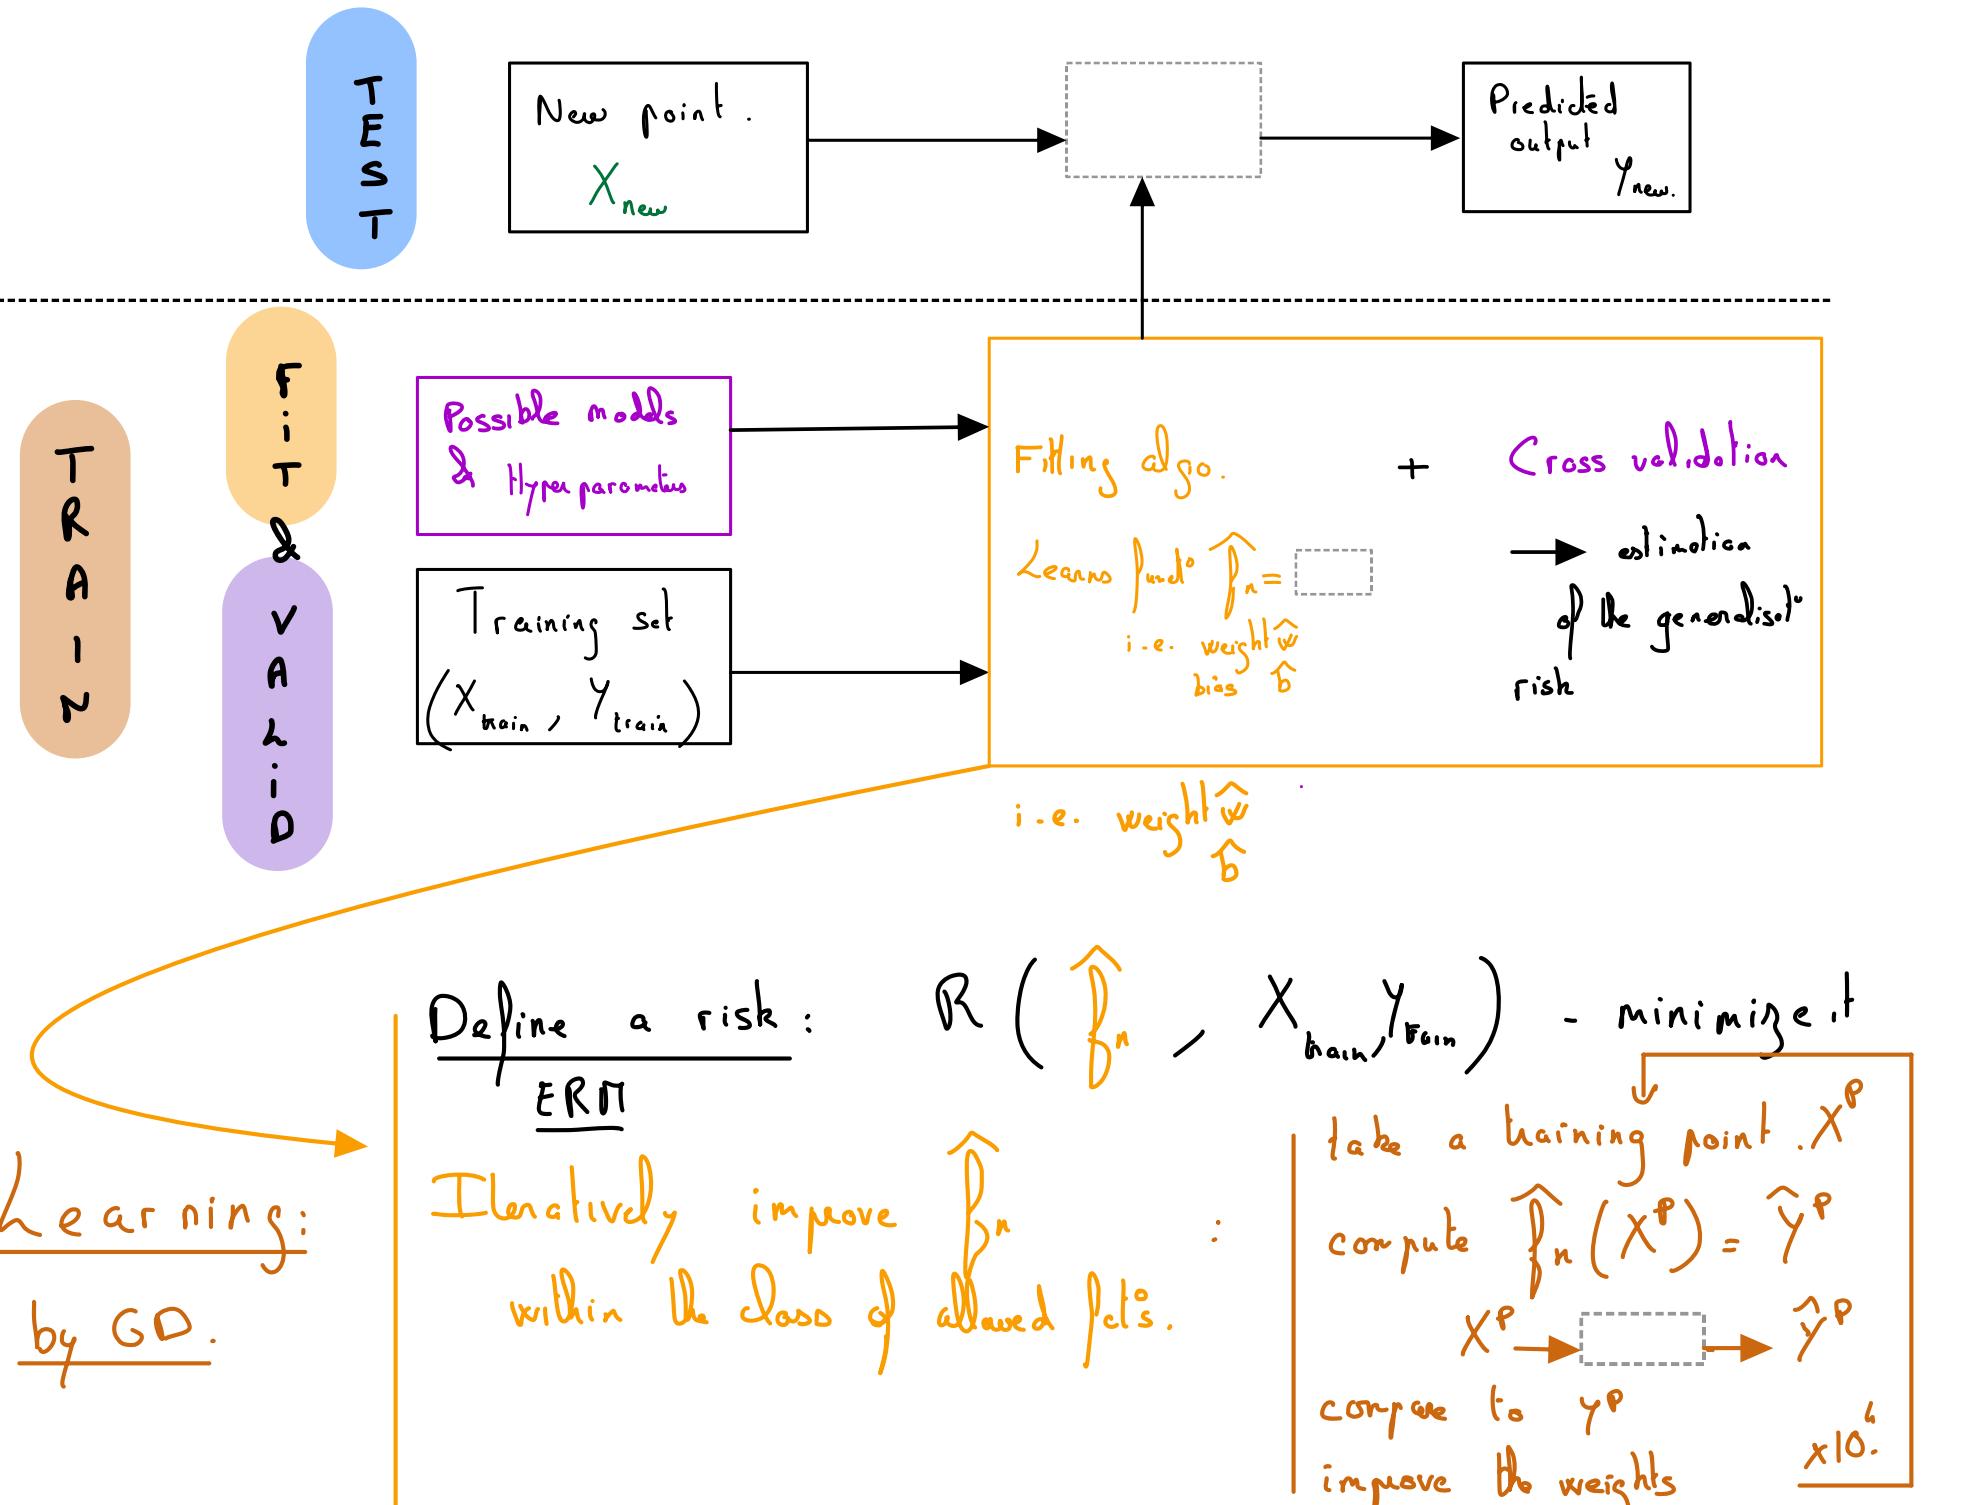
Two “miracles”:

- ① Autodifferentiation: even though the function is very complex and high dimensional, it is possible to compute gradients!
- ② Optimization seems to work ok, though the function is non convex - we do not end up in too-bad local minima. (specialized optim algos: Adam, AdaDelta, etc.)



These 2 miracles make it possible to learn a good regressor/classifier with NN and to benefit from the powerful approximation properties





# Conclusion

Neural networks :

learn features

~~feature design by hand~~

how long remember

- ① learn very complex non linear functions in high dimension ✓ THM ++
- ② can approximate nearly any function ✓
- ③ can be optimized even though highly non convex. ✓

Are thus expected to work:

- When depth or width increases. ✓
- Thus with very large datasets ✓
- Requiring a lot of computational power. ✓

⇒ Explains why they were so successful since 2010 ✓

# Conclusion

Neural networks :

- ① learn very complex non linear functions in high dimension
- ② can approximate nearly any function
- ③ can be optimized even though highly non convex.

Are thus expected to work:

- When depth or width increases.
- Thus with very large datasets
- Requiring a lot of computational power.

⇒ **Explains why they were so successful since 2010**

*What my layer wants me to say:* we haven't talked about many points !!

- Initialization
- Regularization

# Conclusion

Neural networks :

- ① learn very complex non linear functions in high dimension
- ② can approximate nearly any function
- ③ can be optimized even though highly non convex.

Are thus expected to work:

- When depth or width increases.
- Thus with very large datasets
- Requiring a lot of computational power.

⇒ **Explains why they were so successful since 2010**

*What my layer wants me to say:* we haven't talked about many points !!

- Initialization
- Regularization

Next:

- Brief history
- How to use the structure?

# A brief history of NN

- 1943: Neural networks
- 1957: Perceptron
- 1974-86: Backpropagation, RBM, RNN
- 1989-98: CNN, MNIST
- 2009: ImageNet
- 2012: AlexNet,
- 2014: GANss
- 2016: AlphaGo
- 2018: BERT

The Computational power made the major change (+ investment and creativity).

# Directions

- When does it work?

**Large or huge datasets. Structured tasks.**

↗ **Images and text.**

⚠ do not try to apply DL everywhere

Each application requires a special architecture:

- Images : Convolutional Neural Network (CNN).
- Text : Recurrent Neural Networks (RNN).

# Outline

1 Goals

2 The first Neural Network: the Perceptron

3 Convolutional Neural Networks

4 Applications

- Image classification
- Pose, action detection
- Object detection and segmentation
- Scene labeling - Semantic segmentation
- Object tracking - videos

We consider an **image classification task**: we want to recognize which object is on an image.

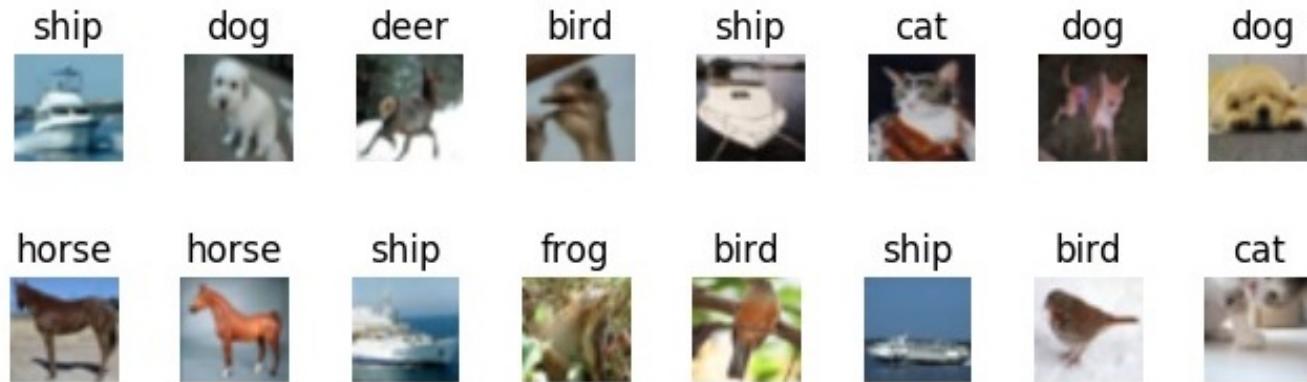


Figure: CIFAR dataset

The input is an image: about  $10^3$  to  $10^7$  pixels: these are our inputs.

Problems: The multi-layer perceptron:

- Does not take into account local information. It would work similarly on any permutation of the pixels!
- Has, just on the first layer, as many parameters as there are inputs! X

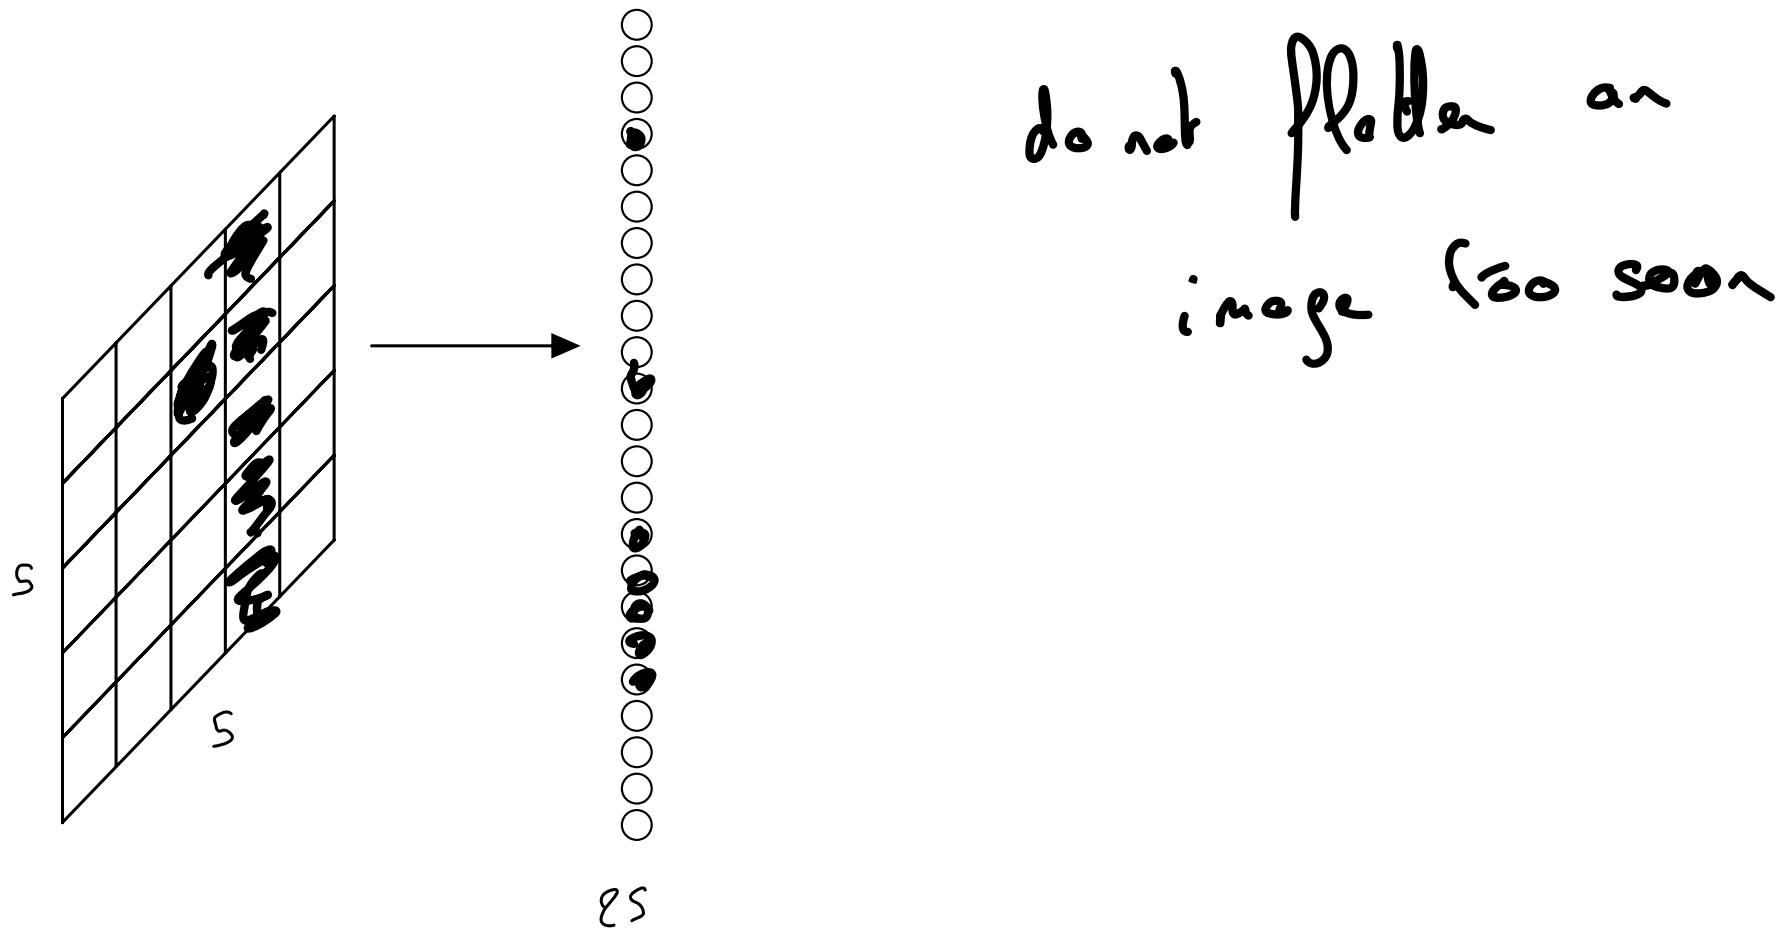
*x nb of units hidden*

## Convolutional Layer

Instead of making a linear combination of all the pixels, we consider a weighted average of a moving window of pixels, of size 3x3, or 5x5...

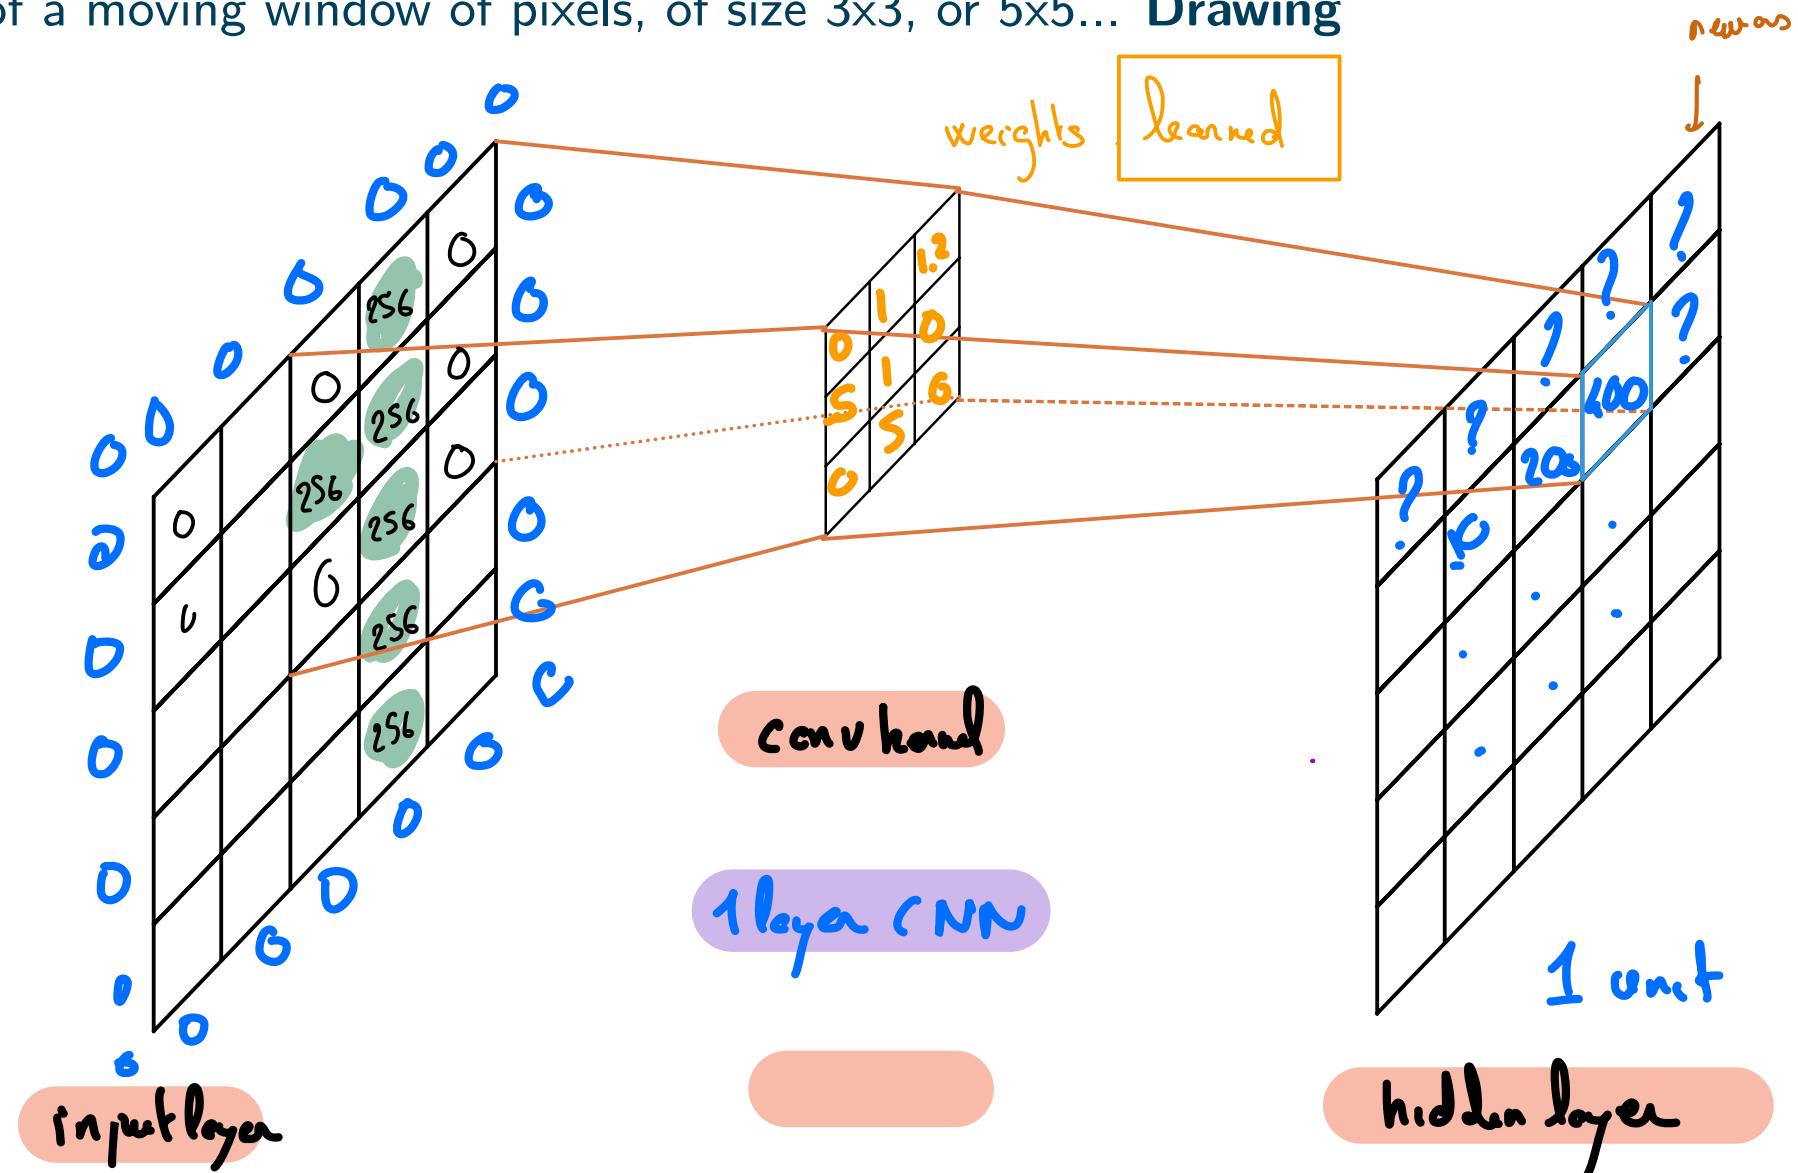
## Convolutional Layer

Instead of making a linear combination of all the pixels, we consider a weighted average of a moving window of pixels, of size  $3 \times 3$ , or  $5 \times 5$ ...



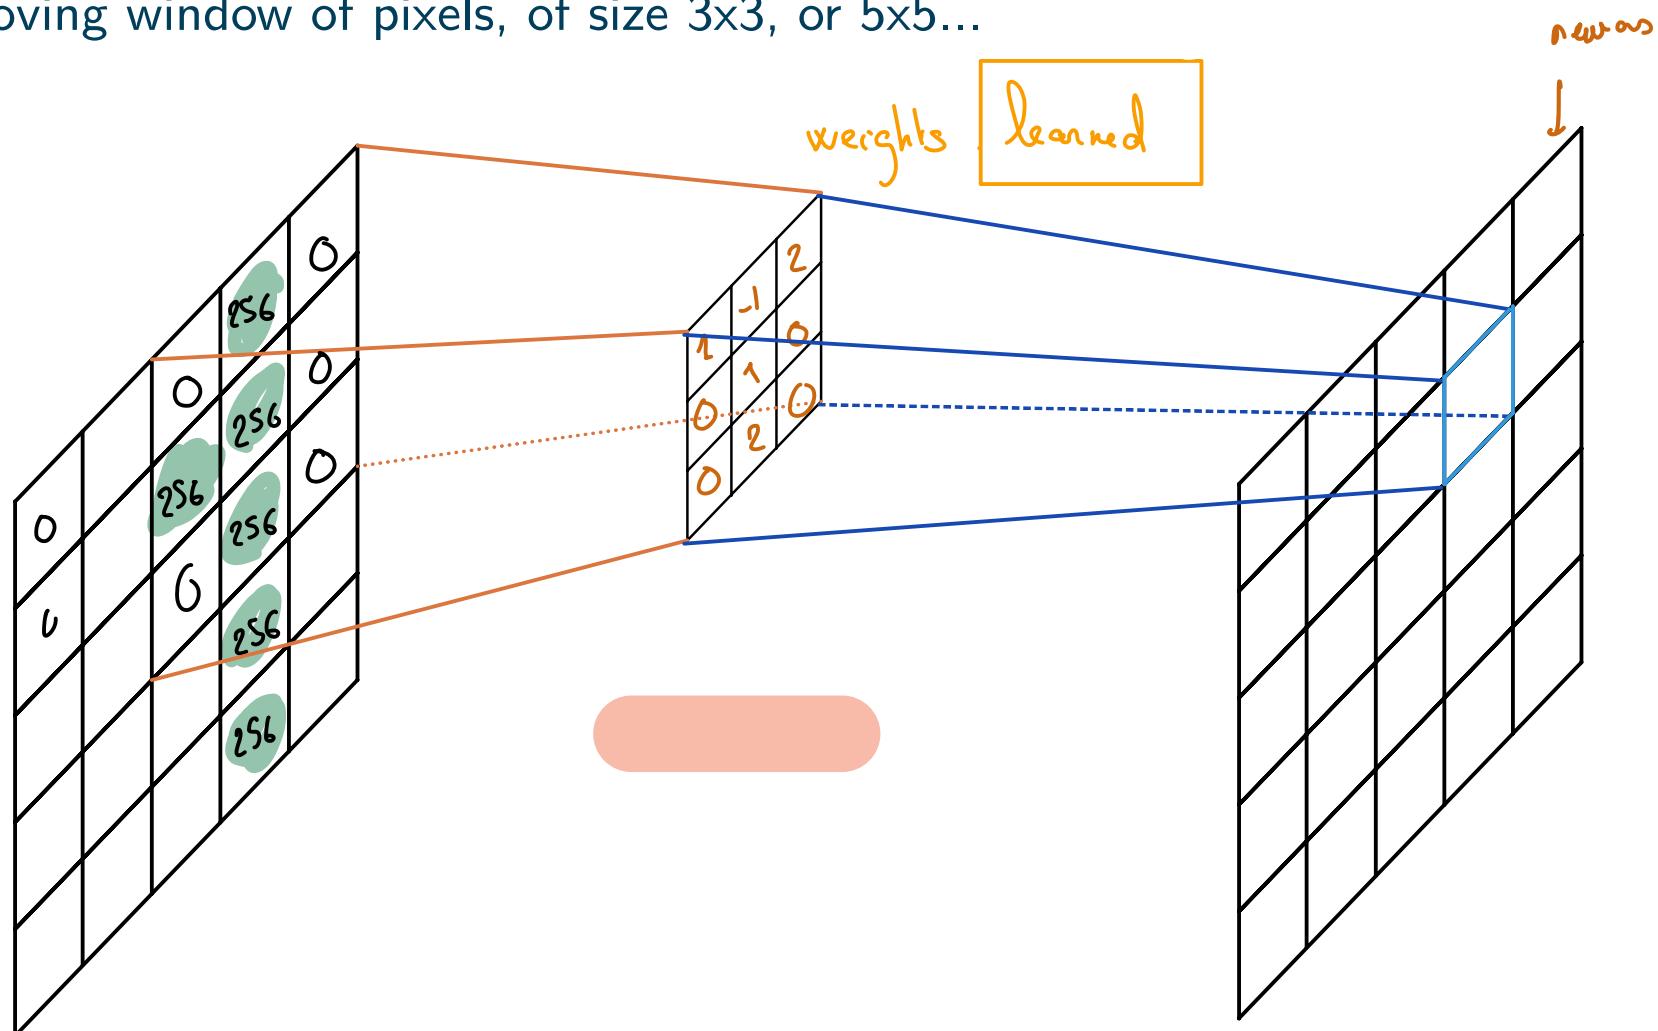
# Convolutional Layer

Instead of making a linear combination of all the pixels, we consider a weighted average of a moving window of pixels, of size 3x3, or 5x5... Drawing



# Convolutional Layer

Instead of making a linear combination of all the pixels, we consider a weighted average of a moving window of pixels, of size 3x3, or 5x5...



# Convolutional Layer

0	0	0	0	0	0	0	...
0	156	155	156	158	158	158	...
0	153	154	157	159	159	159	...
0	149	151	155	158	158	159	...
0	146	146	149	153	158	158	...
0	145	143	143	148	158	158	...
...	...	...	...	...	...	...	...

Input Channel #1 (Red)

0	0	0	0	0	0	0	...
0	167	166	167	169	169	169	...
0	164	165	168	170	170	170	...
0	160	162	166	169	170	170	...
0	156	156	159	163	168	168	...
0	155	153	153	158	168	168	...
...	...	...	...	...	...	...	...

Input Channel #2 (Green)

0	0	0	0	0	0	0	...
0	163	162	163	165	165	165	...
0	160	161	164	166	166	166	...
0	156	158	162	165	166	166	...
0	155	155	158	162	167	167	...
0	154	152	152	157	167	167	...
...	...	...	...	...	...	...	...

Input Channel #3 (Blue)

-1	-1	1
0	1	-1
0	1	1

Kernel Channel #1

1	0	0
1	-1	-1
1	0	-1

Kernel Channel #2

0	1	1
0	1	0
1	-1	1

Kernel Channel #3

308

+

-498

$$156 \times 1 + (+1) \times 155$$

$$153 \times 1 + -1 \times 154$$



$$+ 164 + 1 = -25$$

$$\text{Bias} = 1$$

-25					...
					...
					...
					...
...	...	...	...	...	...

Output

# Convolutional Layer

0	0	0	0	0	0	...
0	156	155	156	158	158	...
0	153	154	157	159	159	...
0	149	151	155	158	159	...
0	146	146	149	153	158	...
0	145	143	143	148	158	...
...	...	...	...	...	...	...

Input Channel #1 (Red)

0	0	0	0	0	0	...
0	167	166	167	169	169	...
0	164	165	168	170	170	...
0	160	162	166	169	170	...
0	156	156	159	163	168	...
0	155	153	153	158	168	...
...	...	...	...	...	...	...

Input Channel #2 (Green)

0	0	0	0	0	0	...
0	163	162	163	165	165	...
0	160	161	164	166	166	...
0	156	158	162	165	166	...
0	155	155	158	162	167	...
0	154	152	152	157	167	...
...	...	...	...	...	...	...

Input Channel #3 (Blue)

-1	-1	1
0	1	-1
0	1	1

Kernel Channel #1



308

+

1	0	0
1	-1	-1
1	0	-1

Kernel Channel #2



-498

0	1	1
0	1	0
1	-1	1

Kernel Channel #3



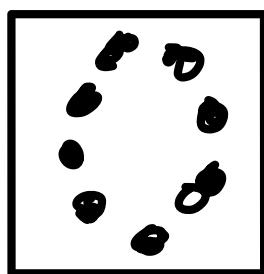
164

$$+ 1 = -25$$

Bias = 1

Output

-25					...
					...
					...
					...
...	...	...	...	...	...



learning  
geometric  
shapes



love image.

# Convolutional Layer

0	0	0	0	0	0	...
0	156	155	156	158	158	...
0	153	154	157	159	159	...
0	149	151	155	158	159	...
0	146	146	149	153	158	...
0	145	143	143	148	158	...
...	...	...	...	...	...	...

Input Channel #1 (Red)

0	0	0	0	0	0	...
0	167	166	167	169	169	...
0	164	165	168	170	170	...
0	160	162	166	169	170	...
0	156	156	159	163	168	...
0	155	153	153	158	168	...
...	...	...	...	...	...	...

Input Channel #2 (Green)

0	0	0	0	0	0	...
0	163	162	163	165	165	...
0	160	161	164	166	166	...
0	156	158	162	165	166	...
0	155	155	158	162	167	...
0	154	152	152	157	167	...
...	...	...	...	...	...	...

Input Channel #3 (Blue)

-1	-1	1
0	1	-1
0	1	1

Kernel Channel #1



310

+

1	0	0
1	-1	-1
1	0	-1

Kernel Channel #2



-170

0	1	1
0	1	0
1	-1	1

Kernel Channel #3



+

325

+ 1 = 466

Bias = 1

Output

-25	466				...
					...
					...
					...
...	...	...	...	...	...

# Convolutional Layer

0	0	0	0	0	0	0	...
0	156	155	156	158	158	158	...
0	153	154	157	159	159	159	...
0	149	151	155	158	159	159	...
0	146	146	149	153	158	158	...
0	145	143	143	148	158	158	...
...	...	...	...	...	...	...	...

Input Channel #1 (Red)

0	0	0	0	0	0	0	...
0	167	166	167	169	169	169	...
0	164	165	168	170	170	170	...
0	160	162	166	169	170	170	...
0	156	156	159	163	168	168	...
0	155	153	153	158	168	168	...
...	...	...	...	...	...	...	...

Input Channel #2 (Green)

0	0	0	0	0	0	0	...
0	163	162	163	165	165	165	...
0	160	161	164	166	166	166	...
0	156	158	162	165	166	166	...
0	155	155	158	162	167	167	...
0	154	152	152	157	167	167	...
...	...	...	...	...	...	...	...

Input Channel #3 (Blue)

-1	-1	1
0	1	-1
0	1	1

Kernel Channel #1



314

+

1	0	0
1	-1	-1
1	0	-1

Kernel Channel #2



-175

0	1	1
0	1	0
1	-1	1

Kernel Channel #3



+ 326 + 1 = 466

Bias = 1

Output

-25	466	466	...
...	...	...	...
...	...	...	...
...	...	...	...
...	...	...	...

# Convolutional Layer

0	0	0	0	0	0	0	...
0	156	155	156	158	158	158	...
0	153	154	157	159	159	159	...
0	149	151	155	158	159	159	...
0	146	146	149	153	158	158	...
0	145	143	143	148	158	158	...
...	...	...	...	...	...	...	...

Input Channel #1 (Red)

0	0	0	0	0	0	0	...
0	167	166	167	169	169	169	...
0	164	165	168	170	170	170	...
0	160	162	166	169	170	170	...
0	156	156	159	163	168	168	...
0	155	153	153	158	168	168	...
...	...	...	...	...	...	...	...

Input Channel #2 (Green)

0	0	0	0	0	0	0	...
0	163	162	163	165	165	165	...
0	160	161	164	166	166	166	...
0	156	158	162	165	166	166	...
0	155	155	158	162	167	167	...
0	154	152	152	157	167	167	...
...	...	...	...	...	...	...	...

Input Channel #3 (Blue)

-1	-1	1
0	1	-1
0	1	1

Kernel Channel #1

1	0	0
1	-1	-1
1	0	-1

Kernel Channel #2

0	1	1
0	1	0
1	-1	1

Kernel Channel #3

↓  
318

+

↓  
-173

+

329

+ 1 = 475

↑  
Bias = 1

Output

-25	466	466	475	...
...	...	...	...	...
...	...	...	...	...
...	...	...	...	...
...	...	...	...	...

# Convolutional Layer

0	0	0	0	0	0	...
0	156	155	156	158	158	...
0	153	154	157	159	159	...
0	149	151	155	158	159	...
0	146	146	149	153	158	...
0	145	143	143	148	158	...
...	...	...	...	...	...	...

Input Channel #1 (Red)

0	0	0	0	0	0	...
0	167	166	167	169	169	...
0	164	165	168	170	170	...
0	160	162	166	169	170	...
0	156	156	159	163	168	...
0	155	153	153	158	168	...
...	...	...	...	...	...	...

Input Channel #2 (Green)

0	0	0	0	0	0	...
0	163	162	163	165	165	...
0	160	161	164	166	166	...
0	156	158	162	165	166	...
0	155	155	158	162	167	...
0	154	152	152	157	167	...
...	...	...	...	...	...	...

Input Channel #3 (Blue)

-1	-1	1
0	1	-1
0	1	1

Kernel Channel #1



298

+

1	0	0
1	-1	-1
1	0	-1

Kernel Channel #2



-491

0	1	1
0	1	0
1	-1	1

Kernel Channel #3



487

+

+ 1 = 295

Bias = 1  
↑

-25	466	466	475	...
295				...
				...
				...
...	...	...	...	...

Output

# Convolutional Layer

0	0	0	0	0	0	...
0	156	155	156	158	158	...
0	153	154	157	159	159	...
0	149	151	155	158	159	...
0	146	146	149	153	158	...
0	145	143	143	148	158	...
...	...	...	...	...	...	...

Input Channel #1 (Red)

0	0	0	0	0	0	...
0	167	166	167	169	169	...
0	164	165	168	170	170	...
0	160	162	166	169	170	...
0	156	156	159	163	168	...
0	155	153	153	158	168	...
...	...	...	...	...	...	...

Input Channel #2 (Green)

0	0	0	0	0	0	...
0	163	162	163	165	165	...
0	160	161	164	166	166	...
0	156	158	162	165	166	...
0	155	155	158	162	167	...
0	154	152	152	157	167	...
...	...	...	...	...	...	...

Input Channel #3 (Blue)

-1	-1	1
0	1	-1
0	1	1

Kernel Channel #1



148

+

1	0	0
1	-1	-1
1	0	-1

Kernel Channel #2



-8

0	1	1
0	1	0
1	-1	1

Kernel Channel #3



646

+

+ 1 = 787

Bias = 1  
↑

-25	466	466	475	...
295	787			...
				...
				...
...	...	...	...	...

Output

# Summary

CNN:

- allow to take spacial information into account
- reduce the number of parameters per layer. We can have deeper networks!

# Summary

## CNN:

- allow to take spacial information into account
- reduce the number of parameters per layer. We can have deeper networks!

## LeNet 1998:

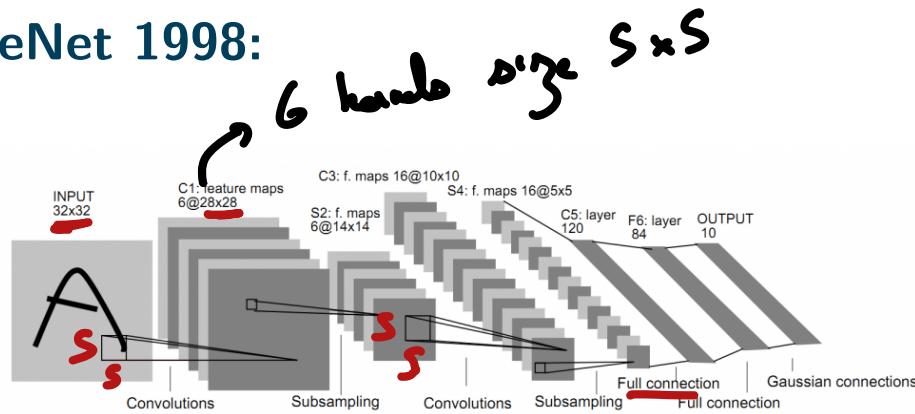
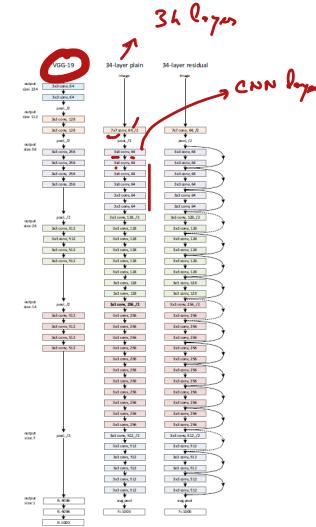


Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

8 layers, 60k parameters.

## Resnet 2016



~ 40 Layers, 140M parameters

# Outline

1 Goals

2 The first Neural Network: the Perceptron

3 Convolutional Neural Networks

4 Applications

- Image classification
- Pose, action detection
- Object detection and segmentation
- Scene labeling - Semantic segmentation
- Object tracking - videos

# Outline

1 Goals

2 The first Neural Network: the Perceptron

3 Convolutional Neural Networks

4 Applications

- **Image classification**
- Pose, action detection
- Object detection and segmentation
- Scene labeling - Semantic segmentation
- Object tracking - videos

# Image classification - CNN Tree

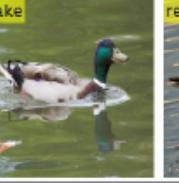
Category	Confusion Set						
tench							
indigo bunting							
red-breasted merganser							
echidna							
shopping basket							

Figure: Confusion set outputs by AlexNet softmax prediction on validation set of ILSVRC 2015.

# Image classification - CNN Tree

Category	Example Validation Images					
barracouta	 coho	 coho	 tench	 tiger shark	 great white shark	 coho
church	 church	 castle	 church	 vault	 castle	 vault
spaghetti squash	 spaghetti squash	 spaghetti squash	 lemon	 spaghetti squash	 butternut squash	 spaghetti squash
espresso	 espresso	 coffee mug	 plate	 espresso	 bakery	 espresso maker
trolleybus	 trolleybus	 trolleybus	 trolleybus	 passenger car	 passenger car	 trolleybus

Figure: Top label is given by basic AlexNet CNN while bottom one is given by CNNTree (green color corresponds to a correct prediction)

# Outline

1 Goals

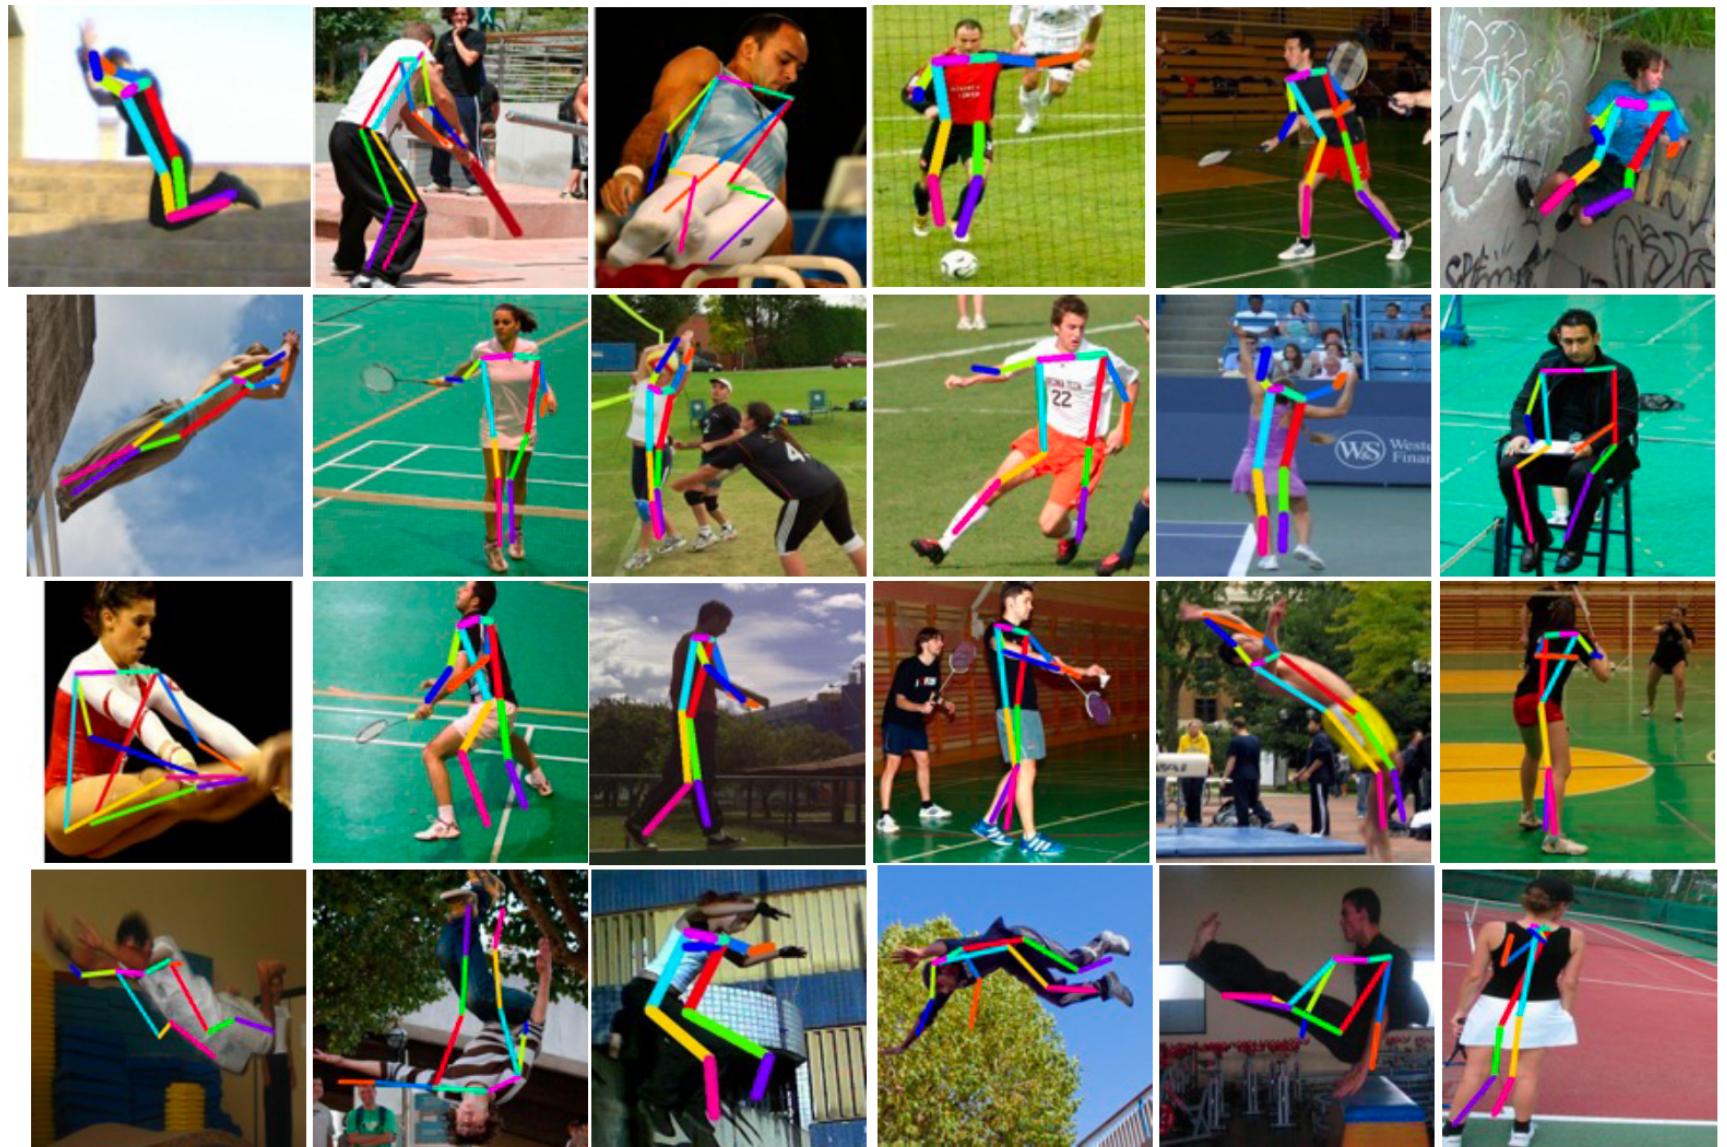
2 The first Neural Network: the Perceptron

3 Convolutional Neural Networks

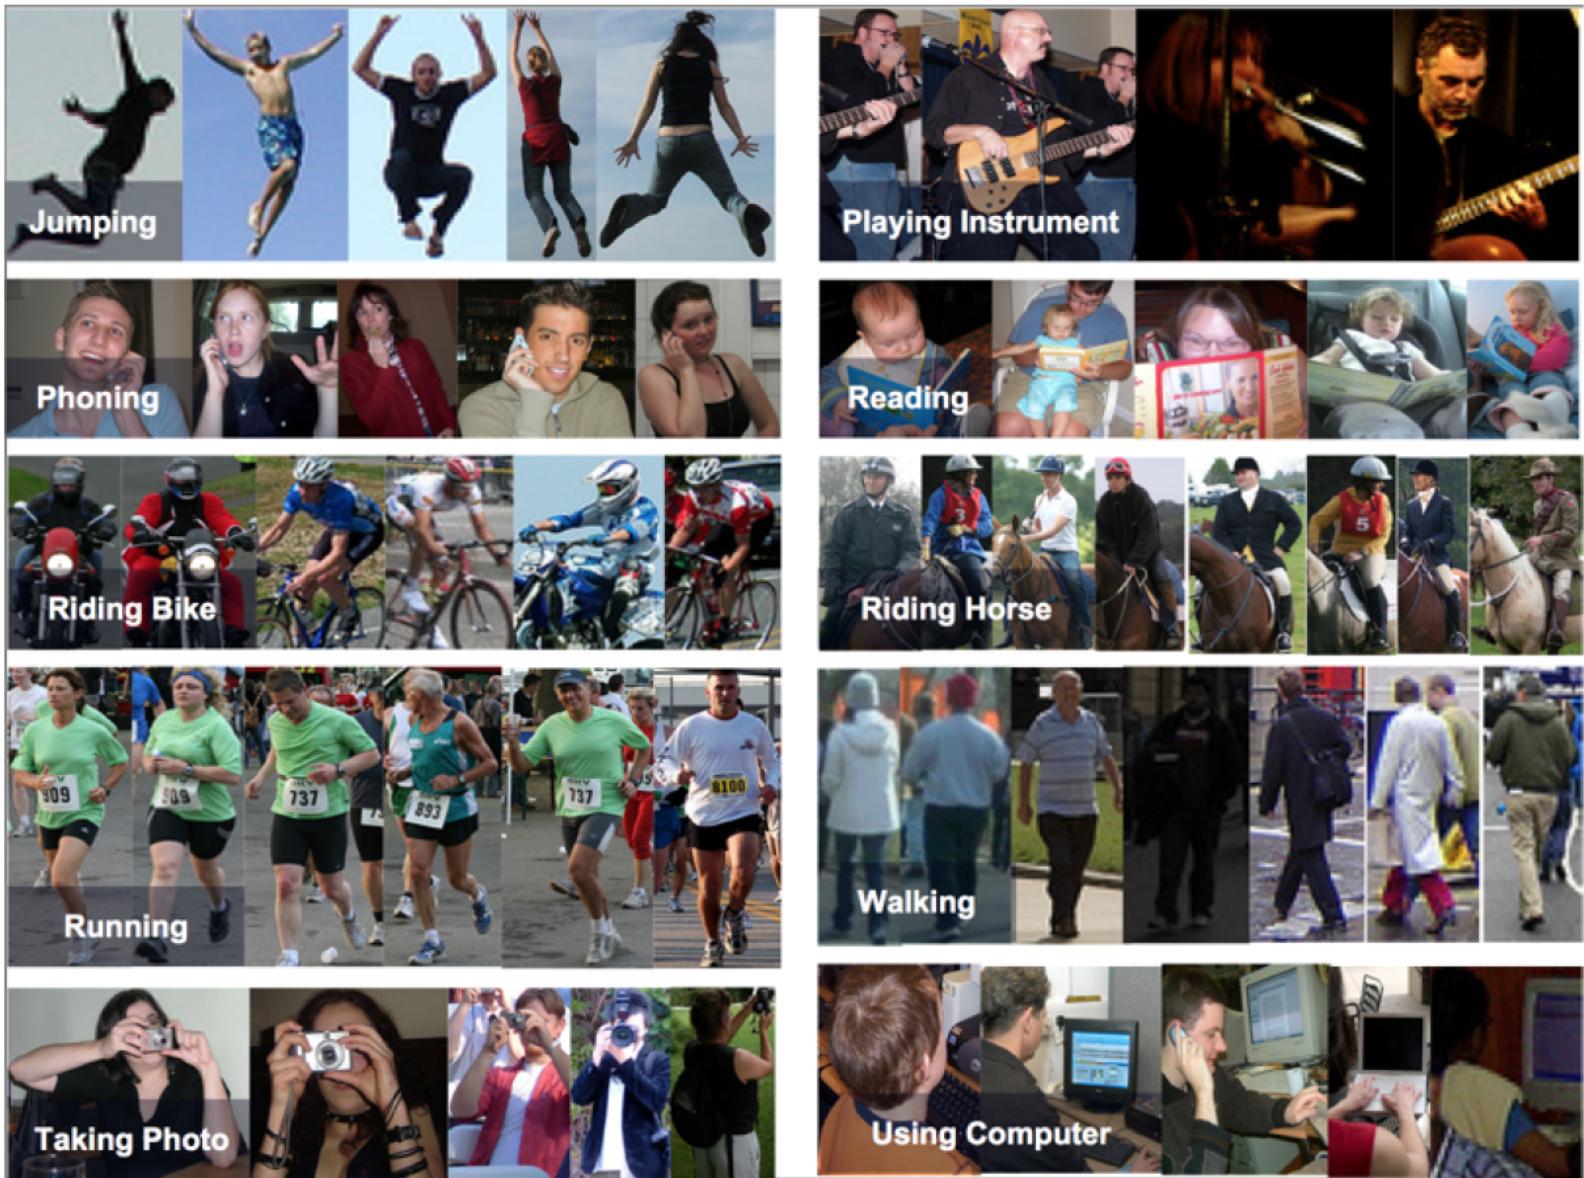
## 4 Applications

- Image classification
- Pose, action detection
- Object detection and segmentation
- Scene labeling - Semantic segmentation
- Object tracking - videos

# Pose estimation - Deeppose



# Action recognition



# Outline

1 Goals

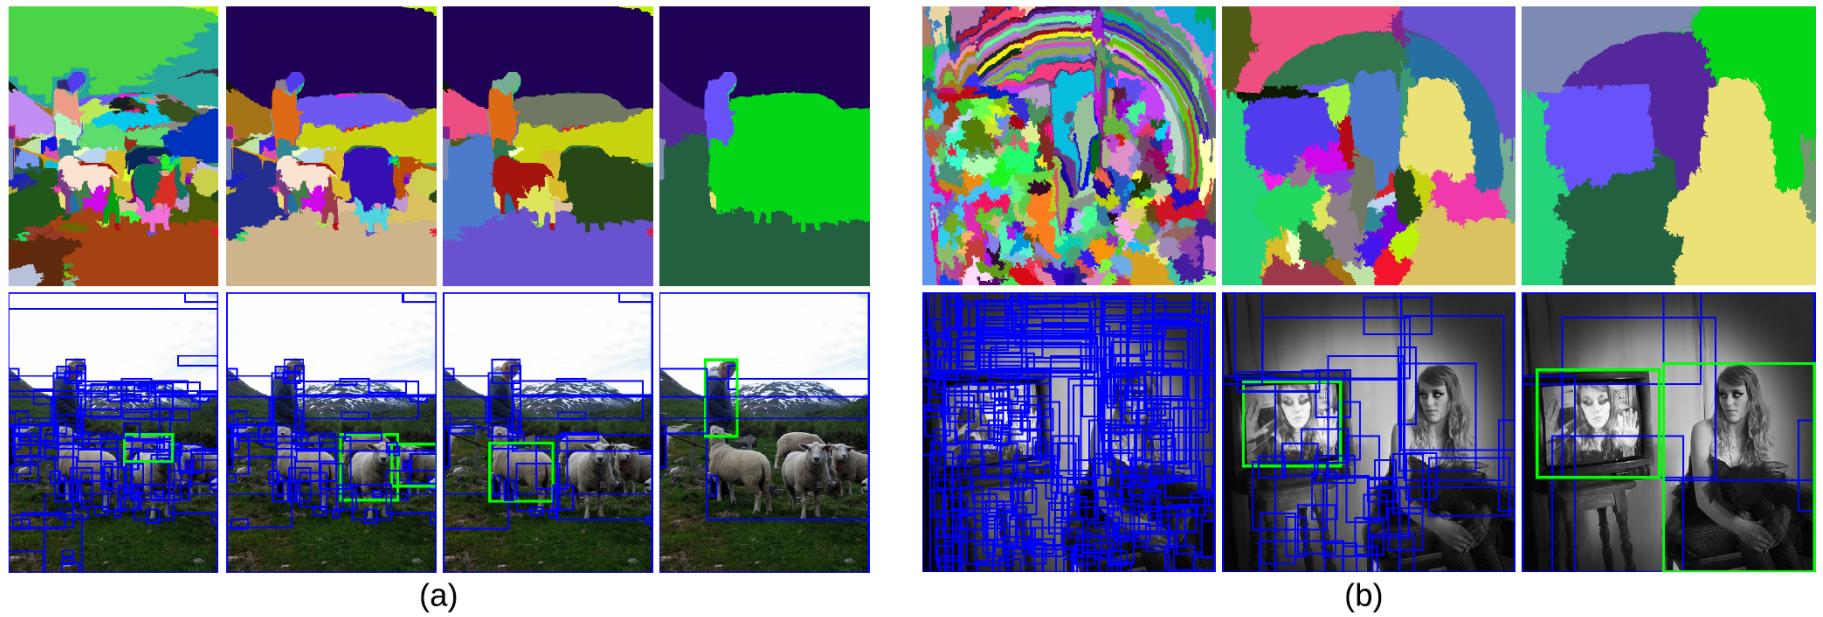
2 The first Neural Network: the Perceptron

3 Convolutional Neural Networks

4 Applications

- Image classification
- Pose, action detection
- **Object detection and segmentation**
- Scene labeling - Semantic segmentation
- Object tracking - videos

# Object detection and segmentation



# Object detection - YOLO, SDD

More recently, YOLO (You Only Look Once) and SSD (Single Shot Detector) allow single pipeline detection that directly predicts class labels.

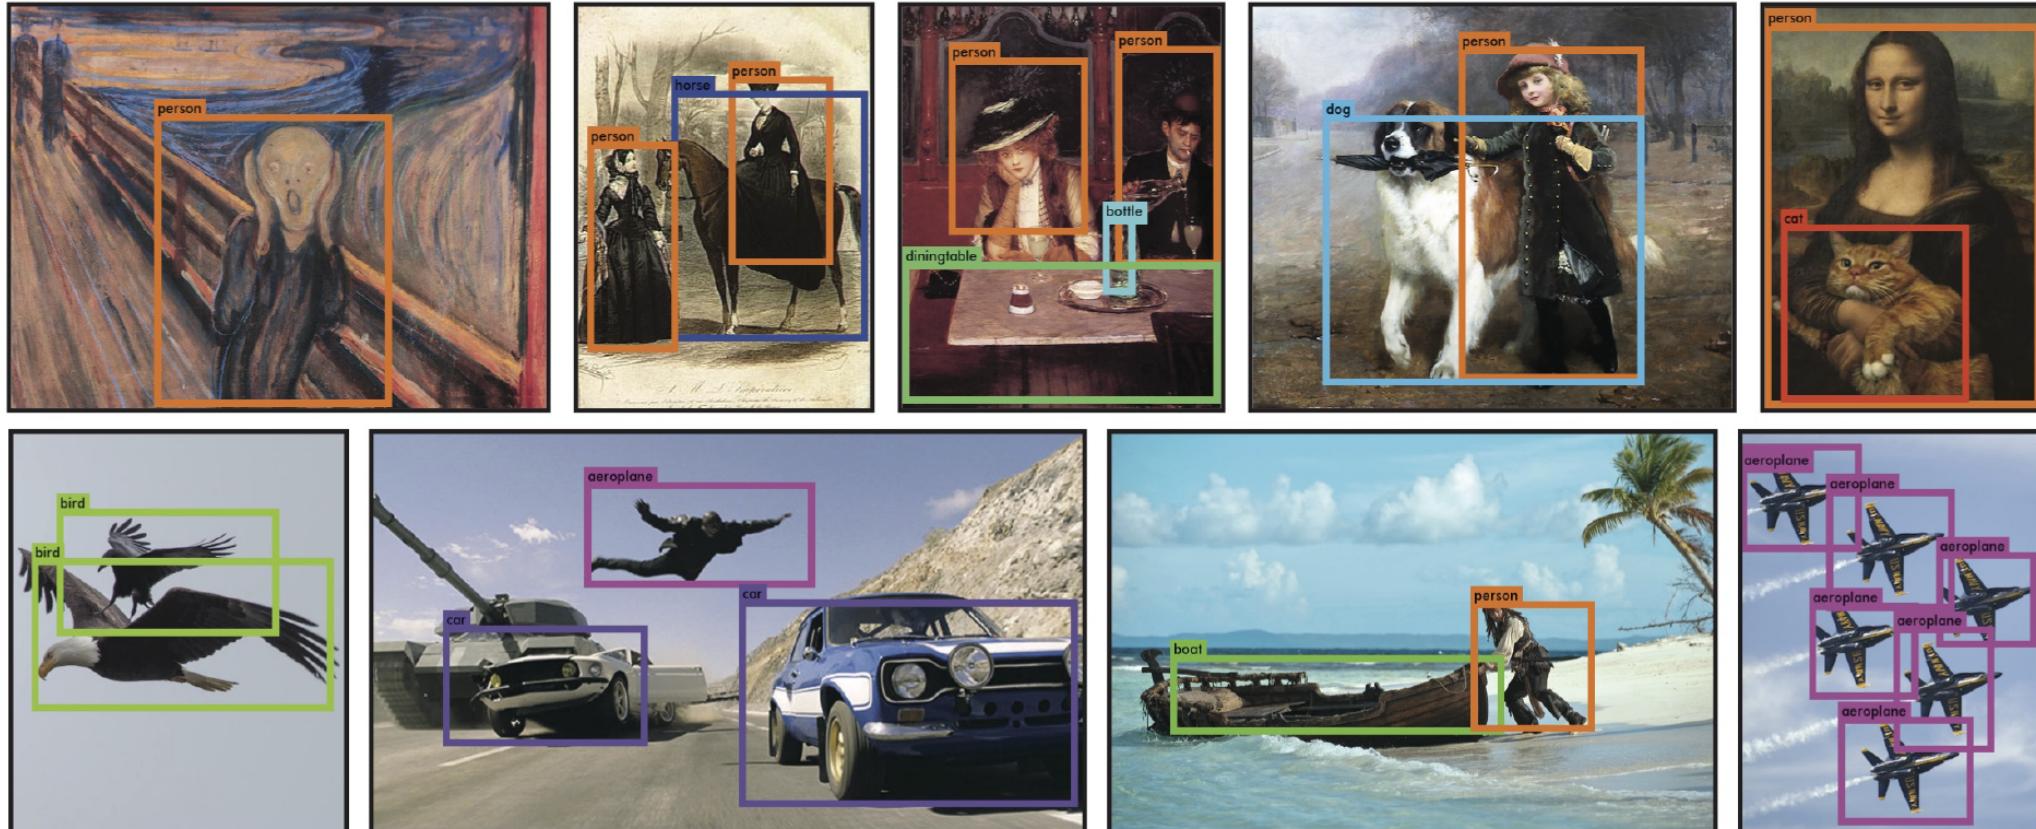
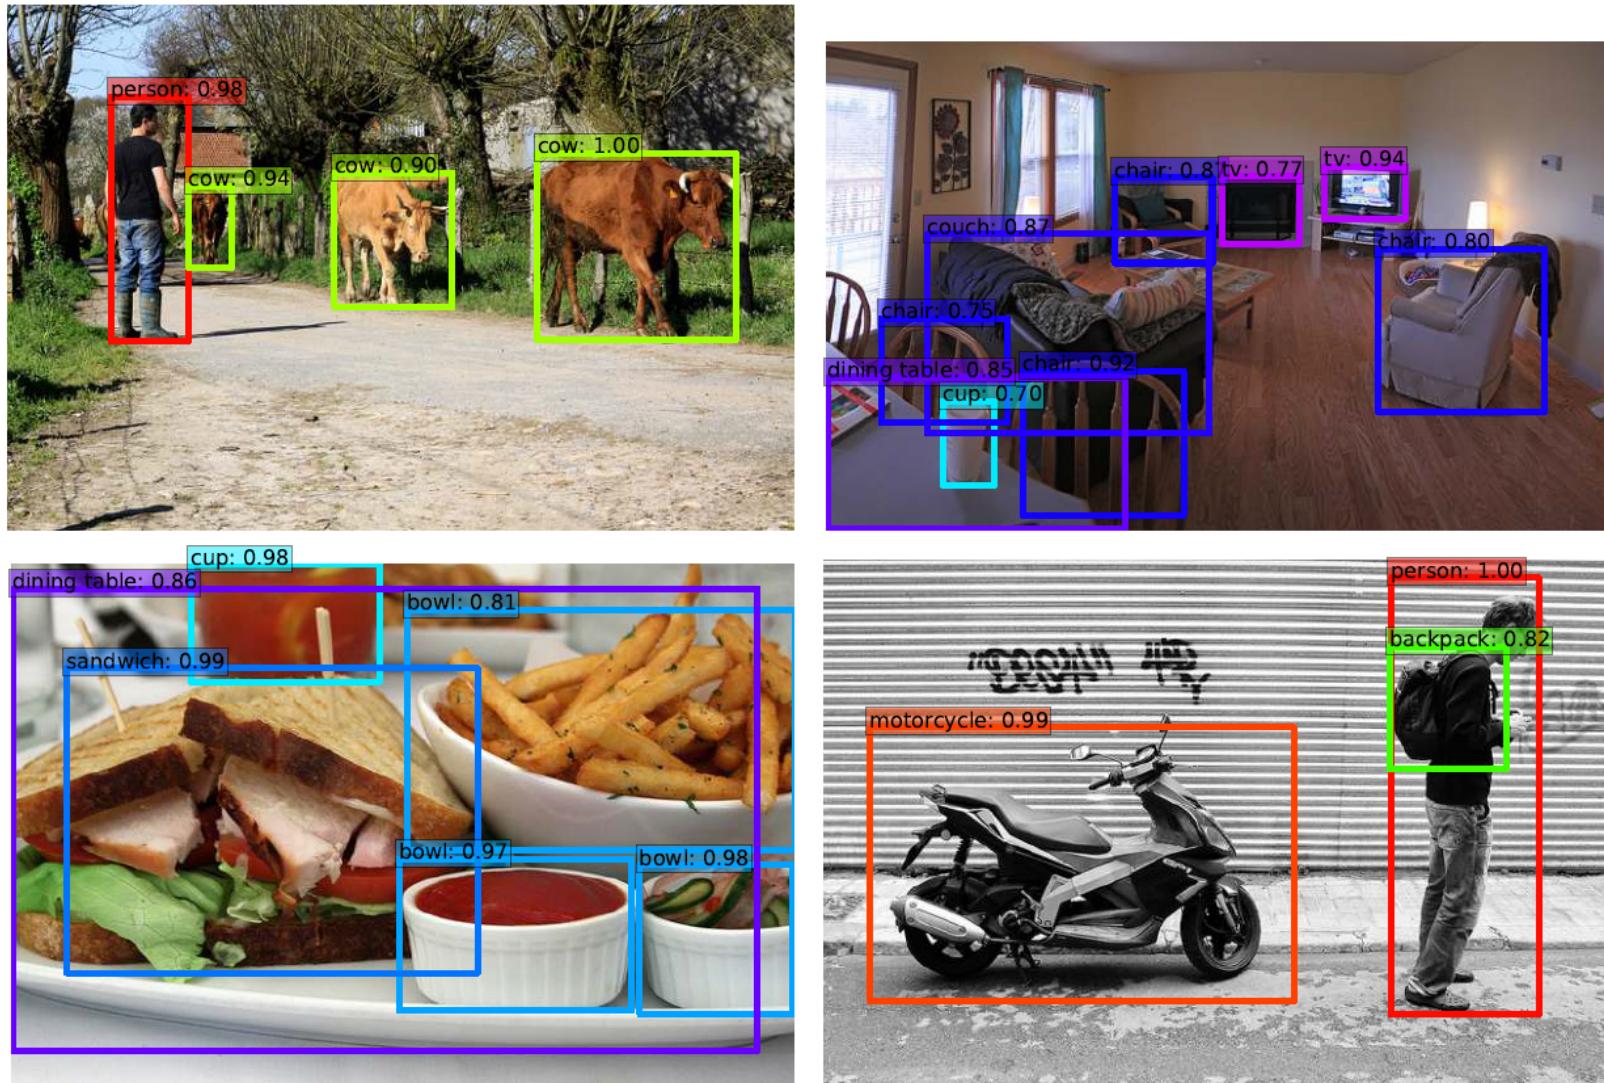


Figure: YOLO results

# SSD



# Outline

1 Goals

2 The first Neural Network: the Perceptron

3 Convolutional Neural Networks

## 4 Applications

- Image classification
- Pose, action detection
- Object detection and segmentation
- **Scene labeling - Semantic segmentation**
- Object tracking - videos

# Scene labeling - DAG-RNN

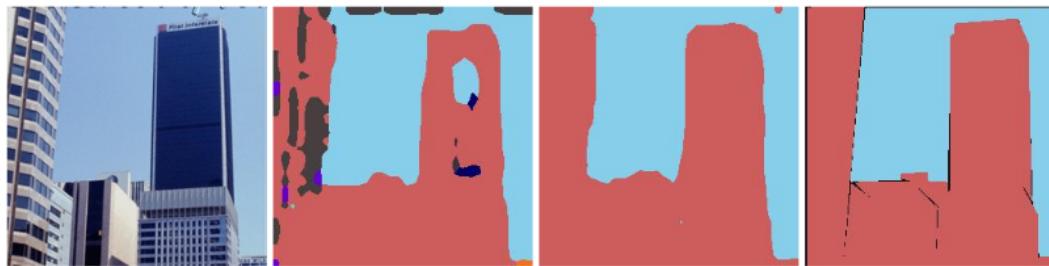


Input Image

CNN

DAG-RNN

Ground Truth



Input Image

CNN

DAG-RNN

Ground Truth

# Outline

1 Goals

2 The first Neural Network: the Perceptron

3 Convolutional Neural Networks

4 Applications

- Image classification
- Pose, action detection
- Object detection and segmentation
- Scene labeling - Semantic segmentation
- Object tracking - videos

# Object tracking

Object Detection in a video:

- YOLO <https://pjreddie.com/darknet/yolo/>
- YOLO - James Bond <https://www.youtube.com/watch?v=V0C3huqHrss>

# Deep Learning - Many many other applications!

- ① Image and video (super-resolution, 3D, Image captioning), medical applications...
- ② Self Driving cars (scene segmentation, etc.)
- ③ In NLP (language, text), automatic translation, voice recognition, text generation, next word completion ... (Virtual assistants)
- ④ Recommendation systems...
- ⑤ For time series, complex datasets too...

But also beyond the supervised settings: using adversarial networks (GANS) for generation of images, faces, voice, etc.

Click on the person who is real.



<http://www.whichfaceisreal.com/>

# Limits and problems

Limits of Deep Learning: → Wooclap!

# Limits and problems

Limits of Deep Learning: → Wooclap!

- Interpretability
- Vulnerability, failures, biases

# Limits and problems

Limits of Deep Learning: → Wooclap!

- Interpretability
- Vulnerability, failures, biases
- Size of the models: Applicability, Ethical, Energy & Ecological concerns (Ai has both a positive and a huge negative carbon footprint)!

# Limits and problems

Limits of Deep Learning: → Wooclap!

- Interpretability
- Vulnerability, failures, biases
- Size of the models: Applicability, Ethical, Energy & Ecological concerns (Ai has both a positive and a huge negative carbon footprint)!
- Theory!

# Limits and problems

Limits of Deep Learning: → Wooclap!

- Interpretability
- Vulnerability, failures, biases
- Size of the models: Applicability, Ethical, Energy & Ecological concerns (Ai has both a positive and a huge negative carbon footprint)!
- Theory!
- Ethical and practical concerns

# Limits and problems

Limits of Deep Learning: → Wooclap!

- Interpretability
- Vulnerability, failures, biases
- Size of the models: Applicability, Ethical, Energy & Ecological concerns (Ai has both a positive and a huge negative carbon footprint)!
- Theory!
- Ethical and practical concerns
- Privacy

# Limits and problems

Limits of Deep Learning: → Wooclap!

- Interpretability
- Vulnerability, failures, biases
- Size of the models: Applicability, Ethical, Energy & Ecological concerns (Ai has both a positive and a huge negative carbon footprint)!
- Theory!
- Ethical and practical concerns
- Privacy
- Source tracking, Ownership, IP

# Limits and problems

Limits of Deep Learning: → Wooclap!

- Interpretability
- Vulnerability, failures, biases
- Size of the models: Applicability, Ethical, Energy & Ecological concerns (Ai has both a positive and a huge negative carbon footprint)!
- Theory!
- Ethical and practical concerns
- Privacy
- Source tracking, Ownership, IP
- Alignment

# Deep Learning: attacks

CNN are not always “robust” to adversarial attacks!

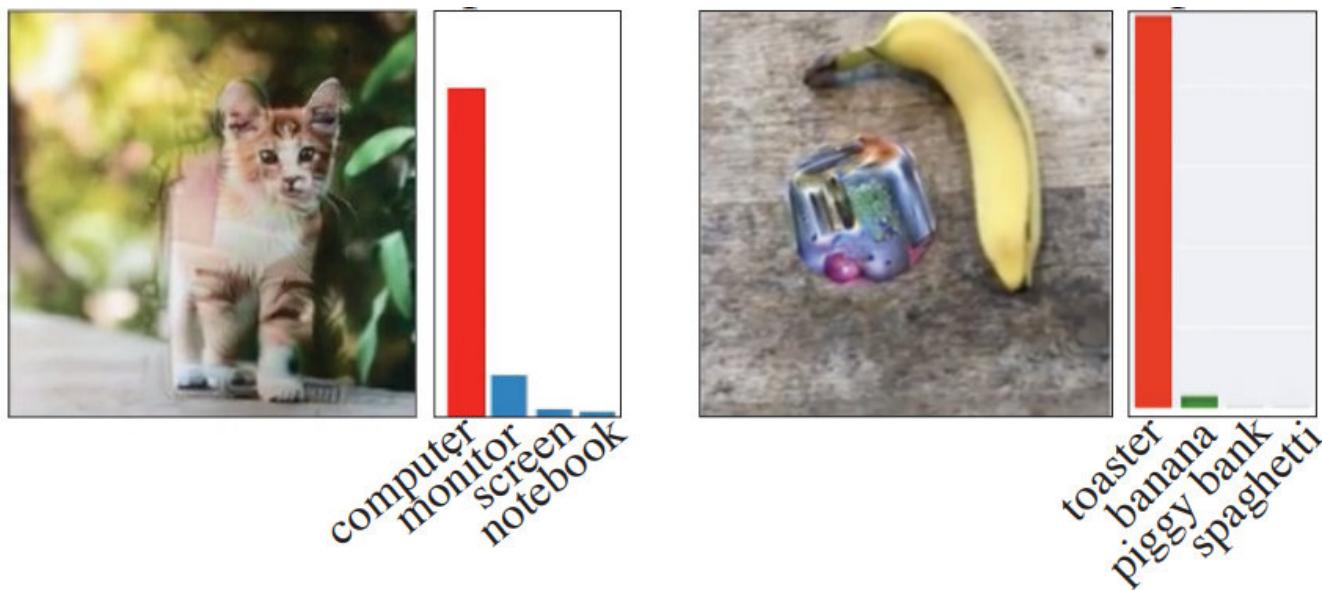


Figure: What happened here??

# Deep Learning: attacks

CNN are not always “robust” to adversarial attacks!



dog

+noise

ostrich

Adding a small well chosen noise can completely fool a CNN!

↗ Can we trust deep networks on cars, medical applications, planes...?

# Deep Learning: failures

## ① Biases in Learning sets that are transferred

- ① Chatbot becomes racist
- ② Google apologises for Photos app's racist blunder

# Deep Learning: failures

- ① Biases in Learning sets that are transferred
  - ① Chatbot becomes racist
  - ② Google apologises for Photos app's racist blunder
- ② Privacy leaks:
  - ① The size of the models often enable memorization of the training sets.
  - ② Need for guarantees and regulation

# Deep Learning: failures

- ① Biases in Learning sets that are transferred
  - ① Chatbot becomes racist
  - ② Google apologises for Photos app's racist blunder
- ② Privacy leaks:
  - ① The size of the models often enable memorization of the training sets.
  - ② Need for guarantees and regulation
- ③ False information/prediction, hallucinations...

# Deep Learning: failures

- ① Biases in Learning sets that are transferred
  - ① Chatbot becomes racist
  - ② Google apologises for Photos app's racist blunder
- ② Privacy leaks:
  - ① The size of the models often enable memorization of the training sets.
  - ② Need for guarantees and regulation
- ③ False information/prediction, hallucinations...
- ④ Overall, lack of control on the model behavior

PIXELS • VIE PRIVÉE

## Nouvelle plainte contre OpenAI pour infraction au RGPD

ChatGPT enfreint le Règlement général sur la protection des données, estime l'association autrichienne None of Your Business (NOYB), car il peut donner de fausses informations concernant un individu, sans possibilité de les faire corriger ou supprimer.

Le Monde avec AFP  
Publié hier à 10h44 - ⏱ Lecture 1 min.

Figure: Link

# Alignment



FAST COMPANY

CO.DESIGN TECH WORK LIFE NEWS IMPACT PODCASTS VIDEO RECOMMENDER INNOVATION FESTI

11-06-17 | PLATFORM WARS

## Netflix CEO Reed Hastings: Sleep Is Our Competition

For Netflix, the battle for domination goes far beyond which TV remote to pick up.



CEO Reed Hastings [Photo: Mondileinchen/Wikimedia Commons]

# Deep Learning: societal risk and questionable applications

Are some applications just bad (they all already exist :( )?

- Deep Fakes: insert someone in a video.
- Underground exploration: finding new fossil resources.
- Voice imitation: what if I cannot check who is calling me.
- Military applications + automatic weapons.
- Mass surveillance.

The likelihood of each of those applications being largely deployed increases every year.

# Deep Learning: societal risk and questionable applications

Are some applications just bad (they all already exist :( )?

- Deep Fakes: insert someone in a video.
- Underground exploration: finding new fossil resources.
- Voice imitation: what if I cannot check who is calling me.
- Military applications + automatic weapons.
- Mass surveillance.

The likelihood of each of those applications being largely deployed increases every year.

# Deep Learning: questionable applications

A challenging problem: [chain of responsibility](#). The same tools are used for positive and negative applications, e.g.:

- mass surveillance
- tumor recognition

are both based on image recognition, and the same ML techniques.

# Deep Learning: questionable applications

A challenging problem: [chain of responsibility](#). The same tools are used for positive and negative applications, e.g.:

- mass surveillance
- tumor recognition

are both based on image recognition, and the same ML techniques.

[Good news!](#) the research community is very aware of the situation.

Some references:

- [Asilomar AI Principles](#) [Ai principles](#)
- [AI Act](#)
- [European guidelines](#)
- [AI for Good](#)



We need you to be aware and active on those challenges!!

# Conclusion: Deep Learning in one slide

- **How does it work:**

- ▶ Automatically learn representations of observations
- ▶ Learn highly non-linear models.

- **What does it require:**

- ▶ Large datasets with structure
- ▶ Computational power

- **Why now:**

- ▶ Combination of the 2 points above
- ▶ investment!

- **Some Applications**

- ▶ Image classification; object / face recognition
- ▶ Self driving cars
- ▶ Automatic Translation, Information extraction
- ▶ Caption Generation
- ▶ Ads, recommendation systems, etc.

**Goals:** Understanding core concepts of Deep Learning.

- When and how it works on paper (data, models, architecture)
- How to implement a simple neural network with Python.
- Overview of some of the main applications and challenges.















You cannot vote anymore

Which of the following things need to be chosen (hyperparameters of the NN)?

1 neuron 13% 5 13%  
 2 activation functions 95% 37 ✓ 2 activation functions 95% 37 ✓  
 3 number of neurons per layer 92% 36 ✓ 3 number of neurons per layer 92% 36 ✓  
 4 weights 18% 7 18%  
 5 input size 46% 18 46%  
 6 number of layers 100% 39 ✓ 6 number of layers 100% 39 ✓  
 7 optimization algorithm 38% 15 ✓ 7 optimization algorithm 38% 15 ✓  
 8 biases 8% 3 8%

**wooclap** Questions 1 / 10 Messages 🔒 100 % 🔍 Exit Tout afficher

california.ipynb

You cannot vote anymore

Which of the following things need to be learned ?

1 neuron 9% 3 9%  
 2 activation functions 3% 1 3%  
 3 number of neurons per layer 3% 1 3%  
 4 weights 97% 32 ✓ 4 weights 97% 32 ✓  
 5 input size 9% 3 9%  
 6 number of layers 0% 0 0%  
 7 optimization algorithm 9% 3 9%  
 8 biases 76% 25 ✓ 8 biases 76% 25 ✓

**wooclap** Questions 2 / 10 Messages 🔒 100 % 🔍 Exit Tout afficher

california.ipynb

You cannot vote anymore

What happens here?

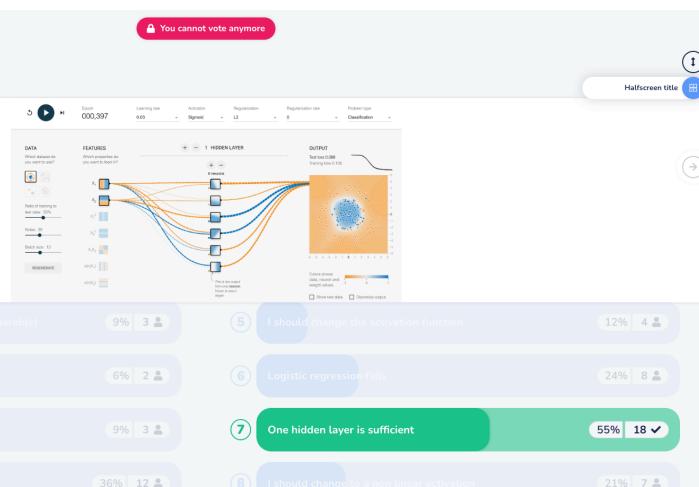
1 Logistic Regression works (data is linearly separable) 58% 15 ✓ 1 Logistic Regression works (data is linearly separable) 58% 15 ✓  
 2 I should use more neurons 0% 0 0%  
 3 The network overfits 4% 1 4%  
 4 I should add more layers 0% 0 0%  
 5 I should change the activation function 4% 1 4%  
 6 Logistic regression fails 0% 0 0%  
 7 One hidden layer is sufficient 31% 8 31%  
 8 I should change to a non linear activation 4% 1 4%

1 Logistic Regression works (data is linearly separable) 13% 3 13%  
 2 I should use more neurons 46% 11 46%  
 3 The network overfits 8% 2 8%  
 4 I should add more layers 54% 13 54%  
 5 I should change the activation function 21% 5 21%  
 6 Logistic regression fails 75% 18 75%  
 7 One hidden layer is sufficient 4% 1 4%  
 8 I should change to a non linear activation 25% 6 25%

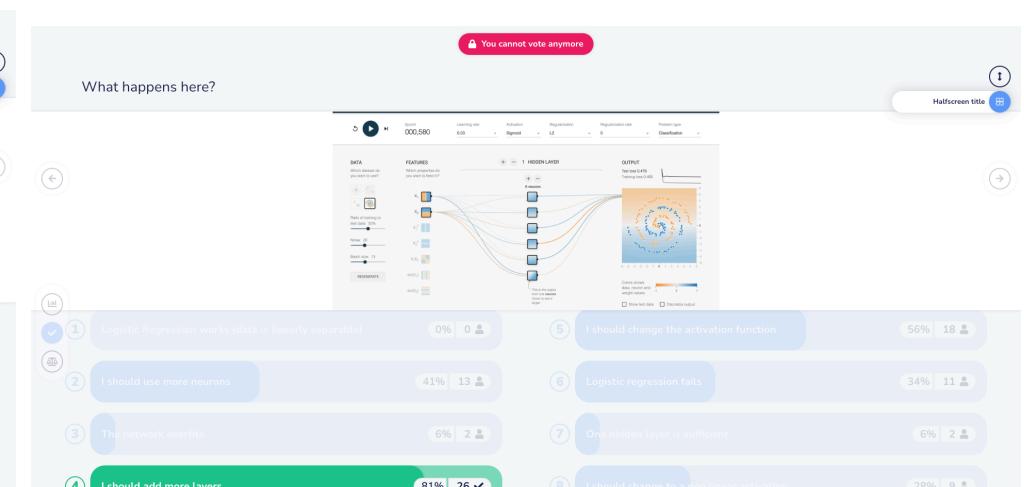
**wooclap** Questions 3 / 10 Messages 🔒 100 % 🔍 Exit Tout afficher

california.ipynb

What happens here?



What happens here?



What happens here?

You cannot vote anymore

This screenshot shows a Wooclap poll titled "What happens here?". At the top, there's a message "You cannot vote anymore". Below it is a neural network diagram with 4 hidden layers. The poll has 8 options, each with a progress bar:

- ① Logistic Regression works (data is linearly separable) - 0% (0 people)
- ② I should use more neurons - 12% (3 people)
- ③ The network overfits - 8% (2 people)
- ④ I should add more layers - 15% (4 people)
- ⑤ I should change the activation function - 85% (22 people) [highlighted in green]
- ⑥ Logistic regression fails - 12% (3 people)
- ⑦ One hidden layer is sufficient - 0% (0 people)
- ⑧ I should change to a non-linear activation - 0% (0 people)

At the bottom, the Wooclap interface shows 7 questions, 100% messages, and 26 / 45 participants.

What happens here?

You cannot vote anymore

This screenshot shows a Wooclap poll titled "What happens here?". At the top, there's a message "You cannot vote anymore". Below it is a neural network diagram with 4 hidden layers. The poll has 8 options, each with a progress bar:

- ① Logistic Regression works (data is linearly separable) - 0% (0 people)
- ② I should use more neurons - 4% (1 person)
- ③ The network underfits - 4% (1 person)
- ④ I should add more layers - 0% (0 people)
- ⑤ I should change the activation function - 83% (19 people) [highlighted in green]
- ⑥ Logistic regression fails - 0% (0 people)
- ⑦ One hidden layer is sufficient - 0% (0 people)
- ⑧ I should change to a non-linear activation - 9% (2 people)

At the bottom, the Wooclap interface shows 7 questions, 100% messages, and 23 / 45 participants.

What happens here?

You cannot vote anymore

This screenshot shows a Wooclap poll titled "What happens here?". At the top, there's a message "You cannot vote anymore". Below it is a neural network diagram with 6 hidden layers. The poll has 8 options, each with a progress bar:

- ① Logistic Regression works (data is linearly separable) - 0% (0 people)
- ② I should use more neurons - 0% (0 people)
- ③ The network overfits - 0% (0 people)
- ④ I should add more layers - 0% (0 people)
- ⑤ I should change the activation function - 0% (0 people)
- ⑥ Logistic regression fails - 0% (0 people)
- ⑦ It works! - 100% (21 people) [highlighted in green]
- ⑧ I should change to a non-linear activation - 0% (0 people)

At the bottom, the Wooclap interface shows 10 questions, 100% messages, and 21 / 45 participants.



